

CONTENTS

TOPICS COVERED	PAGE NO
1. INTRODUCTION	1
1.1 Purpose	2
1.2 Scope	2
1.3 Need for System	2
1.3.1 Existing System	4
1.3.2 Proposed System	5
1.4 System Architecture	6
2. SOFTWARE REQUIREMENT SPECIFICATION AND ANALYSIS	8
2.1 Product Perspective	8
2.2 Product Function	8
2.3 User Characteristics	10
2.4 Modules	11
2.5 Functional and Non-Functional Requirements	12
2.5.1 Functional Requirements	12
2.5.2 Non-Functional Requirements	14
2.6 System Specifications	15
2.6.1 Hardware Requirements	16
2.6.2 Software Requirements	16
2.7 Software Development Life Cycle	17
2.7.1 Waterfall Model	18
2.8 System Study	20
2.9 Methodology and Algorithms	22
2.10 Technologies Used	23

3. SYSTEM DESIGN	29
3.1 ER-Diagram	29
3.2 Normalization	42
3.3 Data Flow Diagram	46
3.4 UML Diagrams	47
4. TESTING	58
4.1 Introduction	59
4.2 Types of Testing	60
4.3 Software Testing Strategy	64
4.4 Test Cases	65
5. IMPLEMENTATION	
5.1 Sample Screens	67

CONCLUSION

Project POs Mapping

Sustainable Development Goals

Justification BIBIOGRAPHY

Appendix – A

 References

Appendix – B

 Glossary

Appendix – C

 List of Table

 List of Figures

 List of Screens

Appendix – D

 Help Document

 Base Paper

ABSTRACT

Abstract

In modern cloud computing, the need for flexible and scalable orchestration of services, combined with robust security measures, is paramount. Here, I propose an innovative approach for managing secure cloud bursting in Kubernetes, combining Attribute-Based Encryption (ABE) with Kubernetes labeling. Our model addresses the challenges of complexity, cost, and data protection compliance by leveraging both Kubernetes and ABE. I introduce an attribute-based bursting component that uses Kubernetes labels for orchestration, and an encryption component that employs ABE for data protection. This unified management model ensures data confidentiality while enabling efficient cloud bursting. Our approach combines the strengths of label- based orchestration with fine-grained encryption, providing a technologically advanced yet user-friendly solution for secure cloud bursting. I present a proof-of-concept implementation that demonstrates the feasibility and effectiveness of our model. Our approach offers a unified solution that complies with security and privacy laws while meeting the needs of contemporary cloud- based systems.

INTRODUCTION

1. Introduction

In recent years, the proliferation of virtualization and containerization technologies has led to a significant increase in the complexity of distributed systems, as cloud computing systems. As organizations strive to achieve efficient resource management and scalability, Kubernetes has emerged as the most popular solution for orchestrating resources in such complex systems. For example, it is integrated into platforms such as Amazon Elastic Kubernetes Service, Google Kubernetes Engine, Azure Kubernetes Service, IBM Cloud Kubernetes Service, and Oracle Container Engine for Kubernetes. It aims at exploring the challenges associated with cloud bursting, which allows private cloud services to use public cloud resources when local resources are exhausted or for any other reasons. Specifically, we address the configuration issues associated with service request load management over a hybrid cloud system including both private and public components. The orchestration of such a heterogeneous system presents a number of challenges, such as optimal management of typically large volume of resources, variable operating conditions, security issues, and compliance with local regulations. In order to address these challenges, resorting to attribute-based management policies are regarded as a valid approach. Attributes are intrinsic characteristics or properties related to the entities, workload, or resources to manage.

1.1 Purpose

The purpose of improving security and privacy attribute-based data sharing in cloud computing is to address the challenges associated with securely sharing sensitive data in cloud environments while ensuring compliance with privacy regulations and organizational policies. This system aims to enhance security measures and privacy protections by secure cloud bursting in Kubernetes, combining Attribute-Based Encryption (ABE) mechanisms, which allow data owners to define access policies based on specific attributes of users and data. By improving the security and privacy of data sharing in cloud computing, this system aims to foster trust among users and organizations, promote data confidentiality, and mitigate the risk of unauthorized access or data breaches.

1.2 Scope

The scope of the system for improving security and privacy attribute-based data sharing in cloud computing encompasses various aspects related to access control, encryption, and privacy-preserving techniques. This includes the development of secure cloud bursting in Kubernetes, combining Attribute-Based Encryption (ABE) models and policies for defining fine-grained access controls based on user attributes such as roles, permissions, or other contextual factors. Additionally, the system may involve the integration of encryption mechanisms to protect data confidentiality during transmission and storage in the cloud. Furthermore, considerations for auditability, scalability, and interoperability across different cloud platforms and data sharing scenarios are essential aspects of the system's scope.

1.3. Need for System

Cloud computing environments often host large volumes of sensitive data, including personal, financial, and proprietary information. Enhancing security and privacy measures is essential to protect this data from unauthorized access, data breaches, and privacy violations. Cloud storage

has become an integral part of data management, but ensuring data integrity associated with third- party auditors. This study addresses key concerns system information includes details about users, workloads, cluster configurations, policies, resource usage, and security attributes. This information is essential for making accurate, automated decisions about when to burst workloads from a private cluster to a public one. Without a proper understanding of system information, it would be impossible to apply the correct policies, enforce secure communication, or ensure compliance with data protection standards. Moreover, knowing the system information helps in defining the functional and non- functional requirements, creating effective test cases, designing database models, and ensuring smooth integration between modules. It allows the project team to anticipate challenges like overloading, unauthorized access, or policy violations and handle them efficiently. Thus, the need for system information is fundamental to building a well-structured, secure, and responsive cloud bursting solution. System information is necessary for monitoring and auditing. Logs, metrics, and status updates form the basis for tracking workload movements, verifying that policies are applied correctly, and investigating any anomalies. This traceability is crucial in environments where data sensitivity and compliance requirements are high. System information also plays a key role in resource optimization. By understanding cluster load, performance trends, and usage patterns, the system can make smarter decisions about scaling and bursting—minimizing cost while maximizing availability. Such as: Eliminating Traditional auditing methods rely on TPAs, which may manipulate, alter, or leak sensitive data. Block chain provides a decentralized and tamper-proof alternative. The quad Merkle hash tree optimizes computation and storage, reducing the overhead associated with conventional integrity verification techniques.

1.3.1 ExistingSystem

Old Systems designed to support unshared multi-owner setting, and cannot be directly applied in the shared multi-owner setting. As existing system does not provide server data for which multi-keyword and multi owner search. It is difficult to identify malicious users who leak the secret keys when more than one data user has the same subset of attributes.

Disadvantages

- Existing System provides only SingleOwner Sharing.
- Existing System takes higher time for getting results from server.
- Privacy issues are raised in submission of query
- Repeatedly send the query, user feel like complex.
- In the existing work, while CPABKS schemes can achieve one- to- many encryption and support expressive access control
- Not capable of identifying data users leaking the secret keys if the ‘culprits’ have the same subset of attributes as other honest data users.
- Data Leak sensitive information cannot be controlled.

It is difficult to identify malicious users who leak the secret keys when more than one data user has the same subset of attributes. The need for a system arises from the goal to solve specific problems, improve efficiency, and ensure consistency in operations. In both manual and automated environments, systems help manage complex tasks, reduce errors, and save time. They provide a structured approach to handling data, processes, and user interactions, allowing organizations to function smoothly and meet their objectives effectively.

1.3.2 ProposedSystem

A framework of secure cloud bursting in Kubernetes, combining Attribute-Based Encryption (ABE) with efficient authority identification is proposed, which guarantees the data confidentiality and protects the user personal privacy as well. In order to avoid unnecessary computations of users in decryption algorithm, we design an authority identification method, which can help the user verify whether he/she is an authorized one and decrypts successfully. The proposed scheme achieves constant private key size, which is independent of user's attribute number. It reduces the cost of transmission and storage. In addition, a compact security analysis by using a sequence of hybrid games is given to show the proposed scheme of how to achieve anonymity, which is lacking in most of the existing works.

Advantages

- Data confidentiality
- Privacy preserving
- Efficient decryption test

The proposed system introduces an attribute-based management approach to enhance the security and efficiency of Kubernetes cloud bursting. It uses attribute-based access control (ABAC) to manage and enforce access to resources based on dynamic attributes such as user roles, department, location, and time. This ensures that only authorized users or services can trigger or access cloud resources during the bursting process. When the local Kubernetes cluster reaches its resource limits, the system securely bursts workloads to the public cloud. It maintains data confidentiality and integrity through encryption and secure transmission methods. The policies used are dynamic and adapt automatically to changing workloads and user contexts, enabling real-time decision-making without manual intervention. This design ensures optimal resource utilization across on-premises and cloud environments, providing a scalable, secure, and policy-driven cloud bursting mechanism.

1.4 Architecture

Attribute-based management of a secure Kubernetes cloud bursting system relies on defining policies and controls based on dynamic attributes such as user roles, workload characteristics, and security contexts. Security is enforced using Attribute-Based Access Control (ABAC), integrating with IAM solutions like AWS IAM, Azure AD, and to ensure fine-grained access control. A Zero Trust model is applied, verifying identity, attributes, and context before granting access. Data security is ensured through encryption at rest and in transit, using TLS and AES-256 standards. The architecture leverages a hybrid cloud model, where an on-premises Kubernetes cluster dynamically extends to public cloud environments when demand exceeds local capacity. Kubernetes Cluster Auto scaler and KEDA facilitate dynamic resource provisioning, scaling workloads based on real-time demand. Service mesh technologies like Istio, Linkerd, or Cilium enable secure, policy-driven communication between services. Here architecture typically includes two Kubernetes environments: one private and one public. When the private cluster reaches certain conditions, such as running out of resources or hitting defined thresholds, the system evaluates whether it's safe and appropriate to shift some of the workload to the public cloud. This decision-making is driven by policies that check attributes in real time. For example, if a workload is tagged as non-sensitive and the user role is trusted, then it might be allowed to run in the public cloud temporarily. A secure communication layer ensures that any data sent between the clusters is encrypted and protected. At the same time, a monitoring system keeps track of everything, logging what was moved, when it was moved, and why. This helps maintain transparency, auditability, and compliance with security standards. The overall architecture balances flexibility and control, allowing organizations to scale when needed while keeping sensitive operations tightly governed.

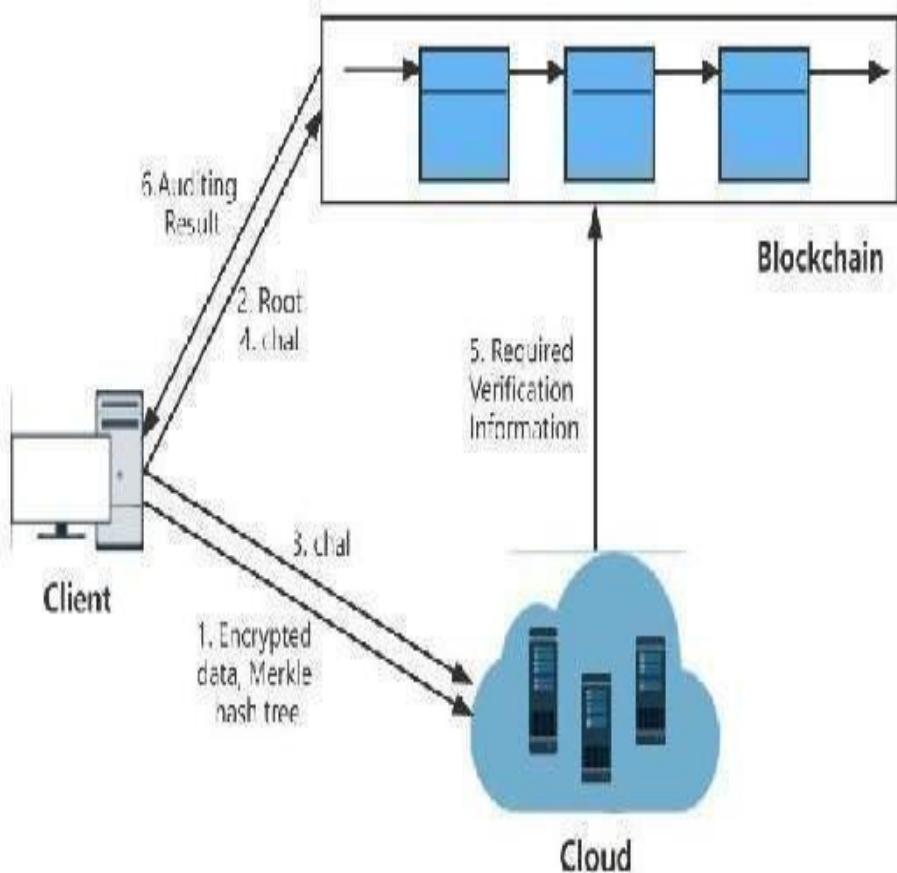


Fig: Architecture

After viewing the architecture diagram of attribute-based management of secure Kubernetes cloud bursting, we can understand how it dynamically controls workload distribution between private and public clusters.

2. Software Requirement Analysis and Specification

2.1 Product Perspective

An essential aspect necessitates assurance, which involves bridging the gap between product management and development. It entails thoroughly defining a product in terms of its associated requirements, encompassing all the necessary specifications that must be explicitly described and continuously available.



Fig. Product Perceptive

2.2 Product Function

In product management, the product requirements specification serves as the central tool, fulfilling the following purposes:

- It articulates how product management addresses partner needs and requests.
- It communicates to development teams what the product or new features to be developed will entail.
- The accompanying figure illustrates these relationships, focusing on the product perspective.

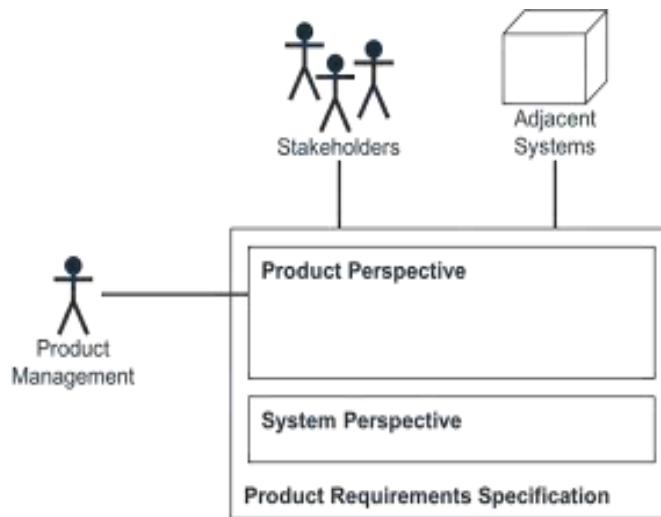


Fig. Characteristics of product perspective

The key characteristics of the product perspective include:

- Describes the external view of the product.
- Answers the question "What"...
- Represents what the product is.
- Describes what the product does.
- Addresses both the business and usage aspects of the product.
- Utilizes problem domain language and concepts.
- Presents a "Black Box" view, focusing on the external behavior and functionality of the product without delving into its internal workings.

2.3 User Characteristics

This section outlines the characteristics of users in three distinct contexts: the end user who interacts with the machine translation system; the end user of the final product of the translation process, which may include post-editing; the organization producing the machine translation system.

Constraints

Pragmatic requirements are product features or functionalities that developers must implement to enable users to accomplish their tasks. It is essential to clearly define them for both the development team and the stakeholders. Typically, pragmatic requirements describe system behavior under unambiguous conditions. For example: Functional requirements can be categorized by various criteria. For instance, they can be grouped based on the roles a given component should perform in the final product. Naturally, these requirements may vary depending on the product under development, but examples of functional requirements:

- Authorization levels
- Compliance with regulations or standards
- External interfaces
- Transaction processing
- Reporting
- Business rules, etc.

They can be technical, such as limited processing power or storage capacity. They might be economic, like budget limits or cost-effectiveness requirements. Legal or regulatory constraints may apply in areas such as data privacy or security compliance. Time constraints can also impact project deadlines and delivery schedules. Constraints are essential to identify early in the project life cycle, as they help shape realistic goals and guide decision-making.

2.4 Modules

Data Owner

In this module, the data owner uploads their data with its chunks in the cloud server. For the security purpose the data owner encrypts the data file and then store in the cloud. The Data owner can have capable of manipulating the encrypted data file.

Cloud Servers

The cloud service provider manages a cloud to provide data storage service. Data owners encrypt their data files and store them in the cloud for sharing with data consumers. To access the shared data files, data consumer.

Data Consumer (End User)

The Cloud User who has a large amount of data to be stored is integrated and downloaded.

2.5 Functional Requirements and Non-Functional Requirements

2.5.1 Functional Requirements

Functional requirements delineate the intended actions of the system and can be segmented into various categories:

- Inputs: Specifies the data the system should accept.
- Outputs: Defines the results or data the system should generate.
- Data Storage: Determines the data that must be stored by the system.

The system should be able to monitor the resource usage within the private Kubernetes environment and detect when workload thresholds are reached. It must evaluate various attributes in real time, such as the type of workload, the user's role, and the classification level of the data involved. Based on predefined attribute-based policies, it should make automated decisions about whether and when to shift workloads to a public cloud cluster. The system should then securely deploy the selected containers to the public environment, ensuring proper encryption and authentication throughout the process. It must also support dynamic scaling, meaning it should automatically expand or shrink public cloud usage depending on changing demands. Communication between private and public clusters must remain secure and uninterrupted. Furthermore, it should track and log all operations, including which workloads were moved, why they were moved, and by whom, allowing for full audit and compliance checks. Finally, it should provide a clear interface for administrators to define, update, and manage attribute-based policies at any time.

The system should provide real-time alerts and notifications to administrators when bursting events occur or when policy violations are detected. It must be capable of rolling back workloads from the public cloud to the private cluster seamlessly when demand decreases or if security concerns arise. Lastly, the solution should be extensible, allowing new attributes and policies to be added easily as organizational requirements evolve.

- Defining input data requirements.
- Organizing and coding data.
- Providing user guidance.
- Implementing input validations and error handling.

Functional requirements describe the specific operations, features, and behaviors a system must support to fulfill its purpose. They outline what the system should do in response to various inputs or interactions, and define the rules under which the system operates. These requirements are often closely tied to user needs and business goals, ensuring that the system delivers the correct outputs, manages data appropriately, and supports necessary user tasks. They form the foundation for designing the system's functionality and serve as a basis for validation and testing.

2.5.2 Non-functional Requirements

Non-functional requirements (NFRs) address fundamental aspects of software systems, focusing on qualities beyond specific functionalities. Failure to address NFRs adequately can result in user dissatisfaction, software conflicts, and increased time and cost to rectify issues. Types of Non-functional Requirements:

- **Scalability:** The system's ability to handle increasing loads without compromising performance.
- **Reliability:** The system's ability to perform consistently and predictably under various conditions.
- **Regulatory Compliance:** Ensuring the system adheres to relevant regulations or standards.
- **Maintainability:** Ease of maintaining and updating the system over time.

Non-functional requirements include security, ensuring data confidentiality and secure access through encryption and attribute-based controls; scalability, allowing the system to handle increased workloads by dynamically bursting to cloud environments; availability, maintaining consistent access to services during peak loads or failures; performance, ensuring fast and efficient resource allocation and workload migration; reliability, with the system functioning correctly under various conditions without failure; and maintainability, allowing easy updates and modifications to policies and configurations without disrupting services. These requirements include performance aspects like response time and scalability, usability factors such as user interface friendliness, reliability measures like system uptime, and maintainability which covers ease of updates and debugging. Security is another important area, ensuring that the system protects data and resists unauthorized access. Unlike functional requirements that describe specific behaviors or functions, non-functional requirements shape the user experience and influence the system's architecture, design decisions, and overall success.

2.6 System Specification

The system must support Kubernetes clusters both in private (on-premises) and public cloud environments. It should monitor workload resource usage and trigger cloud bursting actions based on defined attributes such as CPU usage, memory demand, user role, or data sensitivity. The policy engine must evaluate these attributes dynamically and determine whether to keep workloads local or burst them to the cloud securely. The system specification also includes requirements for secure communication between clusters using encryption, strict role-based access control (RBAC), and logging mechanisms for tracking workload movement. It must ensure high availability, scalability to handle dynamic load, and compliance with organizational or regulatory policies. Additionally, it specifies the use of container orchestration tools like Kubernetes, APIs for policy management, a dashboard for administrators, and integration with identity providers for authentication. These specifications help determine how well a system can perform certain tasks, run applications, or support specific workloads. For example, a high-end gaming system would have different specifications compared to a basic office computer. System specifications are also important when evaluating compatibility with software, ensuring performance efficiency, and planning upgrades. They serve as a guide for users, developers, and IT professionals to understand the system's capabilities and limitations. These requirements ensure that the system can perform its intended functions without issues and provide a good user experience.

2.6.1 Hardware Requirements

- Processor : Intel i3 processor
- RAM : 4GB(min)
- Hard Disk : 500 GB

2.6.2 Software Requirements

- Operating system : Windows 7 Ultimate.
- Coding Language : Java
- Front- End : JSP
- Back-End : MySQL
- Designing : Html, CSS, JavaScript.
- DataBase : MySQL (Wamp Server)

2.7 Software Development Life Cycle Models and Methodologies

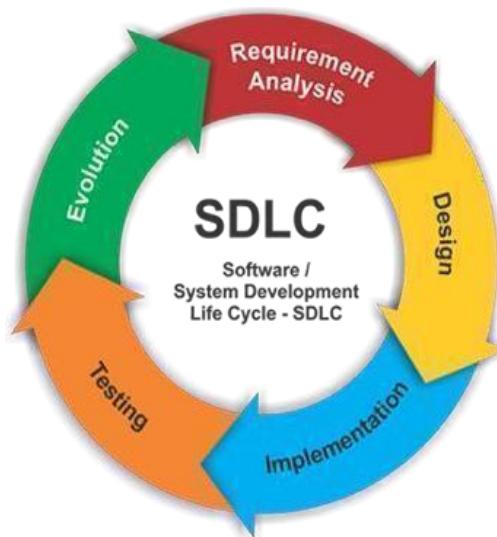


Fig: SDLC

It is essential to select the right SDLC model tailored to the specific concerns and requirements of the project to ensure its success. To explore more about choosing the right SDLC model, you can follow this link for additional information. Furthermore, to delve deeper into software life- cycle testing and SDLC stages, follow the highlighted links here. The exploration will cover various types of SDLC models, their benefits, disadvantages, and when to use them. SDLC models can be viewed as tools to enhance product delivery. Therefore, understanding each model, its advantages, disadvantages, and the appropriate usage is crucial to determine which one suits the project context.

Types of Software developing life cycles (SDLC)

- Waterfall Model
- V-Shaped Model
- Evolutionary Prototyping Model
- Spiral Method (SDM)
- Iterative and Incremental Method
- Agile development

2.7.1 Waterfall Model

The Waterfall Model follows a linear, sequential flow, where progress moves steadily downwards (like a waterfall) through the phases of software development. Each stage in the development cycle begins only after the previous stage is completed. The waterfall approach does not accommodate going back to a previous stage to address changes in requirements. It is the oldest and most well-known method used for software development. The five-stage waterfall model, based on Winston W. Royce's requirements, divides development processes into the following stages:

- Analysis
- Design
- Implementation

Waterfall Model is still valued for its simplicity, clarity in documentation, and ease of management in certain types of projects where a predictable and organized development process is required.

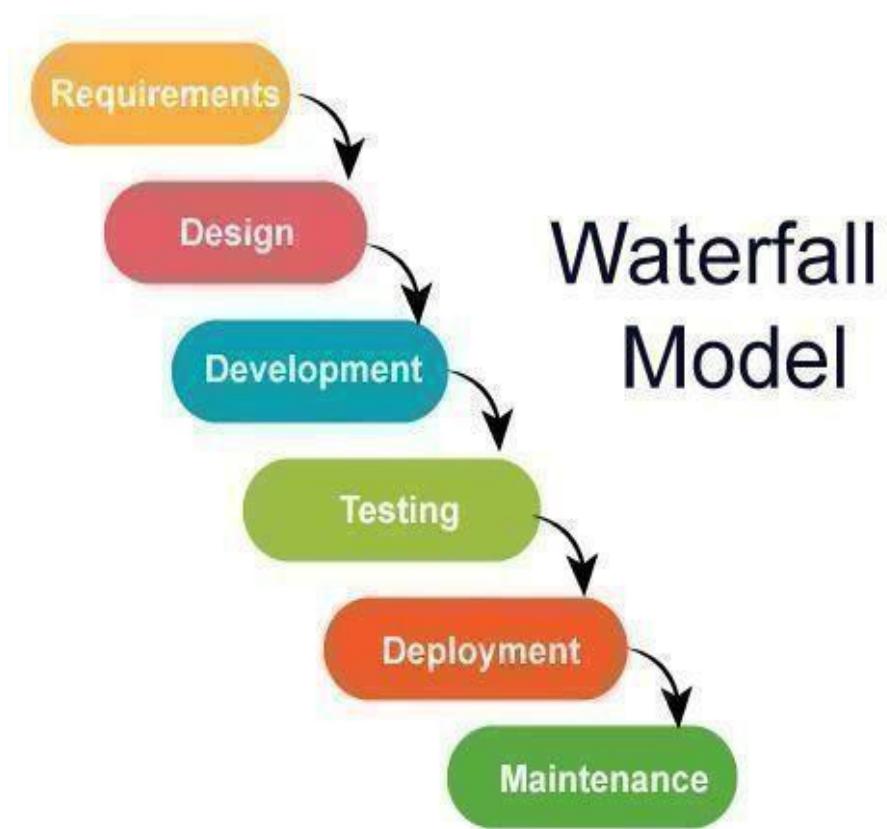


Fig.Water fall model

Advantages

1. Simple to clarify for the clients.
2. Structures approach.
3. Stages and exercises are distinct.

2.8 System Study

The feasibility study evaluates the practicality of the project and enhances fundamental understanding through a comprehensive approach. During system assessment, the integrity evaluation of the proposed structure is crucial to ensure it does not burden the organization. By analyzing these factors, the system study lays the groundwork for defining clear policies, secure communication methods, and automation strategies. This ensures the final architecture not only meets performance demands but also maintains strict security controls, making cloud bursting efficient, flexible, and safe for the organization. The system study explores how attributes such as user roles, workload sensitivity, and system health can be used to automate bursting decisions. It investigates integration points between private and public clusters, including network connectivity, API compatibility, and monitoring tools. By gathering these insights, the system study helps define clear objectives, functional requirements, and constraints. It ensures that the final solution can dynamically and securely burst workloads to the cloud when needed, while maintaining control and visibility across both environments. This comprehensive analysis lays the foundation for designing a scalable, reliable, and secure Kubernetes cloud bursting system. Security aspects form a critical part of the study, focusing on existing authentication methods, access controls, and data protection policies. The study also reviews compliance requirements to ensure any bursting to public clouds aligns with organizational and legal standards.

Three key assessments conducted during feasibility appraisal are:

- Economic Feasibility
- Technical Feasibility
- Social Feasibility

They can be technical, such as limited processing power or storage capacity. They might be economic, like budget limits or cost-effectiveness requirements. Legal or regulatory constraints may apply in areas such as data privacy or security compliance. Time constraints can also impact project deadlines and delivery schedules. This phase helps stakeholders understand how things currently work and what improvements are necessary. It also includes collecting data through observation, interviews, and documentation review. The outcome of system study is a clear picture of user requirements, existing limitations, and the scope for development, forming a solid foundation for system design.

2.8.1 Economic Feasibility

This study examines the financial impact of the system on the organization. It assesses the resources available for system development and justifies expenses. The design should be economically viable, leveraging freely available enhancements wherever possible, with consideration given to necessary purchases.

2.8.2 Technical Feasibility

This study assesses the specific requirements of the system and ensures it does not overly strain existing technical resources. Excessive demands on technical resources can lead to burdens on users. The design should have modest technical requirements, minimizing unnecessary changes.

2.8.3 Social Feasibility

This aspect examines the level of acceptance of the system by users. It involves establishing effective means to familiarize users with the system and ensuring they perceive it as a necessity rather than a threat. User confidence should be bolstered through clear communication and user training. Social feasibility encompasses analyzing how individuals interact within the system or organization and evaluating social impacts to understand their extent and scale.

2.9 Methodology and Algorithms

Pod Security Policies (PSPs) and Attribute Control

Kubernetes Pod Security Policies can enforce security settings for pods, ensuring that pods that burst to another cloud environment adhere to security and compliance rules.

Separation of Workloads

- ABAC policies, can be used to manage access to resources across different clusters.

Algorithms for Secure Kubernetes Cloud Bursting

ABAC Policy Evaluation:

- Policies are evaluated using an algorithm that checks the attributes of a user or workload against a predefined policy.
- The decision to allow or deny access is based on matching the attributes of the subject (e.g., workload) with the policy's criteria.

2.10 Technologies used

Java Technology

Java technology serves as both a programming language and a platform, offering a versatile environment for software development. The Java Programming Language encompasses several key characteristics, including simplicity, architecture neutrality, object orientation, portability, distribution capability, high performance, interpretation, multi threading support, robustness, and dynamism.

The Java Programming Language:

The Java programming language is characterized by a myriad of buzzwords that underscore its versatility and effectiveness:

Architecture Neutral

Java's architecture-neutral design allows it to run on any platform without modification, ensuring compatibility across diverse environments. In the context of attribute-based management of secure Kubernetes cloud bursting, architecture neutral information would describe the core concepts such as monitoring workload conditions, evaluating policies based on attributes, securely transferring workloads between environments, and maintaining audit logs without specifying particular cloud providers, Kubernetes distributions, or encryption technologies. This approach allows architects and developers to understand the fundamental requirements and flow of the system so they can implement it flexibly using the tools and platforms best suited to their organization's needs. It ensures that the architecture remains valid and adaptable regardless of changes in underlying technologies.

Object - oriented

Java follows the object-oriented programming paradigm, emphasizing the creation and manipulation of objects to achieve functionality. These objects interact through well-defined interfaces, sending messages or invoking methods to perform actions like monitoring resources, making decisions based on attributes, and initiating workload transfers. Organizing the system in this way helps in modularization , making it easier to maintain, extend, and understand how different parts collaborate to achieve secure and efficient cloud bursting. There could also be an Attribute Evaluator object that encapsulates the logic to analyze workload and user attributes against policies. This object might interact closely with Policy objects, which define rules as reusable entities that can be updated or extended without affecting the rest of the system. Error handling and security concerns can be managed by objects like Security Manager, responsible for encrypting communications and verifying authentication tokens. Meanwhile, a Logger object records events and actions, providing a centralized way to track system behavior.

High performance

Despite its platform independence, Java offers high performance through efficient byte code execution and optimization techniques.

Robust

Java prioritizes reliability and robustness by incorporating features such as strong memory management, exception handling, and type safety.

Dynamic

Java's dynamic nature allows for runtime adaptation and modification of program behavior, enhancing flexibility and responsiveness.

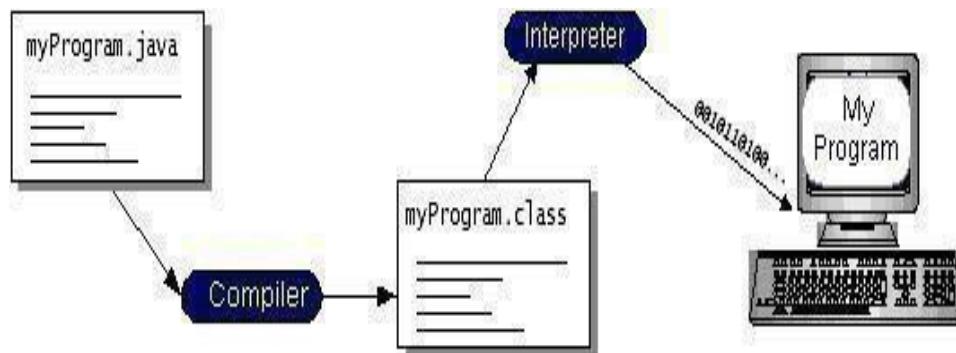


Fig. Working of Java Program

Java byte code serves as the machine code instructions for the Java Virtual Machine (Java VM), enabling the "write once, run anywhere" paradigm. Every Java interpreter, whether it's a development tool or a web browser capable of running applets, functions as an implementation of the Java VM. By compiling a program into byte code, it becomes platform-independent, allowing it to run on any system with a Java VM installed. This flexibility means that a Java program written on one platform, such as Windows 2000, can seamlessly execute on other platforms like a Solaris workstation.

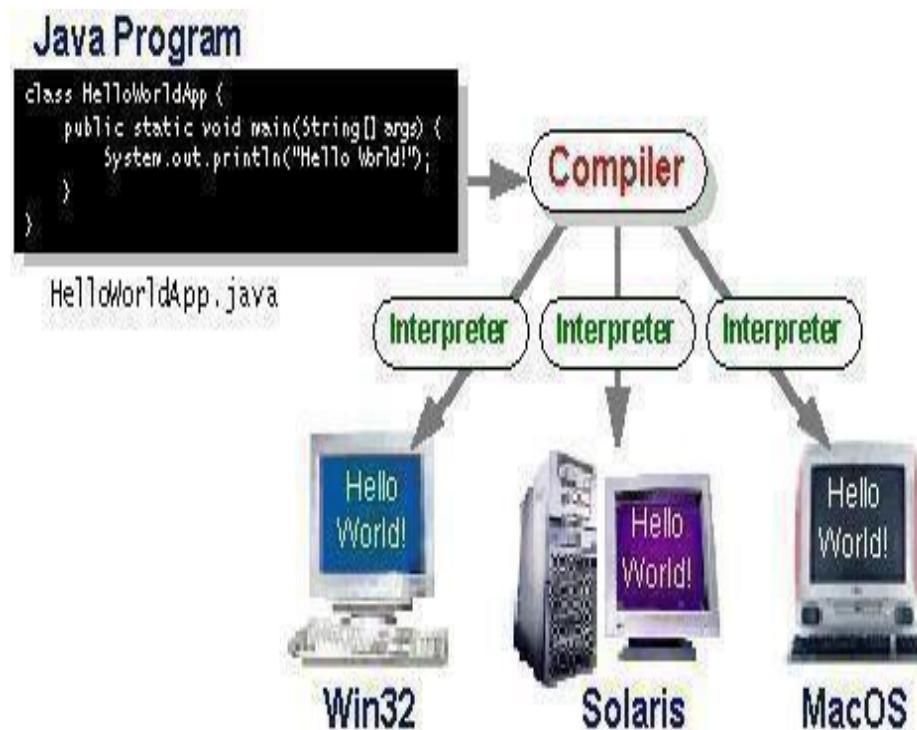


Fig. Implementation of Java Virtual Machine

The Java Platform

A platform refers to the hardware or software environment where a program operates. Common platforms include Windows 2000, Linux, Solaris, and MacOS, which typically combine both operating system and hardware components. The Java platform, however, distinguishes itself by being a software-only platform that operates atop other hardware-based platforms. Comprising two main components, the Java platform consists of:

- The Java Virtual Machine (Java VM):** This serves as the foundation for the Java platform and is adapted to various hardware-based platforms.
- The Java Application Programming Interface (Java API):** This encompasses a vast collection of pre-built software components offering various functionalities, such as graphical user interface (GUI) widgets. The Java API organizes these components into libraries of related classes and interfaces, known as packages.

In the subsequent section, "What Can Java Technology Do?" highlights the functionalities provided by some of the packages within the Java API. The illustration below demonstrates a program operating within the Java platform, with the Java API and the virtual machine shielding the program from direct interaction with the hardware.

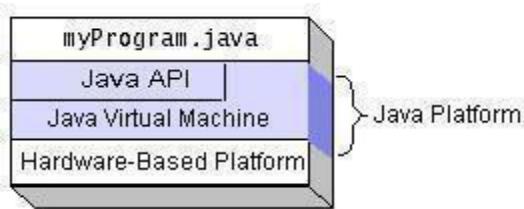


Fig. Program Running on the Java Platform

Native code refers to code that, once compiled, runs directly on a specific hardware platform. In contrast, the Java platform operates as a platform-independent environment, where programs can run across different hardware architectures. While the Java platform may exhibit slightly lower performance compared to native code, optimizations such as smart compilers, well-tuned interpreters, and just-in-time byte code compilers help bridge this performance gap while preserving portability.

What Can Java Technology Do?

Java technology offers a versatile set of programming options. Applets and applications are among the most common types of programs written in the Java programming language. Applets, familiar to web users, adhere to specific conventions allowing them to run within Java-enabled browsers. However, Java's capabilities extend beyond web applets. The Java programming language serves as a robust software platform suitable for various types of applications.

Applications in Java can range from standalone programs to servers that support clients on a network. Servers, such as web servers, proxy servers, mail servers, and print servers, run directly on the Java platform, serving clients' needs. Servlets, another specialized type of program, function

similarly to applets but operate on the server side within Java Web servers, enhancing interactive web applications. The Java API supports a wide array of programs through packages of software components, providing diverse functionalities. A full implementation of the Java platform offers essential features like objects, strings, threads, numbers, input and output mechanisms, data structures, system properties, and date and time functionalities. Furthermore, the Java API includes support for applets, networking (including URLs, TCP, UDP sockets, and IP addresses), internationalization for localized programs, comprehensive security measures (both low-level and high-level), software components known as Java Beans, object serialization for lightweight persistence and communication via Remote Method Invocation (RMI), and Java Database Connectivity (JDBC) for uniform access to relational databases. Moreover, the Java platform provides APIs for 2D and 3D graphics, accessibility, servers, collaboration, telephony, speech, animation, and more. The Java 2 SDK encompasses these functionalities, facilitating the development of diverse applications and enhancing the overall capabilities of the Java platform.

3. SystemDesign

System testing is a critical phase where the entire software system is tested as a whole to ensure that it meets the specified requirements. System testing is a critical phase where the entire software system is tested as a whole to ensure that it meets the specified requirements. It involves verifying both functional and non-functional aspects of the system in an environment that closely resembles the production setup. This type of testing checks the complete end-to-end behavior of the application, including its interactions with external systems, databases, and user interfaces. In system testing, the software is treated as a black box, meaning that the testers do not need to understand the internal code or structure. Instead, they focus on validating whether the system performs its intended functions correctly and efficiently. It includes various forms of testing such as functional testing to verify specific features, performance testing to measure speed and responsiveness, and security testing to ensure data protection and access control. For projects involving attribute-based management of secure Kubernetes cloud bursting, system testing plays a vital role. It helps ensure that workloads can burst from private to public clouds securely and efficiently. It also verifies that attribute-based access controls are enforced correctly, ensuring that only authorized users can initiate or manage cloud bursting operations. Additionally, system testing checks that the bursting process maintains data integrity, meets performance expectations under load, and handles errors or failures gracefully.

3.1 Database Design(E-R Diagram)

An Entity-Relationship (ER) model illustrates the structure of a database using a visual representation known as an Entity-Relationship Diagram (ER Diagram). This model serves as a blueprint for designing the database schema and capturing the relationships between different entities and attributes. The ER model provides a systematic approach to organizing and conceptualizing the data within a database system.

3.1.1 ER model

- The Emergency Room model corresponds to an Entity-Relationship model, serving as a high-level representation of data structures. It is utilized to illustrate the data components and relationships within a defined system.
- It establishes a structured framework for the database. Moreover, it provides a straightforward and easily understandable perspective on the data.
- In Entity-Relationship modeling, the organizational database structure is depicted through a design known as an Entity-Relationship diagram.
- For instance, consider designing a school database. An educational record could be represented as an entity with attributes such as name, ID, age, etc. Similarly, the address could be another entity with attributes like city, street name, zip code, etc., and there would be a relationship between them.

The ERR model is a way of designing databases by visually representing data as entities (such as users or resources), the relationships between them (such as owns, accesses, or manages), and the specific roles each entity plays within those relationships. While the standard ER model shows entities and their associations, the ERR model goes a step further by explicitly showing how an entity behaves or functions in that association. For example, in a cloud bursting system, a user might play the role of "initiator" in a "burst request" relationship with a cloud resource. This approach helps in clearly defining constraints and access control, which is especially useful in attribute-based access control systems. By capturing not just who is connected but also how they are involved, the ERR model offers a deeper and more flexible structure for designing secure and functional systems.

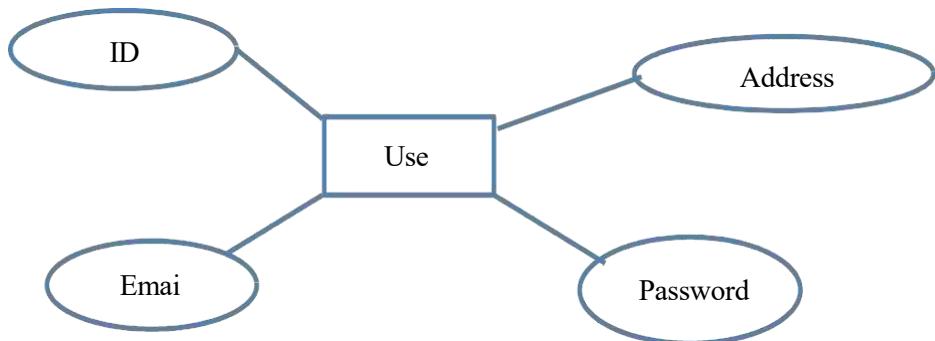
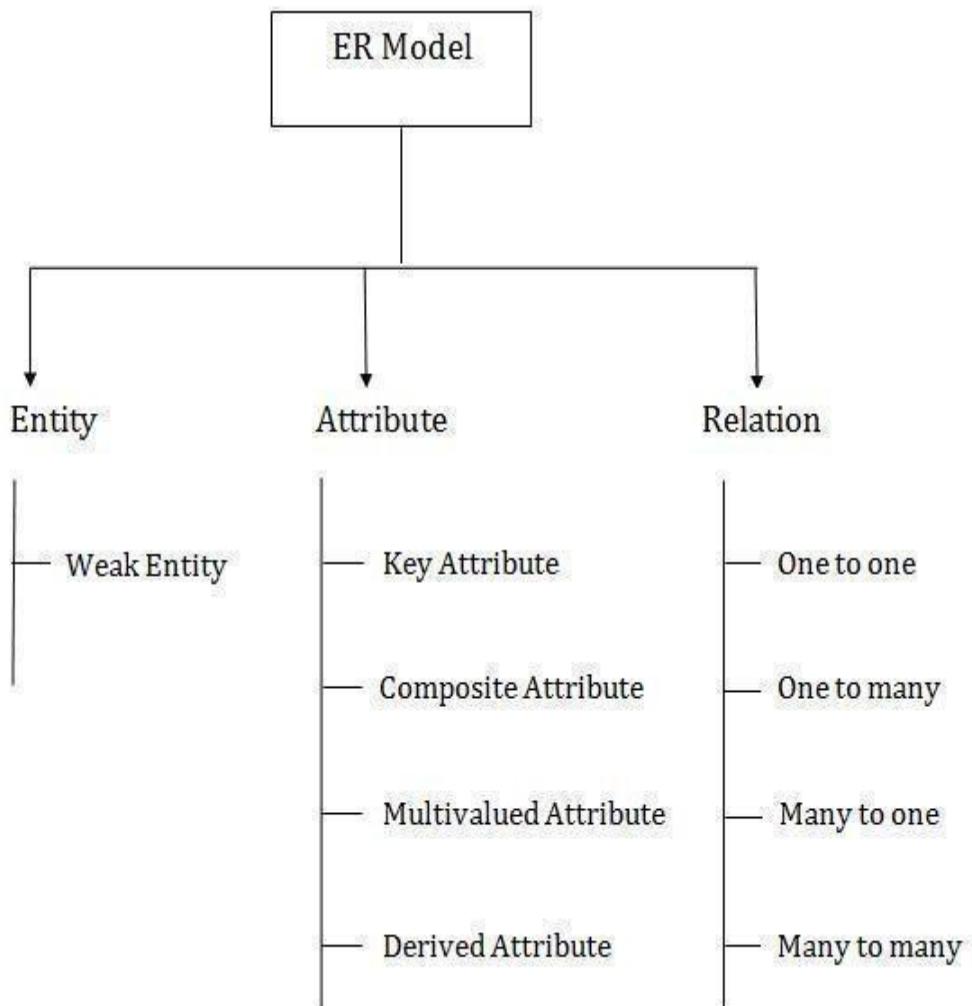


Fig. ER-Model

Component of ER Diagram



Entity

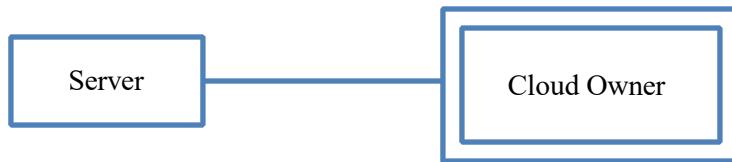
A substance may be anything, class, individual or spot. In the ER frame, a substance can be tended to as square shapes. Think about a relationship as a delineation chief, thing, specialist, office, etc. can be taken as a substance. An entity is something that exists independently and is clearly distinguishable. It can be a person, object, organization, or even a concept, depending on the context. In business, an entity refers to a legally recognized organization such as a company or partnership. In computer science, especially in databases, an entity represents a real-world object or concept that data is stored about, like a student or product. In natural language processing, entities are names of people, places, or organizations that are identified from text. In law, an entity is anything that can enter into legal agreements or be held accountable, like individuals, companies, or institutions. The meaning of "entity" adapts based on the field in which it's being used.



Powerless Entity

A substance that depends upon another component called a frail substance. The frail element contains no critical trait of its own. The feeble substance is addressed by a twofold square shape. A powerless entity refers to something or someone that exists as a distinct being or unit but lacks the ability to influence, control, or change its situation or environment. It may have an identity or role, but no real authority, strength, or decision-making power. In a social or political context, a powerless entity could be a group, organization, or individual that is recognized but has no voice or control like a symbolic committee or a marginalized community.

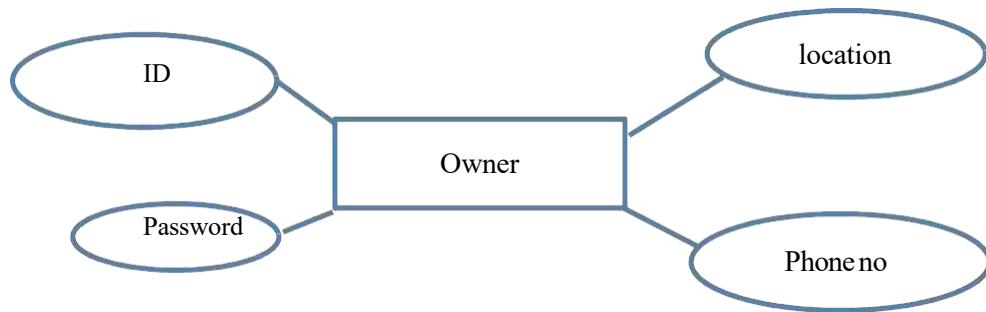
In a business context, it might be a subsidiary company that legally exists but is entirely controlled by a parent company, having no real autonomy. In a fictional or philosophical sense, it could represent a character or spirit that is present and aware but unable to act or intervene in events. Being powerless doesn't mean being non-existent—it means existing without the capacity to affect outcomes.



Characteristic

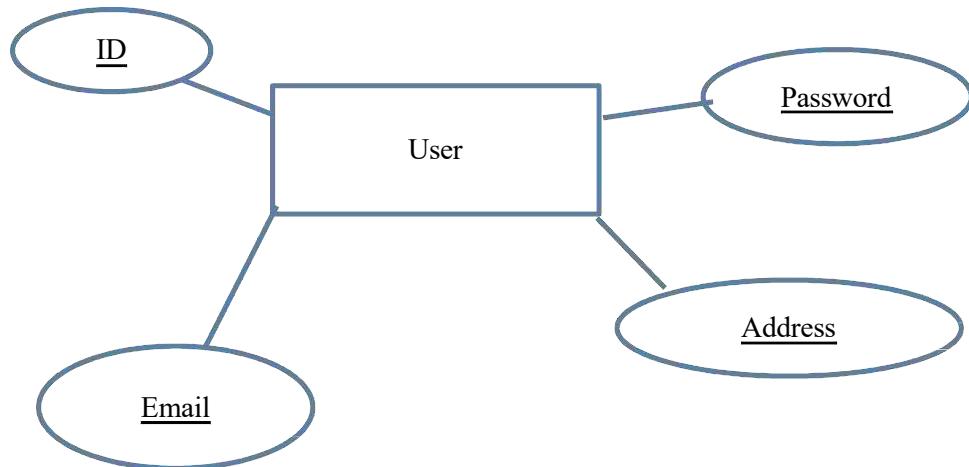
The quality is utilized to depict the property of a section. Obscure is utilized to address a quality.

For example, id, age, contact number, name, etc. can be attributes of a student. It is often recognized or acknowledged as existing but lacks the ability to make decisions or influence outcomes. It may hold a formal role or identity but is unable to act independently or assert control. Its presence is symbolic or passive rather than active or impact . Such an entity may rely on others for action, and its limitations could be due to external control, lack of resources, suppression, or inherent structural weakness. It often evokes a sense of vulnerability, dependence, or insignificance within a larger system.



Key Attribute

The key quality is used to address the essential ascribes of a substance. It tends to a fundamental key. The key property is tended to by a circle with the text underlined. The key attribute of a powerless entity is its lack of control or influence over its own actions or circumstances. While it may exist and be recognized, it cannot effectively change, decide, or impact anything around it due to restrictions, dependence, or absence of authority.

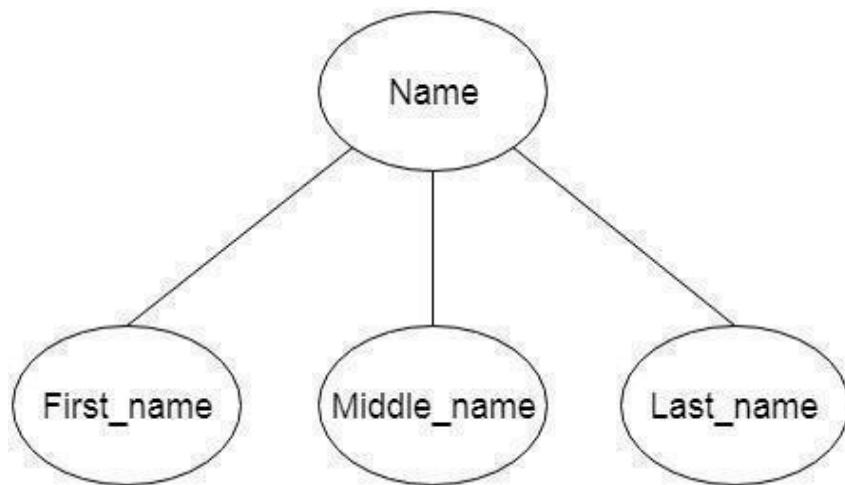


Composite Attribute

A property that is made from various attributes is known as a composite quality. The composite trademark is tended to by an oval, and those circles are related with a circle. A composite attribute is an attribute that can be divided into smaller sub-parts, each of which represents a more basic attribute with its own meaning.

For example, in a database or data modeling context:

- The attribute "Name" can be a composite attribute because it can be broken down into First Name, Middle Name, and Last Name.
- The attribute "Address" is another common example; it can be divided into Street, City, State, and Postal Code.



Multivalued Attribute

A quality can have more than one worth. These qualities are known as a multi valued property. The twofold oval is used to address multi valued property. For example, a student can have more than one phone number.

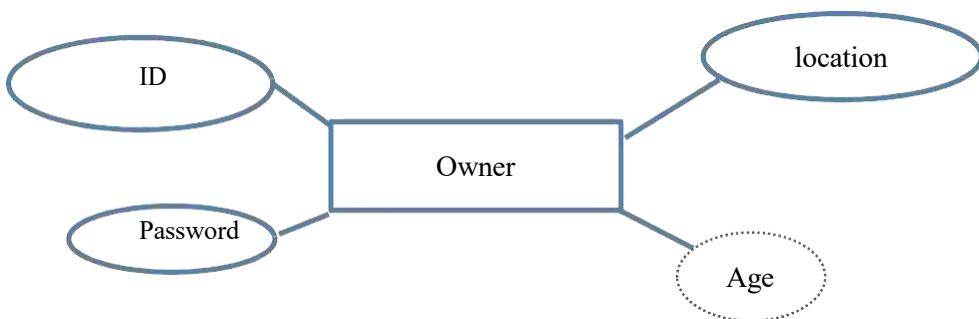


Determined Attribute

A property that can be gotten from another quality is known as a decided attribute. It will in general be tended to by a ran circle. For example, a singular's age changes long term and can be gotten from one more quality like Date of birth. A determined value refers to a value that can be uniquely identified or derived based on another attribute or set of attributes. In database or data modeling terms, if one attribute's value is dependent on another, then its value is said to be determined by that other attribute.

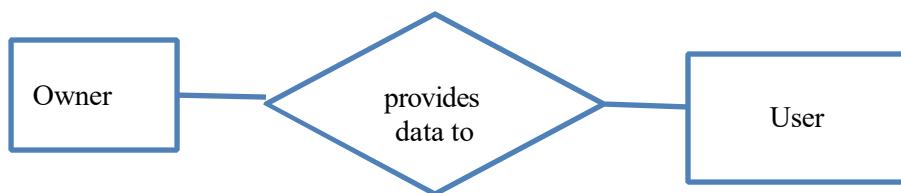
For example, in a student database:

- If the Roll Number determines the Student Name, then the Student Name is a determined value, because it can be identified based on the unique roll number.
- A determined value exists when there is a functional dependency between two attributes. This means the value of one attribute (called the dependent attribute) is directly based on the value of another (called the determinant). If you know the value of the determinant, you can uniquely find the value of the determined attribute. A determined attribute is one whose value is identified or derived from another attribute, typically a key. In databases, it is often used in the context of functional dependency. For instance, if a student ID is known, then the student's name can be determined from it. In this case, the student name is the determined attribute because its value depends on the student ID. This concept is important in database normalization, where attributes that depend on non-key or part of a composite key are separated to maintain data consistency and eliminate redundancy.



Relationship

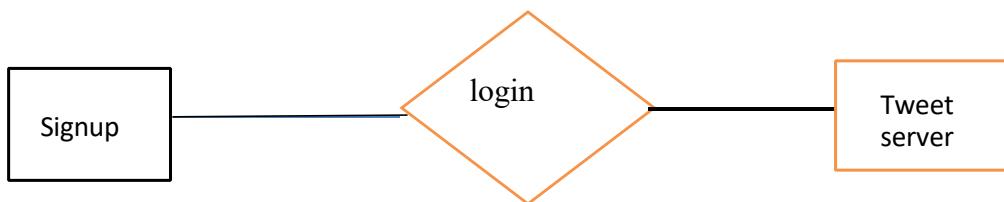
A relationship is used to depict the connection between substances. Important stone or rhombus is utilized to address the relationship. An entity relationship refers to the association between different entities in a database. An entity represents a real-world object or concept, such as a student, teacher, or course. Each entity has attributes that describe its properties. The relationship shows how entities are connected to one another. For example, a student enrolls in a course, or a teacher teaches a subject. In an ER (Entity-Relationship) diagram, a relationship represents the way in which two or more entities are connected within a system. It shows how data in one part of the system relates to data in another, helping to define the structure and logic of the underlying database. These relationships can vary in complexity, ranging from one-to-one, where a single entity is connected to exactly one other entity, to more flexible associations like one-to-many or many-to-many.



Sorts of relationship are as per the following:

One-to-One relationship

At the point when just a single instance of a component is connected with the relationship, then it is known as facilitated relationship. For instance, A female can wed to one male, and a male can wed to one female. A one-to-one relationship in object-oriented design means that one instance of a class is directly linked to exactly one instance of another class. This relationship is important when modeling systems because it clearly defines how objects are connected and ensures data integrity and clarity in interactions.



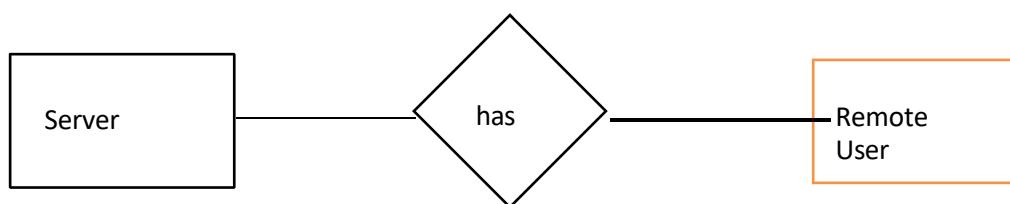
A one-to-one relationship in a system or database design means that one instance of an entity is associated with exactly one instance of another entity, and vice versa. This kind of relationship is useful when you want to separate data for security, modularity, or clarity while maintaining a direct link between two parts of the system.

- Enforce strict associations between specific components.
- Simplify data access when entities are closely tied.
- Improve data organization by splitting sensitive or frequently updated information into a separate, but linked, table or object.

One-to-many relationship

Exactly when simply a solitary illustration of the substance on the left, and more than one event of a component on the right associates with the relationship then this is known as a one-to-various connections.

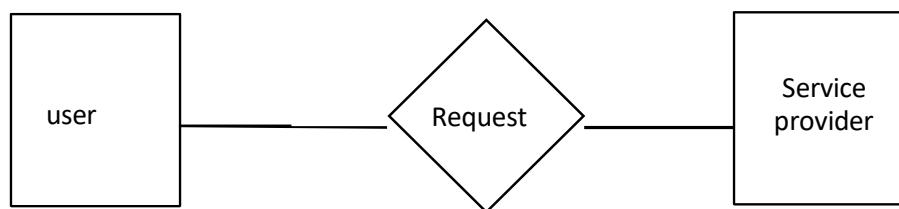
For example, Scientist can envision various manifestations, but the improvement is done by the really express analyst.



Many-to-one relationship

Exactly when more than one event of the component on the left, and simply a solitary event of a substance on the right associates with the relationship then it is known as a many-to-one relationship.

For example, Student enrolls for only a solitary course, but a course can have various students.

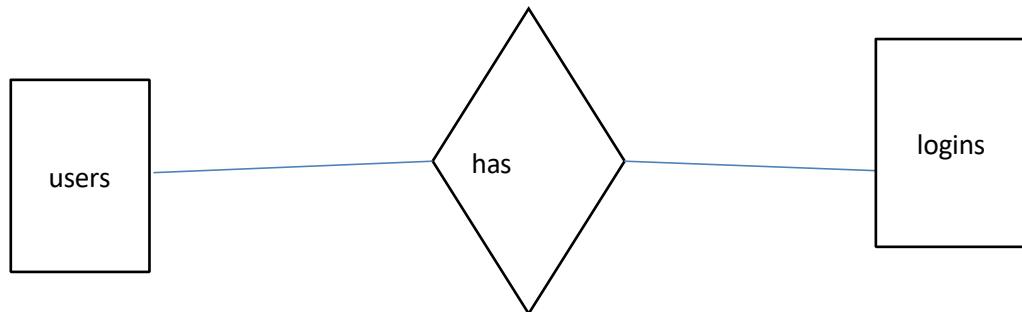


Many-to-many relationship

At the point when more than one event of the substance on the left, and more than one event of a component on the right associates with the relationship then it is known as a many-to-various connections. Employee can allot by numerous exercises and project can have various specialists. A many-to-many relationship in object-oriented design occurs when multiple instances of one class can be associated with multiple instances of another class, and vice versa. This is useful when both entities need to be linked in a flexible and interconnected way.

For example, in the context of attribute-based management of secure Kubernetes cloud bursting:

You might have a Workload class and a Policy class. One workload could be governed by multiple policies (e.g., security, resource limits, compliance), and one policy could apply to multiple workloads. This creates a many-to-many relationship between Workloads and Policies.



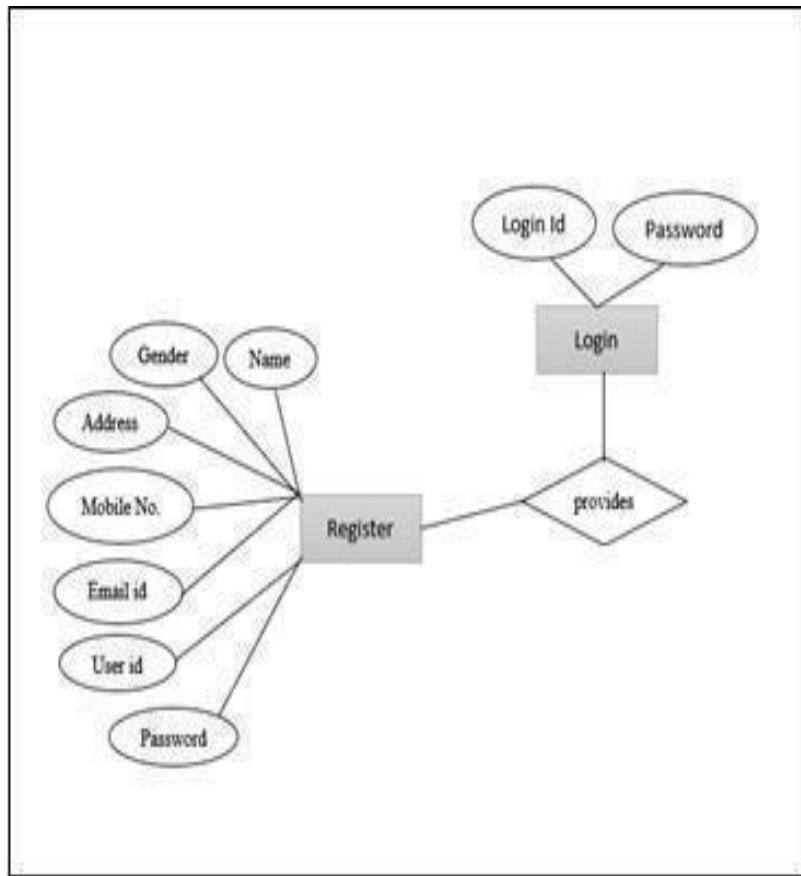


Fig: E-R Diagram

This diagram helps in designing the database schema or back end structure in a way that supports secure and efficient attribute-based cloud bursting. It provides a clear understanding of data dependencies, constraints, and relationships, making it easier for developers and architects to build a robust, scalable, and maintainable system.

3.2 Normalization

Normalization is the primary method for optimizing data in a database to fulfill two essential criteria: Data dependencies are logical, ensuring that all related data items are stored together. Normalization is crucial for various reasons, primarily because it enables databases to occupy minimal disk space, resulting in enhanced performance. Normalization is also referred to as data standardization. The three primary types of normalization are outlined below. Note: "NF" stands for "normal form. "When data is not well-organized, the same information may be repeated in many places (redundancy), and updates can lead to inconsistencies (like changing a phone number in one place but not another). Normalization helps fix this by ensuring that each piece of data is stored only once and related properly.

First typical structure (1NF)

Tables in 1NF should comply with certain standards:

1. Every cell should contain just a solitary (nuclear) esteem.
2. Each part in the table ought to be astoundingly named.
3. All characteristics in a part ought to connect with a comparative region.

Consider a scenario where a Kubernetes administrator defines access policies for users who need to burst workloads from an on-premises cluster to multiple cloud environments. Initially, the access policy table might store attributes in a single field, such as "Region: US, Cloud: AWS" or "Region: EU, Cloud: GCP." This structure makes querying difficult and leads to data redundancy. By restructuring the table so that each attribute is stored in its own row with a clear key-value relationship, the database achieves 1NF. Now, instead of having a single record with combined attributes, each policy is associated with multiple rows, ensuring each field holds atomic values. In Kubernetes cloud bursting, access policies often contain multiple attributes such as user roles, regions, cloud providers, and resource limits.

If these attributes are stored in a single field as a JSON object or a comma-separated list, the database violates 1NF. To comply with 1NF, these attributes must be split into separate rows. 1NF improves data clarity and search ability. It eliminates confusion caused by storing multiple values in a single field and lays the foundation for more advanced normalization steps (2NF, 3NF, etc.). It's the baseline requirement for any well-structured relational database table.a fundamental rule in relational database design that ensures data is stored in a clear and consistent structure. It requires that each table in the database has a well-defined structure where each row represents a single record and each column represents a single attribute of that record. In 1NF, all the values in each column must be atomic, meaning they cannot be divided into smaller parts. This eliminates the use of lists, arrays, or groups within a single field. Additionally, every column should contain values of the same type, and each record in the table must be unique, usually guaranteed by a primary key. The purpose of applying 1NF is to remove redundancy, prevent anomalies in data manipulation, and lay the groundwork for further normalization steps, which help improve the overall efficiency and integrity of the database system. Important aspect of 1NF is the elimination of repeating groups. A table should not have multiple columns representing the same type of data for a single entity. Each record (row) in the table should be unique, typically maintained using a primary key. Applying 1NF helps prevent redundancy and inconsistency, making the database easier to manage, query, and scale. It also sets a strong foundation for higher levels of normalization, such as Second and Third Normal Forms, which further refine the structure and relationships in a database. 1NF, each column in a table should contain only indivisible values, meaning that multi-valued or composite attributes are not allowed. All entries in a column must be of the same type, and each record must be unique. By applying 1NF, the structure of a table becomes more straightforward and easier to manage, reducing data redundancy and improving consistency. This step sets the foundation for more advanced normalization forms that further refine the database design.

Table name : Access Policies

Policy id	User	Role	Attribute key	Attribute value
P001	Alice	Admin	Region	US
P001	Alice	Admin	Cloud	AWS
P002	Bob	Dev	Region	EU
P002	Bob	Dev	Cloud	GCP
P003	Alice	Admin	Region	AP
P003	Alice	Admin	Cloud	AZURE

Table: 1NF**Second typical structure (2NF)**

Tables in 2NF ought to be in 1NF and not have any most of the way dependence (e.g., each non-prime quality ought to be dependent upon the table's fundamental key). Builds on the foundation of First Normal Form by further organizing data to eliminate redundancy caused by partial dependencies. A table is in 2NF if it is already in 1NF and every non-key attribute is fully functionally dependent on the entire primary key. This means that if the primary key is made up of multiple columns (a composite key), no non-key attribute should depend on just part of that key; instead, each non-key attribute must depend on the whole key. The goal of 2NF is to remove partial dependencies, which occur when some attributes depend only on a portion of the composite key rather than the entire key. By doing this, 2NF reduces duplication of data and helps ensure that changes to data only need to be made in one place, which improves data integrity and consistency. Achieving 2NF typically involves splitting tables to separate data that

depends on different parts of the key, making the database structure more efficient and easier to maintain.

Table Name: Policy Attribute

Policy id	User id	Attribute key	Attribute value
P001	U001	Region	US
P001	U001	Cloud	AWS
P002	U002	Region	EU
P002	U002	Cloud	GCP
P003	U001	Region	AP
P003	U001	Cloud	Azure

Table: 2 NF

Second Normal Form (2NF) ensures that all non-key attributes in a table depend fully on the entire primary key, eliminating partial dependencies. This step reduces data redundancy and helps maintain data integrity by organizing the data more logically. Achieving 2NF makes the database more efficient, easier to update, and prepares it for further normalization to address other types of dependencies.

3.3 Data Flow Diagrams

A graphical tool used to describe and analyze the moment of data through a system manual or automated including the process, stores of data, and delays in the system. Data Flow Diagrams are the central tool and the basis from which other components are developed. The transformation of data from input to output, through processes, may be described logically and independently of the physical components associated with the system. It shows how input data is transformed into output data through processes, illustrating the data movement, storage, and processing. DFDs are widely used in system analysis and design to understand and visualize how information moves through a system, often in the early stages of software development DFDs are the model of the proposed system. They clearly should show the requirements on which the new system should be built. Later during design activity this is taken as the basis for drawing the system's structure charts. The Basic Notation used to create a DFD's are as follows:

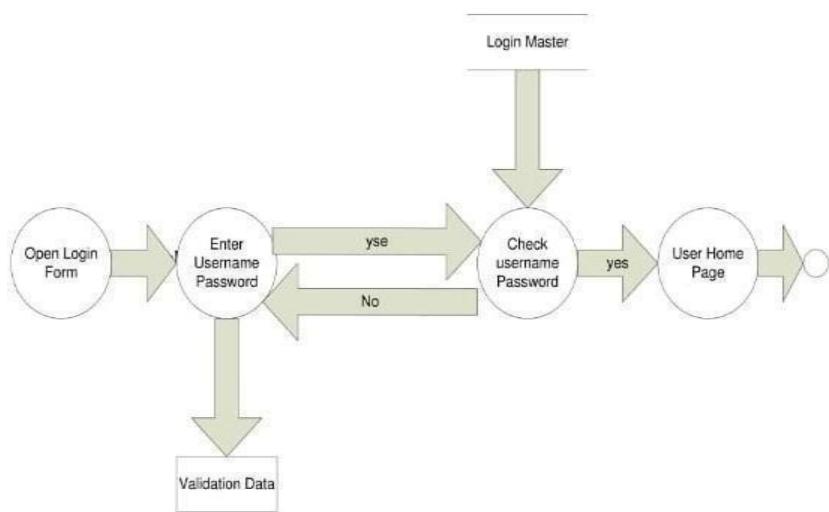


Fig: Data Flow Diagram

3.4 UML Diagrams

The UML diagrams are arranged into fundamental charts, social frameworks, and besides correspondence frame graphs. The diagrams are logically organized in the going with figure:

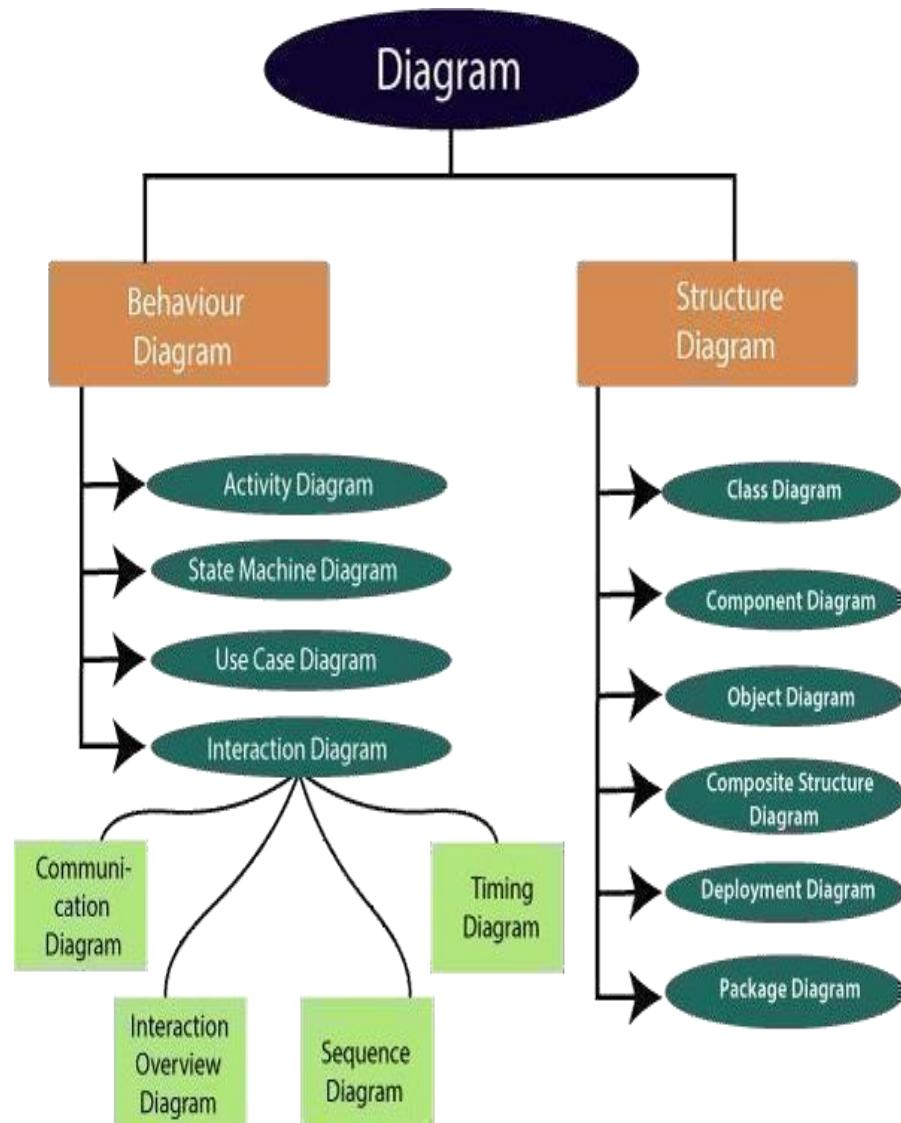


Fig: UML Diagram

Use Case Diagrams

A use case diagram in the Unified Modeling Language (UML) is a type of behavioral diagram described and expressed using a Use-case analysis. Its objective is to present a graphical representation of the value provided by a system in terms of actors, their goals (represented as use cases), and any relationships between those use cases. The primary purpose of a use case diagram is to illustrate which system functions are performed for which actor.

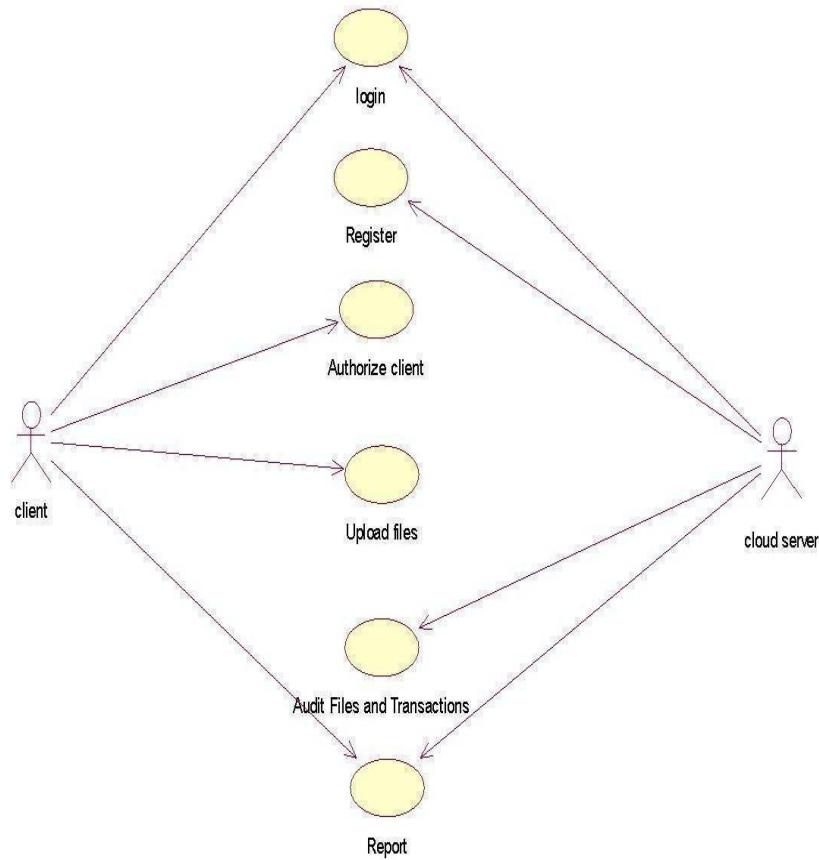


Fig. Use Case Diagram

ClassDiagram

In computer programming, a class diagram in the Unified Modeling Language (UML) is a type of static structural diagram that illustrates the architecture of a system by displaying the structure's classes, their properties, operations (or methods), and the relationships among the classes. It delineates which class holds data. A Class is a blueprint for an object. Objects and classes go hand in hand. We can't talk about one without talking about the other. And the entire point of Object-Oriented Design is not about objects, it's about classes, because we use classes to create objects. So a class describes what an object will be, but it isn't the object itself. In fact, classes describe the type of objects, while objects are usable instances of classes. Each Object was built from the same set of blueprints and therefore contains the same components (properties and methods). The standard meaning is that an object is an instance of a class and object

- Objects have states and behaviors.

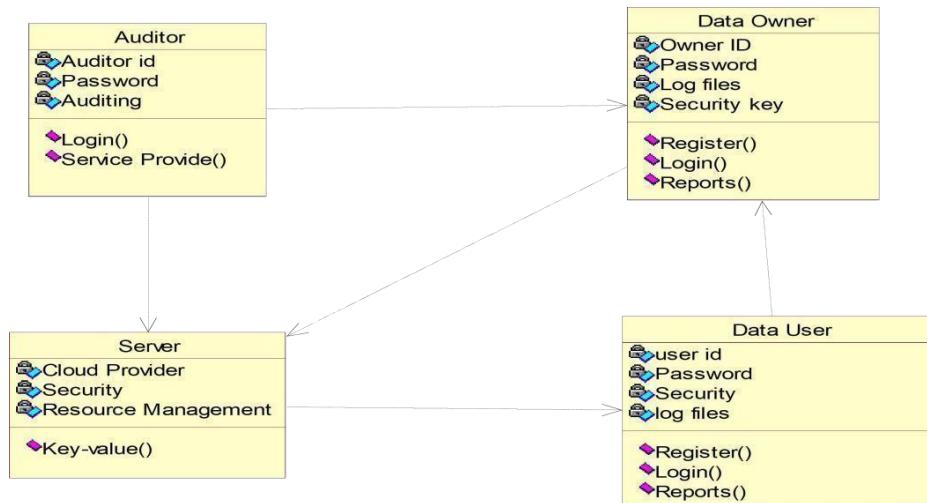
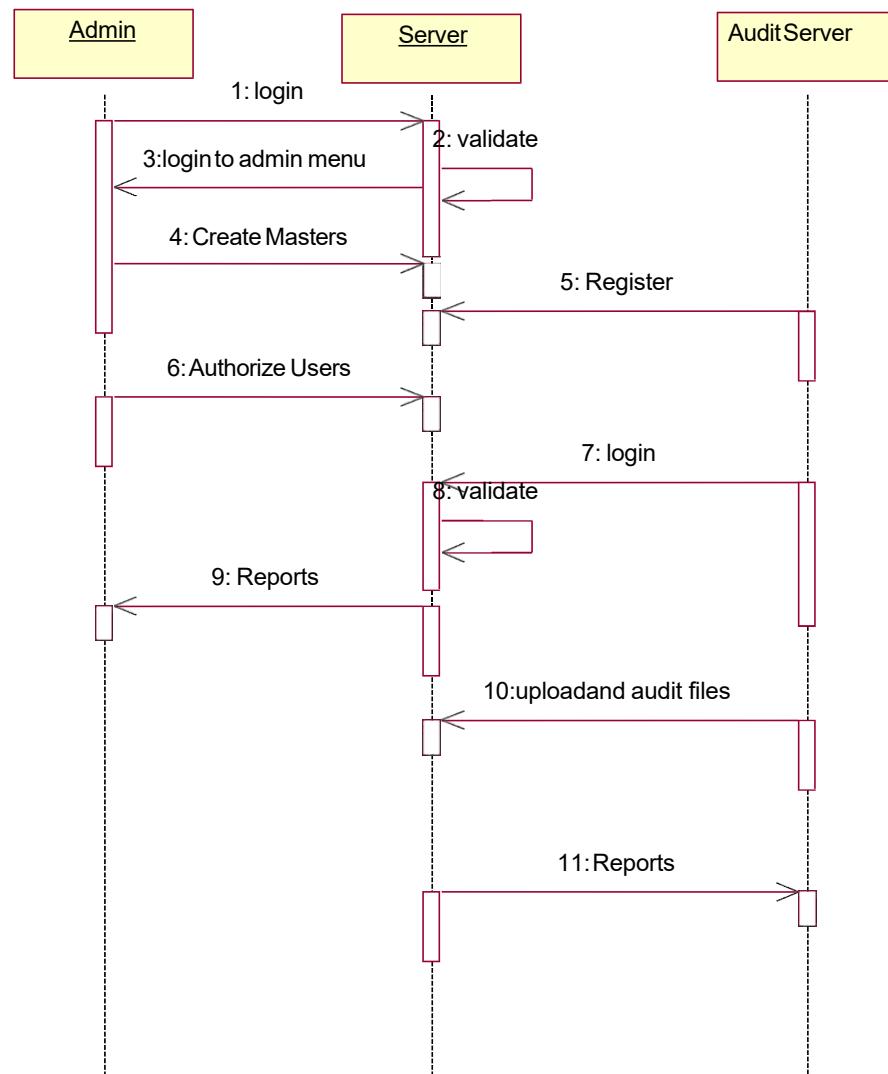


Fig. Class Diagram

SequenceDiagram

A sequence diagram in Unified Modeling Language (UML) is a type of communication diagram that illustrates how processes interact with each other and the order in which they occur. It is a variation of a Message Sequence Chart. Sequence diagrams are sometimes referred to as event diagrams, event scenarios, or timing diagram,. When a user places an order, the front-end sends the request to the Order Service. The Order Service interacts with the Inventory Service to verify if the requested items are in stock. The Inventory Service queries the database for stock levels and responds with the availability status. If the items are available, the Order Service proceeds to initiate a transaction through the Payment Service. Upon receiving payment confirmation, the Order Service instructs the Inventory Service to update the stock and updates the order status in the database. Finally, the Order Service returns a confirmation to the front-end, which presents it to the user. The front end gathers this information and sends it to the authentication service. The service queries the user database to determine if the email is already in use. If not, the service hashes the user's password and creates a new user entry in the database. A response is returned to the front end to confirm the successful registration, which is then displayed to the user. A blueprint for an object. Objects and classes go hand in hand. We can't talk about one without talking about the other. And the entire point of Object-Oriented Design is not about objects, it's about classes, because we use classes to create objects. So a class describes what an object will be, but it isn't the object itself.

**Fig. Sequence Diagram**

Activity Diagram

UML is extremely useful for visualizing and documenting software systems, but the terminology can be somewhat overwhelming for someone unfamiliar with UML. An activity diagram is essentially a flowchart that shows activities performed by a system. If you're new to UML diagramming software used to represent the flow of control or the sequence of activities in a system or process. They visually model the dynamic aspects of a system by showing the flow from one activity to another, similar to a flowchart. An activity diagram gives a clear picture of how a task or system proceeds step by step. It begins with a starting point, flows through a series of actions or activities, and ends when the goal is reached. Each activity shows a specific task being performed. Sometimes the process may face a decision where the flow splits depending on the outcome like "yes" or "no" choices. These decisions guide the direction of the process. There can also be moments where two or more activities happen at the same time, which helps in modeling real-life situations where tasks run in parallel. As the activities complete, they might join again to continue in a single flow. This kind of diagram is useful when explaining workflows such as user registrations, online purchases, or service requests, as it visually maps out every step and alternative paths. It's not just for technical people anyone can understand what's happening by following the flow. An activity diagram continues to be helpful when you want to show how something works over time, especially when there are many steps, decisions, or people involved. It doesn't just show what tasks happen it shows when, why, and how each task connects to the next. For example, in an online shopping system, the diagram might start with a user browsing products, then go to adding items to the cart, followed by payment and finally order confirmation. If payment fails, it might go in a different direction, like showing an error or prompting a retry.

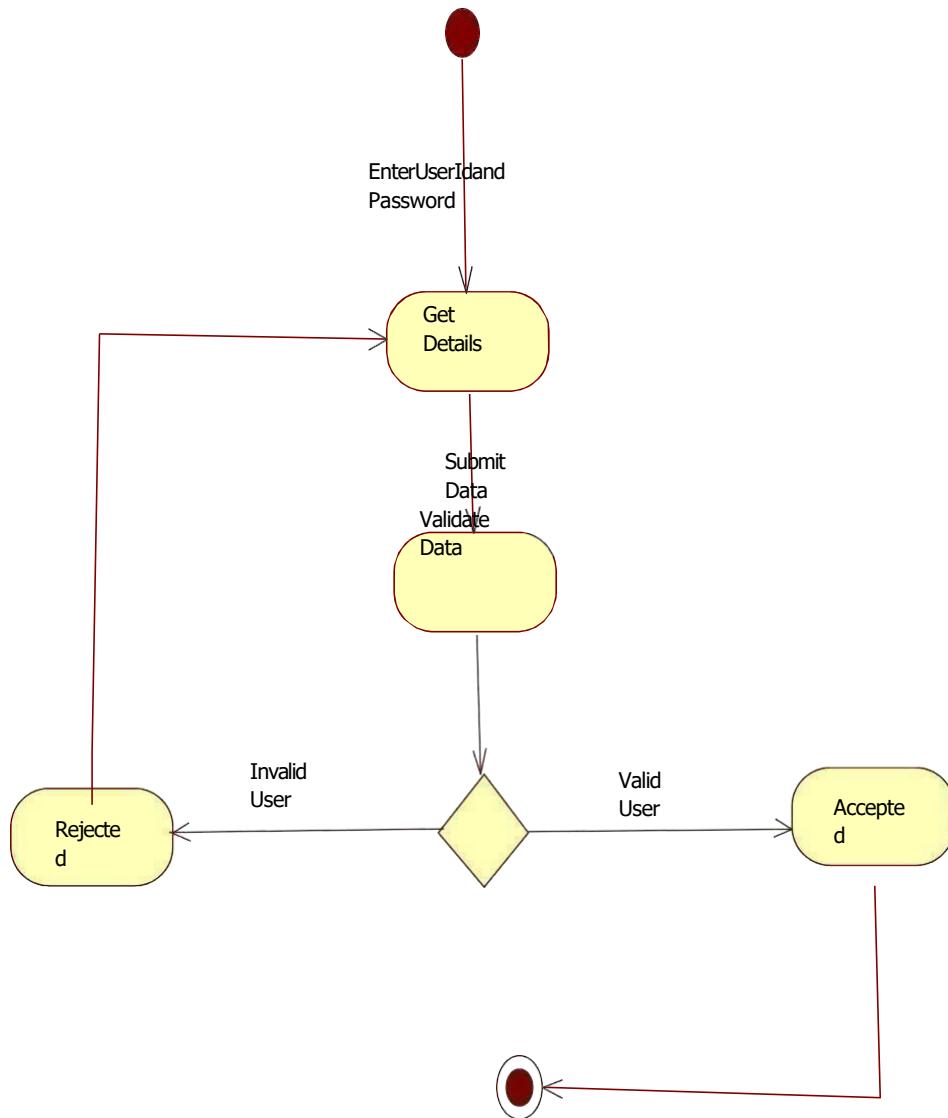


Fig. Activity Diagram

Component Diagram

A component diagram is used to break down a large object-oriented system into smaller components, making them more manageable. It provides a physical view of a system, such as executables, files, libraries, etc., that reside within the node. It visualizes not only the relationships but also the organization between the components present in the system. It assists in forming an executable system. A component is a single unit of the system, which is replaceable and executable. The implementation details of a component are hidden, and it requires an interface to interact and execute a function. It operates like a black box whose behavior is understood by the inputs and required interfaces.

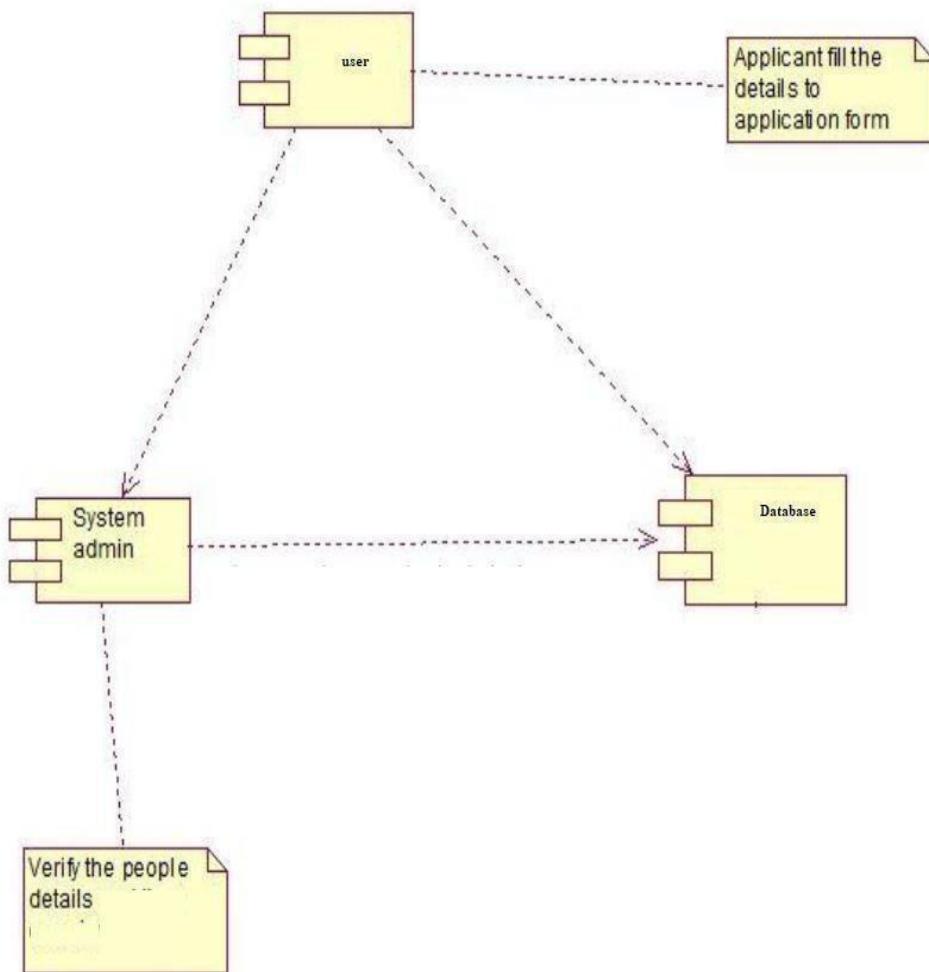


Fig. Component Diagram

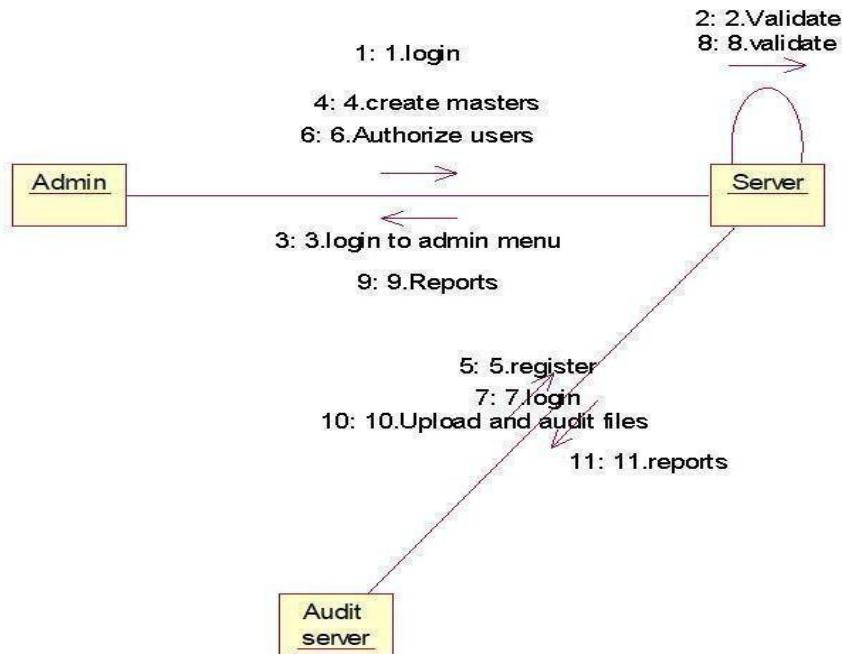
Collaboration Diagram

In a collaboration diagram, the sequence of method calls is depicted using a numbering technique to illustrate how methods are called sequentially. This numbering indicates the order in which the methods are invoked. Let's consider the order management system to describe the collaboration diagram. The method calls in a collaboration diagram resemble those in a sequence diagram. However, the key difference lies in the fact that while the sequence diagram focuses solely on method invocation, the collaboration diagram provides additional information by illustrating the organization of objects involved in the interactions. A collaboration in system design refers to how different objects, roles, or parts of a system work together to achieve a specific goal. It focuses on the interaction between multiple elements rather than what a single element does alone. In UML, collaboration is often shown through diagrams like communication diagrams or collaboration diagrams, where each element (like a person, system, or object) interacts by sending and receiving messages.

For example, in an online shopping system, collaboration might involve:

- A customer placing an order
- The system checking inventory

Each of these components works together through defined roles and message exchanges to complete the full task. Collaboration diagrams help visualize these interactions clearly, especially when multiple participants are involved in completing a single process. It's useful for understanding responsibilities and the flow of communication in a system.

**Fig: Collaboration Diagram**

This visual representation helps us understand the dynamic behavior of the system how attribute-based policies are checked, how decisions are made to migrate workloads, and how security is enforced during communication. By mapping these interactions, the collaboration diagram reveals the dependencies and sequence of operations, making it easier to spot potential bottlenecks or areas for optimization. Overall, it highlights the teamwork between objects that enables flexible, secure, and policy-driven cloud bursting.

Deployment Diagram

The deployment diagram illustrates the physical hardware on which the software will be deployed. It represents the static deployment view of a system, depicting the nodes and their relationships. It details how software is distributed across the hardware. The deployment diagram maps the software architecture created in design to the physical system architecture, where the software will be executed as a node. As it involves multiple nodes, the relationships are depicted using communication pathways.

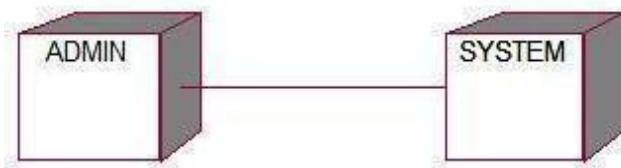


Fig. Deployment Diagram

4. Testing

Testing plays a crucial role in ensuring the reliability and performance of any system or product. It involves examining various aspects of the software or hardware to identify potential issues, inconsistencies, or failures before the final release. Rather than simply checking if things work, testing is about understanding how well they work under different conditions. It often starts early in the development process and continues through various phases, adapting to changes and improvements along the way. A thoughtful testing approach includes understanding the requirements, anticipating possible user behaviors, and simulating real-world scenarios to reveal hidden problems. Whether it's manual testing done by a human or automated testing performed through scripts and tools, the ultimate goal is to deliver a smooth and dependable experience to the end user. Testing doesn't just assure quality it builds confidence in what's being delivered. Testing is not just a phase it's a mindset that ensures the end product meets expectations. When done thoughtfully, it helps catch errors early, reducing the risk of failures after deployment. Testing contributes to a smoother user experience and helps maintain trust in the product. Teams that prioritize testing are better prepared to handle changes and improve continuously based on feedback and results. Testers bring a valuable perspective to the development process. They don't just look for what's broken they look for what could go wrong and how to make it better. Their feedback often leads to important changes that improve both the product and the process. In many ways, testers act as advocates for the end user, helping to ensure the product is not only functional but also reliable and enjoyable to use. Testing can be challenging, especially when working with complex systems or tight deadlines. Sometimes issues are hard to reproduce, or bugs appear only under specific conditions. Testers must be detail-oriented, curious, and patient. They often have to communicate clearly with developers and stakeholders, making sure that everyone understands the nature and impact of any issues found.

4.1. Introduction

Software Testing is a critical element of software quality assurance and represents the ultimate review of specification, design and coding, Testing presents an interesting anomaly for the software engineer.

- Testing is a process of executing a program with the intent of finding an error.
- A good test case is one that has a probability of finding an as yet undiscovered error.
- A successful test is one that uncovers an undiscovered error.
- These above objectives imply a dramatic change in view port.

The testing phase for the Attribute-Based Management of Secure Kubernetes Cloud Bursting system is crucial to ensure its reliability, security, and performance. This system dynamically extends Kubernetes clusters to external cloud environments based on user attributes and predefined policies. The testing aims to validate that all components attribute verification, policy enforcement, resource scaling, and data protection function as intended under various scenarios. The testing strategy includes functional testing to verify the correct implementation of attribute-based access control (ABAC) mechanisms, integration testing to ensure seamless cloud bursting across hybrid environments, performance testing under different workloads, and security testing to confirm that sensitive operations and data transmissions are protected. Emphasis is placed on simulating real-world use cases, including dynamic user roles and varying workload demands, to ensure the system behaves correctly and securely in production environments. Testing cannot show the absence of defects, it can only show that software errors are present.

4.2 Types of Testing

Testing is a critical phase in the software development process that ensures the system or application functions as expected, meets requirements, and is free from significant defects. It involves executing software to identify bugs, verify performance, and validate that all components interact correctly. There are various types of testing, each focusing on different aspects of the system. These can broadly be divided into functional testing (which checks what the system does) and non-functional testing (which checks how the system performs). Depending on the stage of development, testing can be done at different levels such as unit testing, integration testing, system testing, and acceptance testing. Different testing types help ensure quality from different angles, such as security, usability, performance, and compatibility. Choosing the right types of testing for a project improves software reliability, user satisfaction, and long-term maintainability.

The following are the Testing Methodologies:

- Unit Testing.
- Integration Testing.
- User Acceptance Testing.
- Output Testing.
- Validation Testing.

Unit Testing

Unit testing puts together check exertion with respect to the smallest unit of Software plan that is the module. Unit testing rehearses unequivocal courses in a module's control development to ensure all out consideration and most outrageous screw up distinguishing proof. This test revolves around each module independently, ensuring that it limits suitably as a unit. Subsequently, the naming is Unit Testing. During this testing, each module is attempted solely and the module points of communication are affirmed for the consistency with plan assurance.

Coordination Testing

Coordination testing settle the issues related with the twofold issues of affirmation and program improvement. After the item has been incorporated a bunch of high request tests are directed. The primary goal of testing process is to take unit attempted modules and collects a program structure that has been coordinated by plan. Coordination testing is a type of software testing that focuses on verifying how well different components, modules, or systems work together. It ensures that when multiple parts interact whether within the same system or across different systems they do so correctly and smoothly. Instead of testing individual units separately, coordination testing looks at the communication, data exchange, and timing between components. It checks whether actions that depend on other parts are executed in the right order and with the expected results.

Coming up next are the kinds of Integration Testing:

Top-Down Integration

This method is a languid procedure for managing the progress of program structure. Modules are composed by moving lower through the control request, beginning with the chief program module. The module subordinates to the essential program module are coordinated into the development in either a significance first or breadth first way. The product is tried from primary module and individual stubs are supplanted when the test continues downwards.

Base up Integration

This system begins the new turn of events and testing with the modules basically level in the program structure. Since the modules are facilitated from the base up. The base up fuse system may be executed with the going with progresses, Bottom-Up Integration (also called Base-Up Integration) is an approach in software testing and development.

- The bunch is tried.
- Drivers are disposed of and bunches are joined moving vertical in the program structure.
- The granular perspectives test every module independently and afterward every module will be module is incorporated with a primary module and tried for usefulness.

This means that the foundational parts of the system, such as utility functions or core services, are developed and verified before integrating the modules that depend on them. In the context of a complex system like attribute-based management of secure Kubernetes cloud bursting, bottom-up integration helps ensure that the fundamental building blocks such as workload monitoring, attribute evaluation, and secure communication modules are thoroughly tested and stable before combining them with higher-level components like policy controllers or user interfaces. It ensures that when multiple parts interact whether within the same system or across different systems they do so correctly and smoothly. Instead of testing individual units separately, coordination testing looks at the communication, data exchange, and timing between components.

Validation Testing

At the end of integration testing software is completely assembled as a package. Validation testing is the next stage, which can be defined as successful when the software functions in the manner reasonably expected by the customer. Reasonable expectations are those defined in the software requirements specifications. Information contained in those sections form a basis for validation testing approach. Reasonable expectation is defined in the software requirement specification document that describes all user-visible attributes of the software. The specification contains a section titled “Validation Criteria”. Information contained in that section forms the basis for a validation testing approach.

Validation Test Criteria

Software validation is achieved through a series of black-box tests that demonstrate conformity with requirement. A test plan outlines the classes of tests to be conducted, and a test procedure defines specific test cases that will be used in an attempt to uncover errors in conformity with requirements. met. After each validation test case has been conducted, one of two possible conditions exists: The function or performance characteristics conform to specification and are accepted, or a deviation from specification is uncovered and a deficiency list is created. Deviation or error discovered at this stage in a project can rarely be corrected prior to scheduled completion. It is often necessary to negotiate with the customer to establish a method for resolving.

Client Acceptance Testing

Client Acceptance of a structure is the indispensable part for the advancement of any system. The structure practical is pursued for client affirmation by constantly keeping in touch with the arranged system clients at the hour of making and making changes any spot required. The system made gives a well-disposed UI that can undoubtedly be another seen even by an individual to the framework.

Output Testing

In the wake of playing out the endorsement testing, the ensuing stage is yield attempting of the proposed system, since no structure could be useful if it doesn't convey the essential outcome in the predefined plan. Getting some data about the setup expected by them tests the outcomes made or displayed by the structure feasible. Consequently, the outcome configuration is considered in 2 ways one is on screen and one more in printed plan.

4.3 Testing Methodology

A system for structure testing organizes structure examinations and plan methodologies into an overall organized series of steps that results in the productive advancement of programming. The testing procedure must collaborate test organizing, try setup, test execution, and the resultant data grouping and appraisal. A strategy for programming testing ought to oblige low-level tests that are critical to affirm that a little source code part has been precisely executed as well as irrefutable level tests that endorse huge system limits against client.

Frame work testing

Programming once supported ought to be gotten together with other structure parts (for instance Equipment, people, informational collection). Structure testing affirms that all of the parts are real and that overall system work execution is achieved. It moreover tests to find irregularities between the structure and its novel objective, current subtleties and system documentation.

Unit testing

This philosophy begins the new turn of events and testing with the modules basically level in the program structure. Unit testing is central for affirmation of the code conveyed during the coding stage, and hence the goals to test within reasoning of the modules. Including the distinct arrangement portrayal as an assistant, critical Conrail ways are attempted to reveal bungles inside the constraint of the modules. This testing is finished during the programming stage itself. Sort of testing step, every module was viewed as working agreeably as respects to the normal result from the module. In Due Course, most recent innovation head ways will be thought about. As a feature of specialized develop a large number of the systems administration framework will be nonexclusive in nature so future undertakings can either utilize or cooperate with this. What the future holds a ton to propose to the new development and refinement of this endeavor. In system testing, coordination testing passed parts are taken as data.

4.4 Test Cases

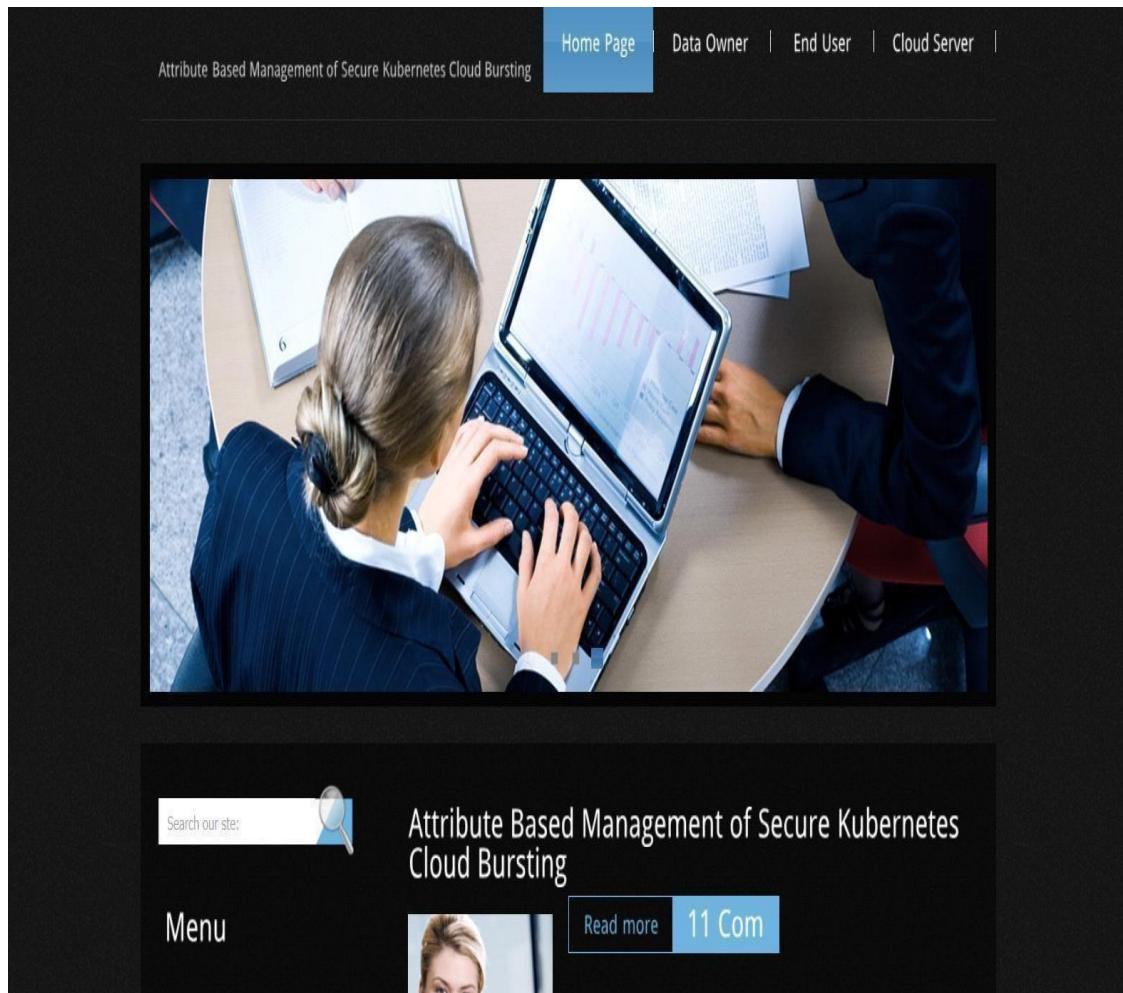
An experiment is a report, which has a lot of test data, preconditions, expected results and post conditions, made for a particular test circumstance to really take a look at consistence against a specific need. Try goes probably as the early phase for the test execution, and ensuing to applying a lot of information regards, the application has a legitimate outcome and leaves the structure at some end point or generally called execution post condition. Each test case includes specific inputs, expected outputs, execution steps, and the actual result. For example, a test case might check whether a workload with high CPU usage and a low-security level is correctly burst to a public cloud cluster when the private cluster reaches capacity. Another might verify that a sensitive workload with a strict security attribute is prevented from bursting, even under high load. Test cases should verify that Kubernetes clusters both private and public handle workload scheduling, migration, and scaling as expected based on attribute- driven policies. For example, a test case might ensure that when a Kubernetes node in the private cluster reaches a CPU usage threshold, the system triggers the bursting process to deploy pods in the public cluster without downtime. Another important test verifies that Kubernetes Role-Based Access Control (RBAC) settings correctly enforce user permissions during workload migration, preventing unauthorized users from initiating bursts. Additionally, test cases should validate that Kubernetes networking and service discovery continue to function seamlessly as workloads move between clusters. Other test cases might include verifying pod health monitoring during bursting, ensuring persistent volumes are handled correctly if state full workloads are involved, and confirming that Kubernetes APIs respond properly to policy updates or scaling commands. Testing should also cover failure scenarios, such as what happens if bursting fails midway does the system roll back safely? and how Kubernetes handles node failures or network partitions during bursting operations. These Kubernetes focused test cases ensure that the cloud bursting architecture integrates smoothly with Kubernetes features, maintaining security, reliability, and performance.

S.no	Test Cases	Input	Expected Result	Actual Result	Status
1	User registration	Enter all fields	User gets registered	Registration is successful	Pass
2	User registration	If user miss any fields	User not registered	Registration is unsuccessful	Fail
3	Cloud login	Give the user name and password	Cloud home page should be opened	Cloud home page has been opened	Pass
4	Upload data set	Test whether the data set is uploaded or not into the system	If dataset is uploaded	We can do further operations	Pass

5. Implementation

5.1 Sample Outputs

Home Page



Screen no 5.1.1: Home Page

Description: Secure & scalable Kubernetes cloud bursting using attribute-based access control.

Data Owner Registration

The screenshot shows a web application window titled "Attribute Based Management" with a sub-page "Data Owner Registration". The URL in the address bar is "http://localhost:8090/kuber/Register.html".

The page has a dark header with a search bar and a logo. The main content area has a dark background.

Left sidebar (Menu):

- Home
- Data Owner
- End User
- Cloud Server

Main Content:

Data Owner Registration

User Name (required)

Password (required)

Email Address (required)

Mobile Number (required)

Cloud Providers

Your Address

Date of Birth (required)

Enter Location (required)

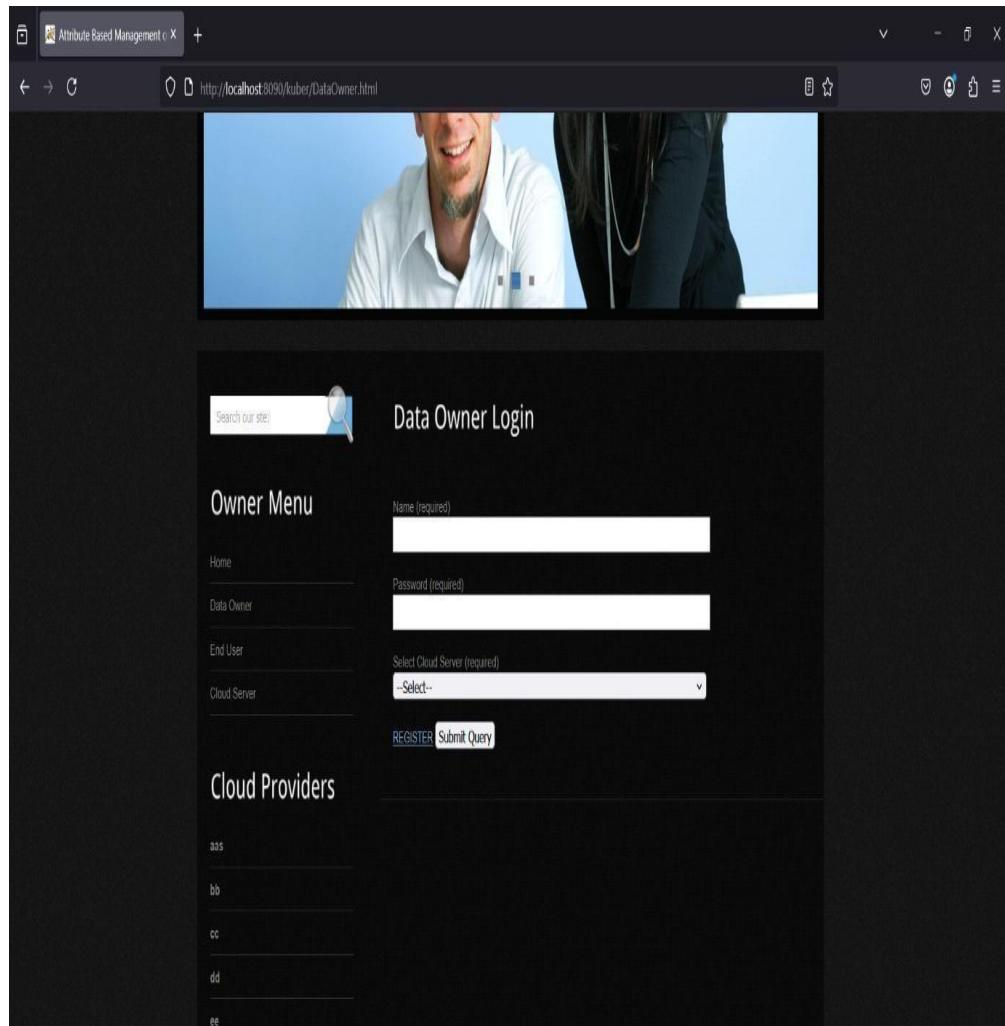
Select Cloud Server (required)

Service Level Agreement Period-SLA (required)

Screen no 5.1.2: Data Owner Registration

Description: Register data owner to manage secure Kubernetes cloud bursting with attribute-based control.

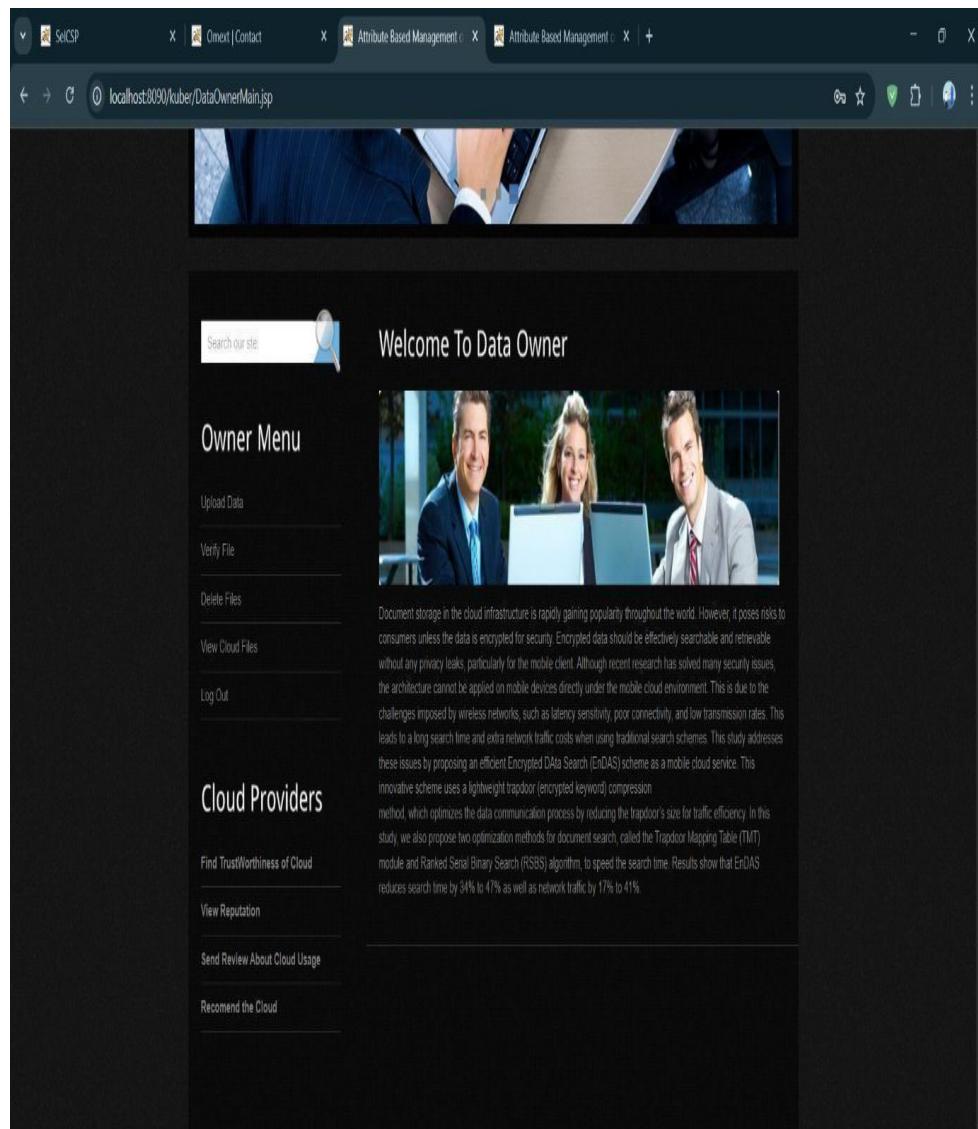
Data Owner Login



Screen no 5.1.3: Data Owner Login

Description: Login for Kubernetes cloud bursting using attribute based access control.

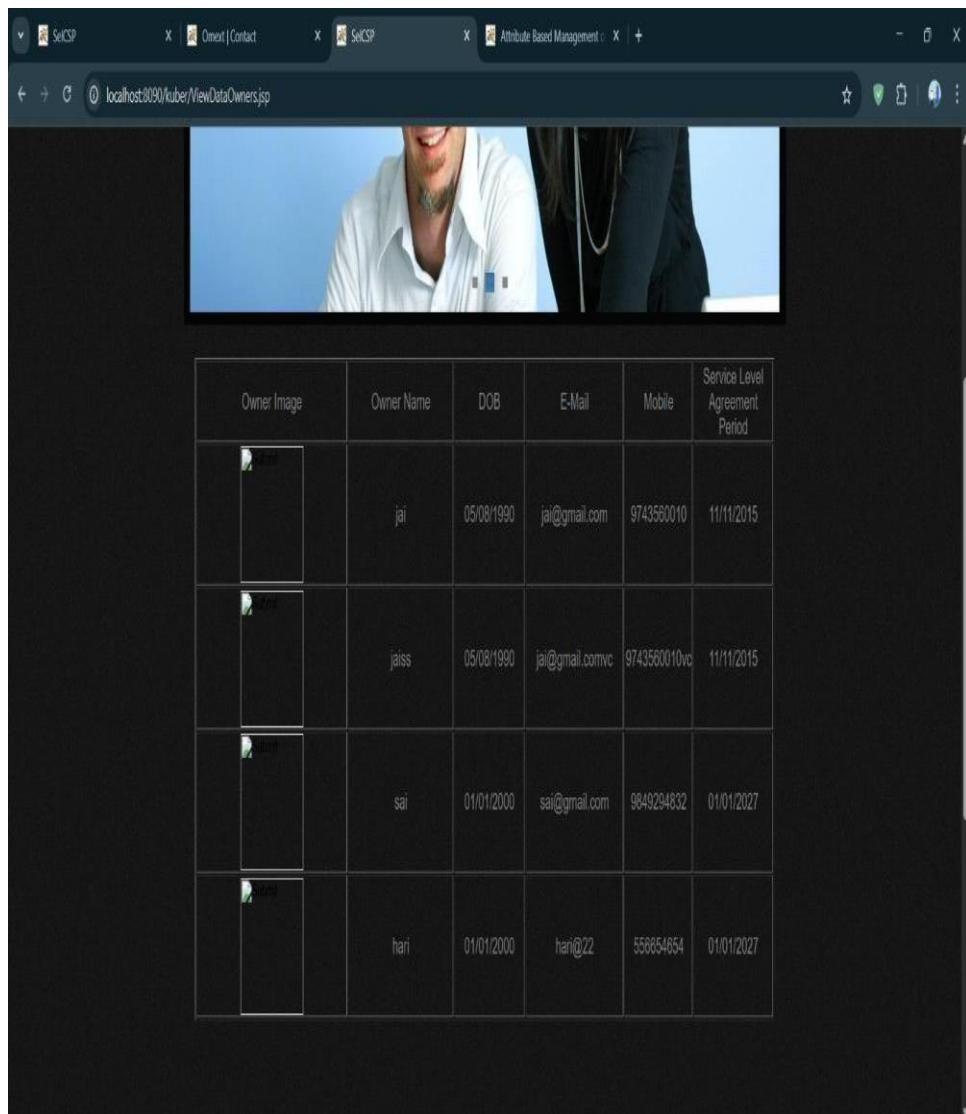
Welcome to Data Owner



Screen no 5.1.4: Welcome to Data Owner

Description: Welcome message to data owner.

Owner Details



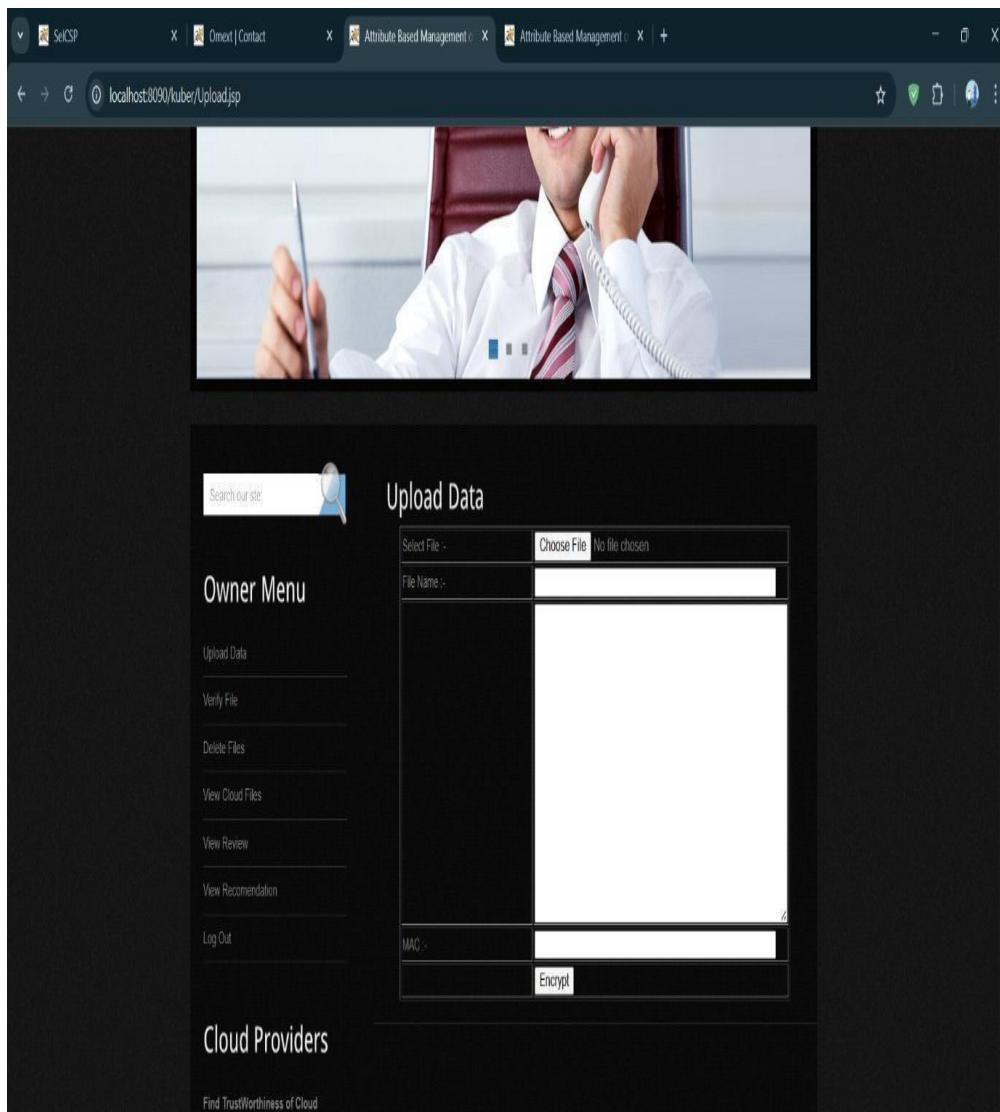
The screenshot shows a web browser window with four tabs open. The active tab is titled "View Data Owners" and has the URL "localhost:8090/kuber/ViewDataOwners.jsp". The page displays a grid of owner information. At the top left is a placeholder image for a profile picture. To its right are columns for "Owner Name", "DOB", "E-Mail", "Mobile", and "Service Level Agreement Period". There are four rows of data, each corresponding to a different owner.

Owner Image	Owner Name	DOB	E-Mail	Mobile	Service Level Agreement Period
[Placeholder Image]	jai	05/08/1990	jai@gmail.com	9743560010	11/11/2015
[Placeholder Image]	jaius	05/08/1990	jai@gmail.comvc	9743560010vc	11/11/2015
[Placeholder Image]	sai	01/01/2000	sai@gmail.com	9849294832	01/01/2027
[Placeholder Image]	hari	01/01/2000	hari@22	556654654	01/01/2027

Screen no 5.1.5: Owner Details

Description: View and manage your details for secure Kubernetes cloud bursting with attribute -based access.

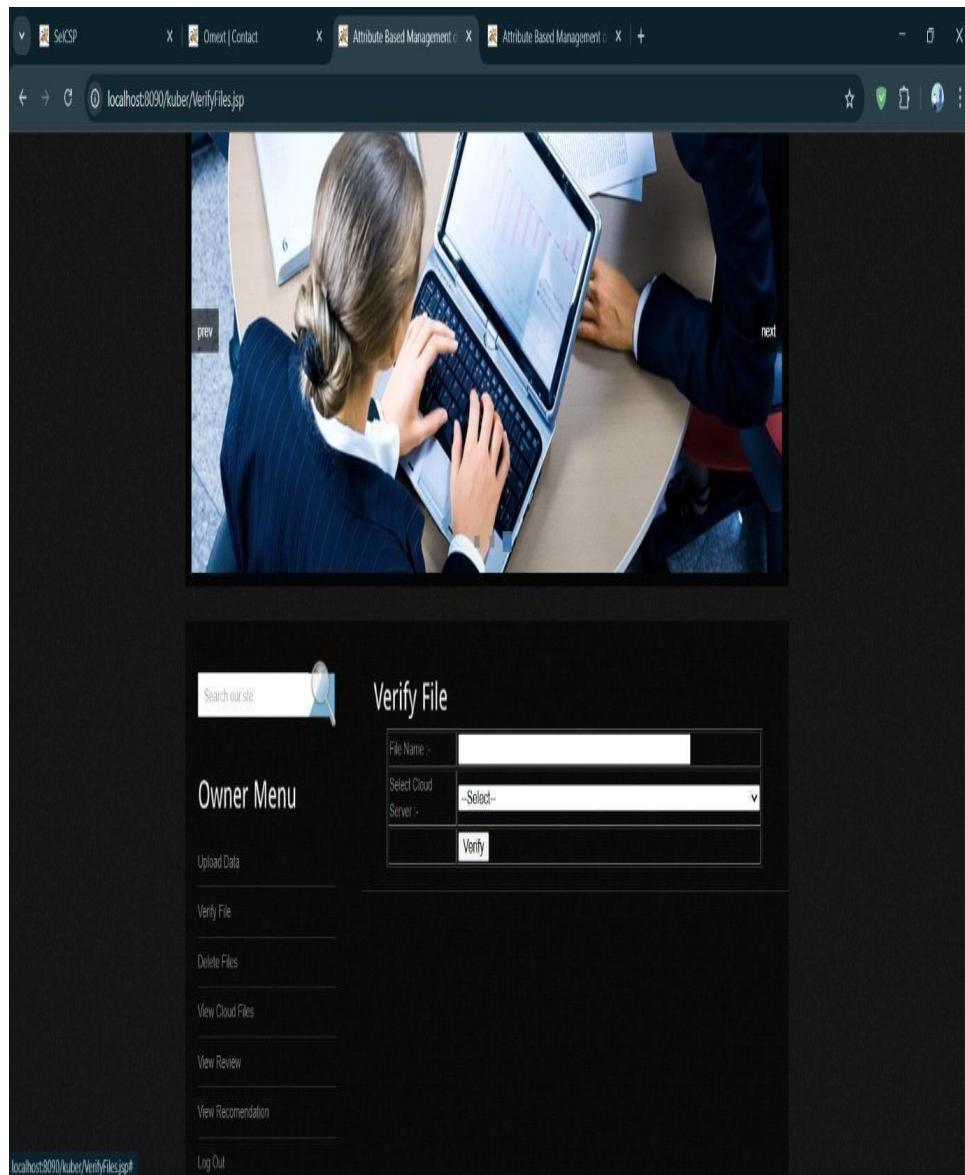
Upload Data



Screen no 5.1.6: Upload Data

Description: Easily uploading data for secure Kubernetes cloud bursting with attribute based access control.

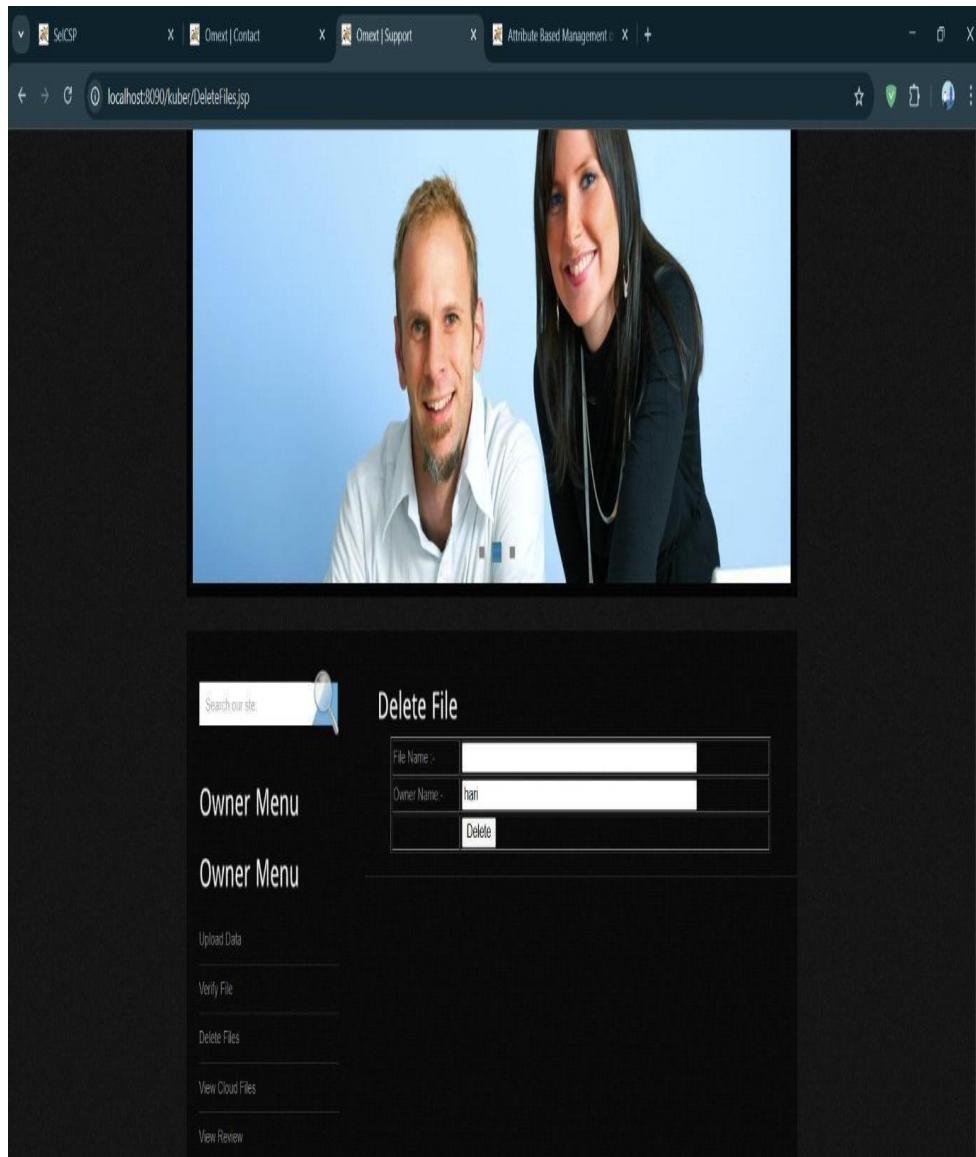
Verify Files



Screen no 5.1.7: Verify Files

Description: Verify uploaded files securely.

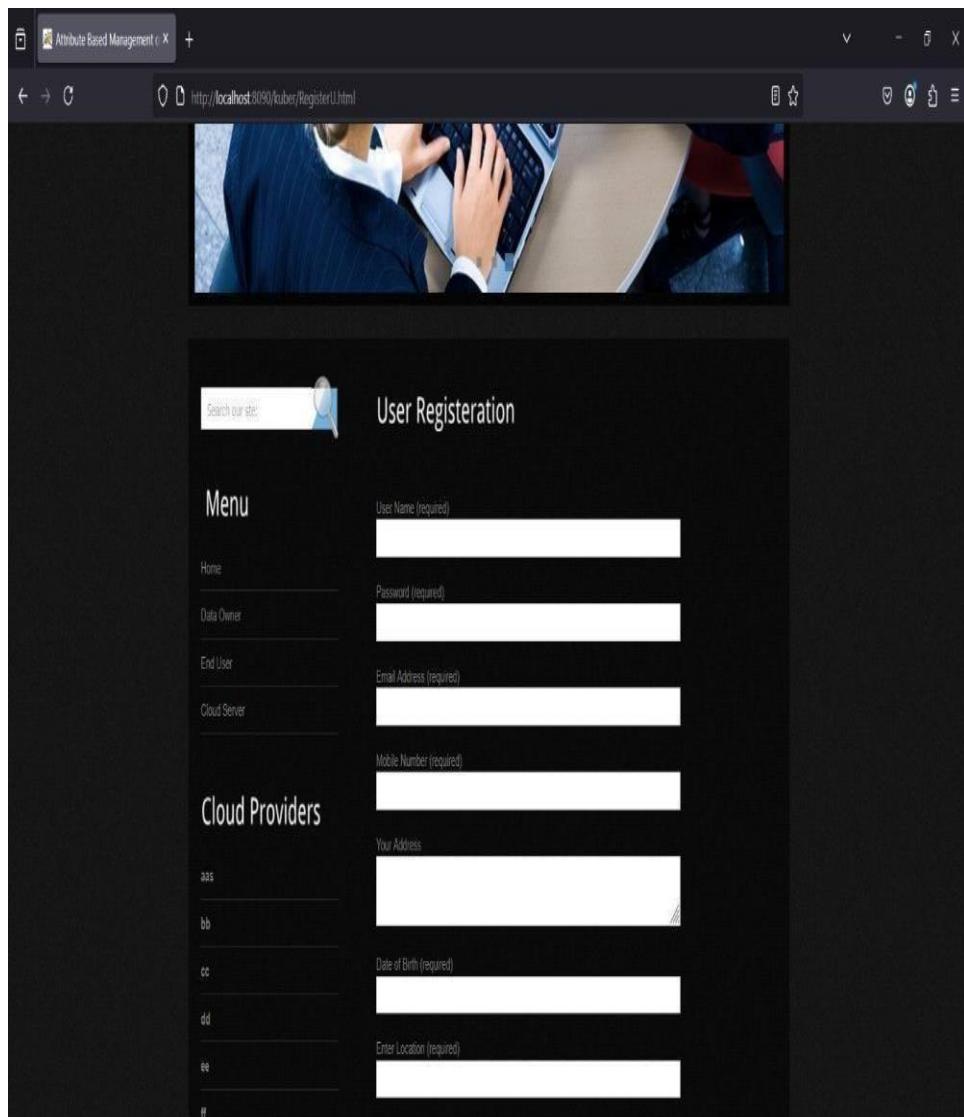
Delete Files



Screen no 5.1.8: Delete Files

Description: Securely delete files with attribute-based access control in Kubernetes cloud bursting.

User Registration



The screenshot shows a web browser window with the title "Attribute Based Management" and the URL "http://localhost:8080/kuber/RegisterU.html". The page has a dark background with a photo of a person in a suit at a keyboard at the top. On the left, there's a sidebar with a search bar and a menu with options: Home, Data Owner, End User, and Cloud Server. The main content area is titled "User Registration". It contains several input fields:

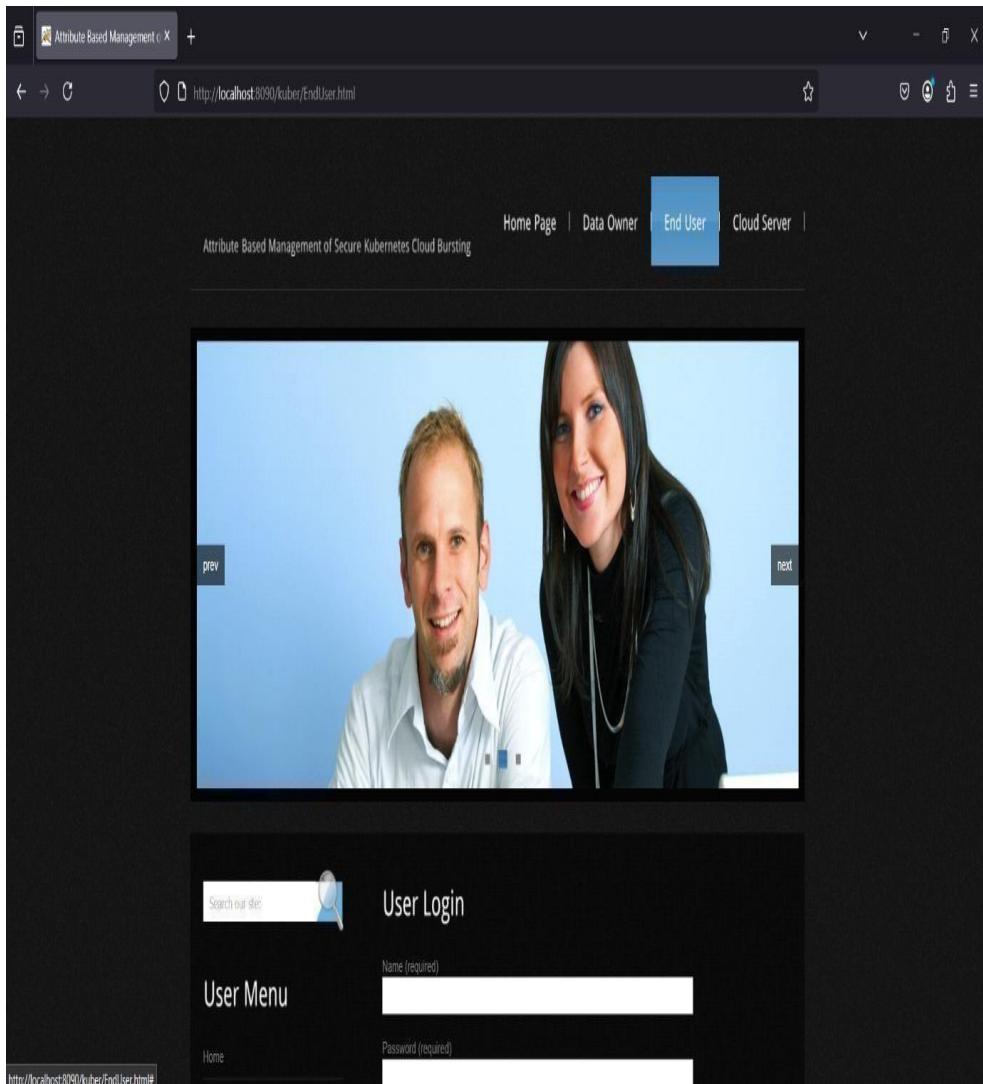
- User Name (required) - input field
- Password (required) - input field
- Email Address (required) - input field
- Mobile Number (required) - input field
- Your Address - input field
- Date of Birth (required) - input field
- Enter Location (required) - input field

Below the address input field, there are two columns of short entries: "aa", "bb", "cc", "dd", "ee", and "ff".

Screen no 5.1.9: User Registration

Description: Register users securely for Kubernetes.

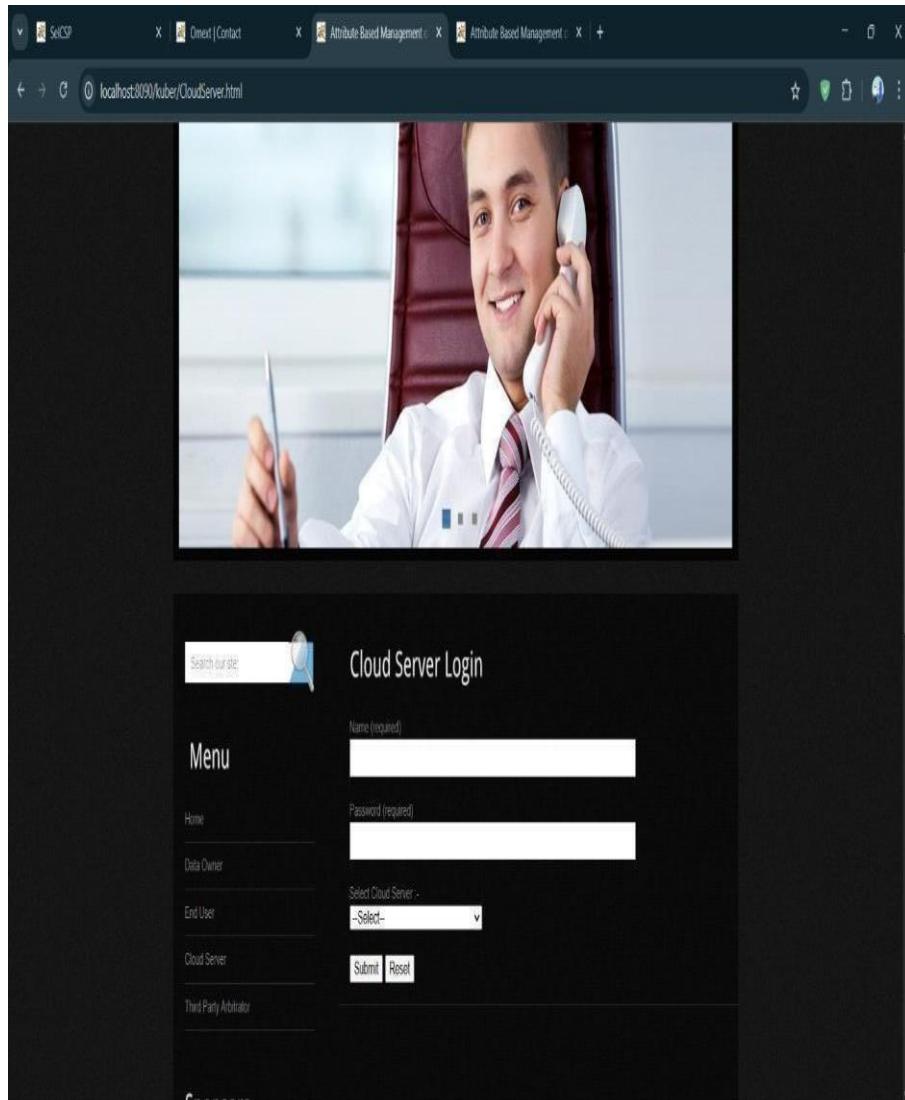
User login



Screen no 5.1.10: User login

Description: Login securely for Kubernetes.

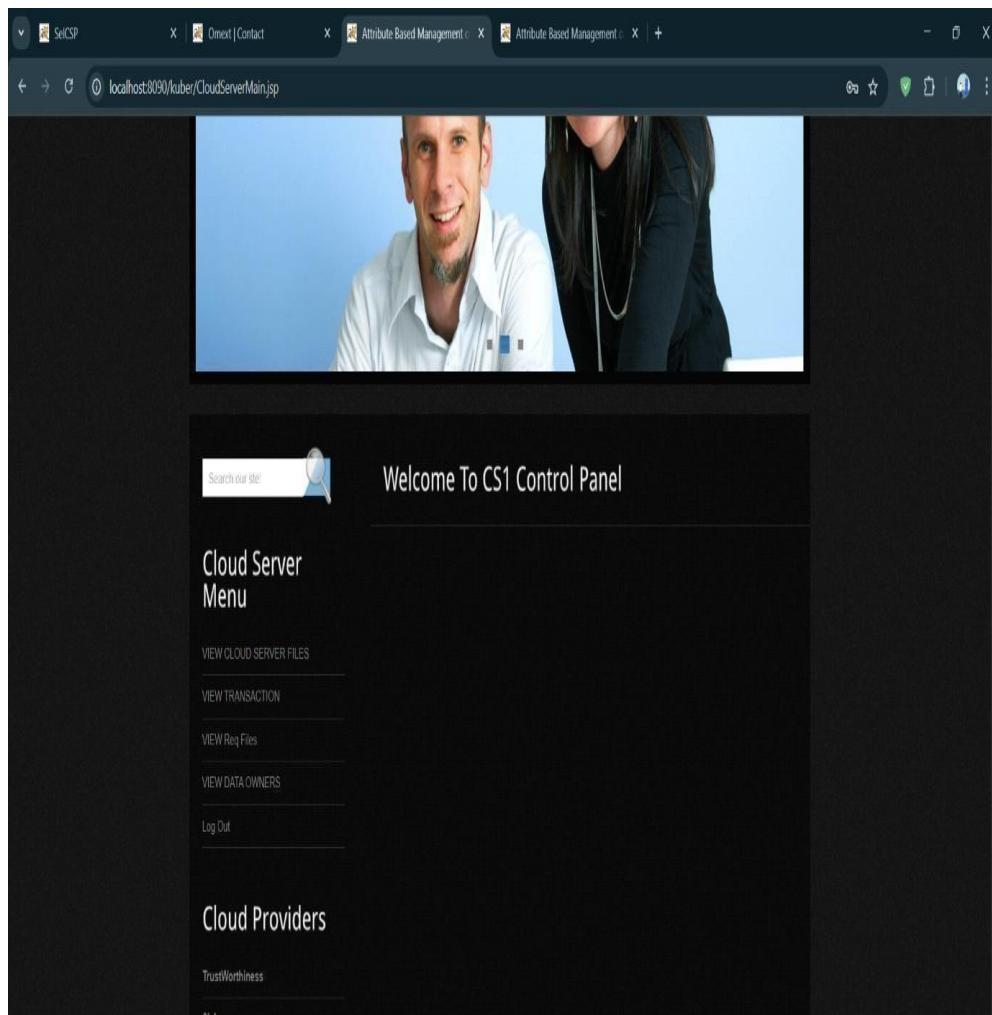
Cloud Server Login



Screen no 5.1.11: Cloud Server Login

Description: Login securely to cloud servers for Kubernetes.

Welcome to CS1 Control Panel



Screen no 5.1.12: Welcome to CS1 Control Panel

Description: Welcome to CS1 for secure Kubernetes cloud bursting with attribute-based access control.

User available files

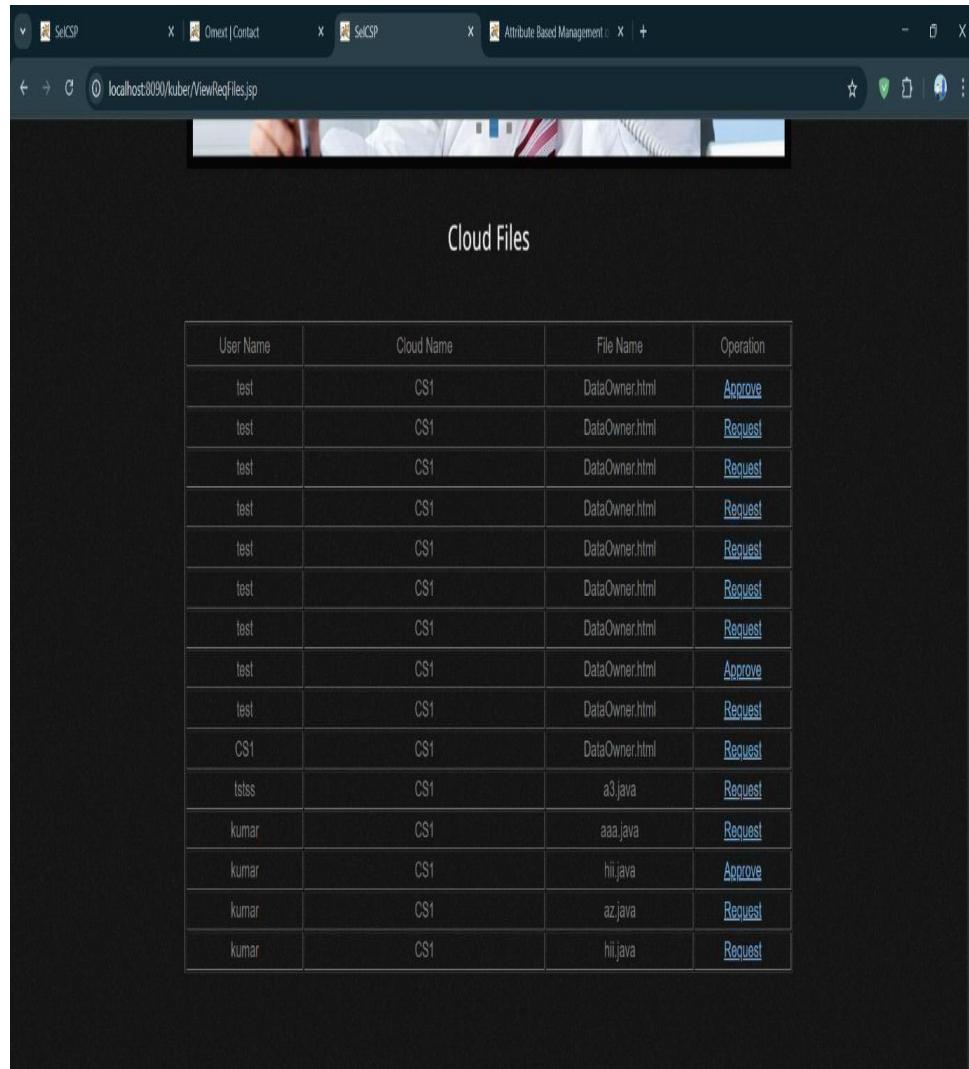
The screenshot shows a web browser window with the URL <http://localhost:8900/kuber/availablefile.jsp>. The page displays a photograph of a man and a woman smiling. Below the image, the text "Available Files of null" is visible. A table lists the available files:

File Name	Cloud Name
tpauth.jsp	CS1
userauth.jsp	CS1
test.java	CS1
tst.java	CS1
aaa.java	CS1
bbb.java	CS1
bbb.java	CS1

Screen no:5.1.13 User available files

Description : User now can see which files are available in the server.

Cloud Files



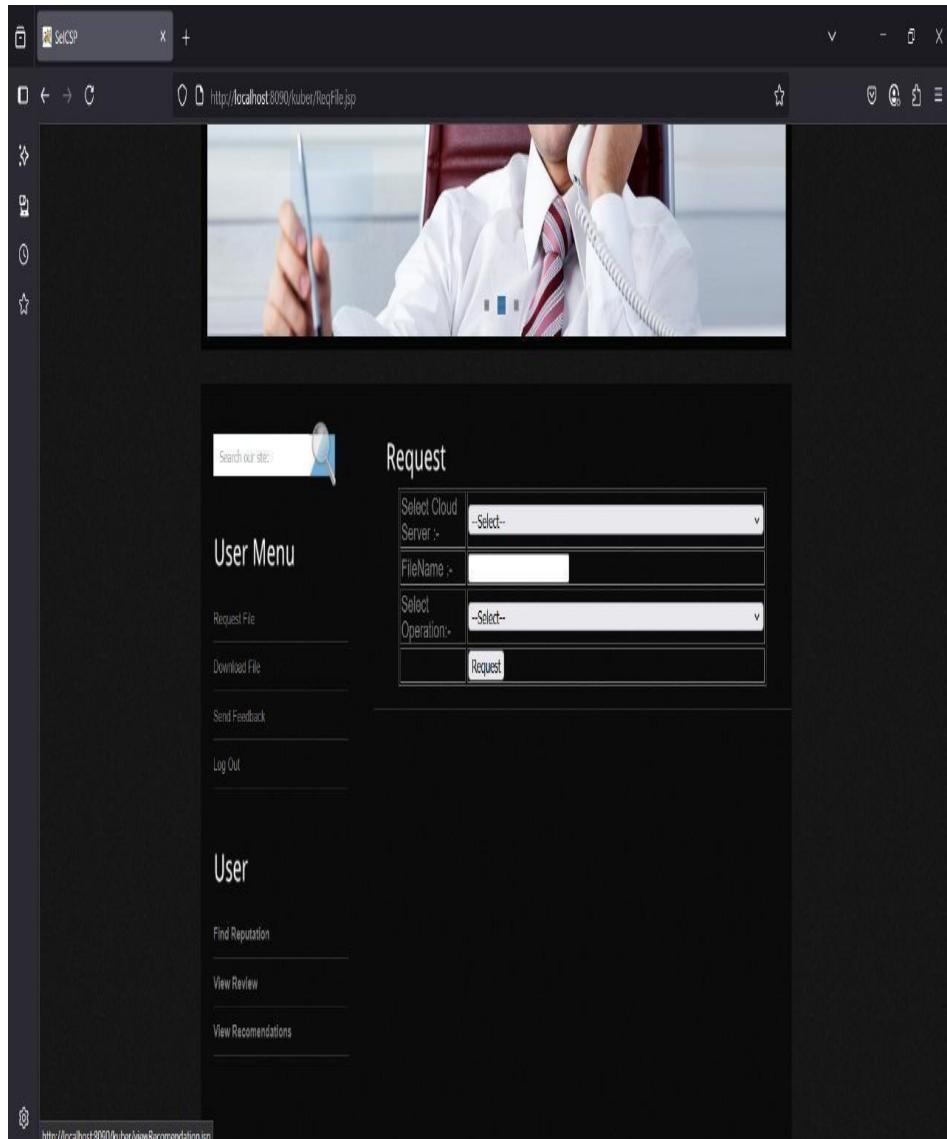
The screenshot shows a web browser window with three tabs open: 'SelCSP', 'Omext | Contact', and 'Attribute Based Management'. The active tab is 'Attribute Based Management' with the URL 'localhost:8090/kuber/ViewReqFiles.jsp'. The page title is 'Cloud Files'. Below the title is a table with the following data:

User Name	Cloud Name	File Name	Operation
test	CS1	DataOwner.html	Approve
test	CS1	DataOwner.html	Request
test	CS1	DataOwner.html	Request
test	CS1	DataOwner.html	Request
test	CS1	DataOwner.html	Request
test	CS1	DataOwner.html	Request
test	CS1	DataOwner.html	Approve
test	CS1	DataOwner.html	Request
CS1	CS1	DataOwner.html	Request
tsstss	CS1	a3.java	Request
kumar	CS1	aaa.java	Request
kumar	CS1	hi.java	Approve
kumar	CS1	az.java	Request
kumar	CS1	hi.java	Request

Screen no 5.1.14: Cloud Files

Description: Manage and secure cloud files for Kubernetes.

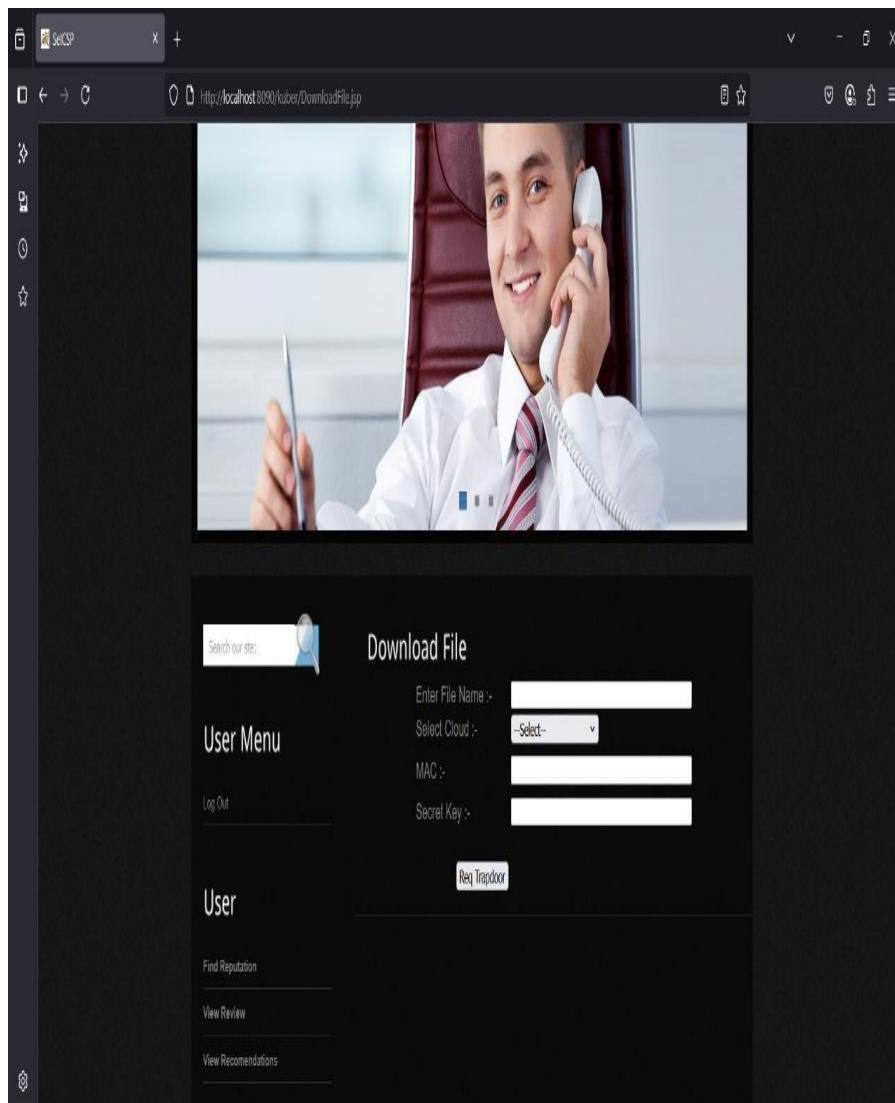
Request Files



Screen no 5.1.15: Request Files

Description: User Requesting Files from cloud.

Download File



Screen no 5.1.16: Download Files

Description: User now can download the files from cloud.

Conclusion

Finally, I conclude that to introduce a robust orchestration scheme tailored for secure cloud bursting, and to maintain the complexities, cost challenges, and stringent data protection compliance requirements by harnessing the combined capabilities of K8s and ABE. By incorporating ABE, our approach achieves granular encryption and provides cloud resources with suitable confidentiality. At the same time, cloud bursting empowers the extension of computational tasks beyond the scope of a local primary cloud environment. The synergy of these two technologies establishes a cohesive management framework, guaranteeing secure access to bursting services and streamlined deployment of excess work loads to the cloud, all facilitated by Kubernetes. Significantly, our contributions encompass the design blueprint for an attribute-based cloud bursting authorization system within Kubernetes.

Project Work Mapping With PO's

Pos	PO1	PO2	PO3	PO4	PO5	PO6	PO7	PO8	PO9	PO10	PO11	PO12
Mapping Level	3	3	3	3	3	2	1	2	3	2	3	2

PO1: Computational Knowledge

PO2: Problem Analysis

PO3: Design/Development of Solutions

PO4: Conduct Investigations of Complex Computing Problems

PO5: Modern Tool Usage

PO6: Professional Ethics **PO7:**

Life-Long Learning

PO8: Project Management and Finance

PO9: Communication Efficacy

PO10: Societal and Environmental concern

PO11: Individual Team Work

PO12: Innovation and Entrepreneurship

Project Mapping with Sustainable Development Goals

S. No	Name of the Sustainable Development Goal	Description of the Goal	Mapping	Justification
1	No Poverty	It could support it indirectly by enabling cost-effective, secure infrastructure.		
2	Zero Hunger	Initiatives by providing a secure, scalable back end infrastructure for food-tech.		
3	Good Health and Well-being	Providing a secure, scalable infrastructure for hosting health-related digital platform.		
4	Quality Education	Secure and scalable cloud environment to support digital learning platforms and educational technologies.	✓	This infrastructure allows educational platforms to handle large user loads securely and efficiently, ensuring uninterrupted access to learning resources.
5	Gender Equality	Providing secure and scalable infrastructure for digital platforms that promote inclusive access and empowerment.	✓	It enables reliable and secure digital services that can be used to support initiatives focused on women's empowerment, education, and equal participation.

6	Clean Water and Sanitation	Indirectly supports Clean Water and Sanitation by providing secure, scalable infrastructure.		
7	Affordable and Clean Energy	Enabling secure, scalable cloud infrastructure for energy management and monitoring systems.		
8	Decent Work and Economic Growth	Enabling secure, scalable cloud infrastructure that boosts business agility and innovation.		
9	Industry, Innovation, and Infrastructure	advances Industry, Innovation, and Infrastructure by providing secure, scalable, and efficient cloud infrastructure management through attribute-based access controls and dynamic cloud bursting.	✓	It fosters innovation and resilient infrastructure by enabling businesses to securely scale their applications and optimize cloud resource use.
10	Reduced Inequality	Access to digital services through controlled and attribute-based cloud resource management.		
11	Sustainable Cities and Communities	providing secure, scalable infrastructure for smart city applications and urban data platforms.	✓	It enables reliable and secure digital systems that can support urban planning, transportation, public safety, and other smart city

				services.
12	Responsible Consumption and Production	Optimizing cloud resource usage through secure and efficient workload management.	✓	It promotes efficient use of computing resources, reducing energy waste and supporting sustainable IT operations.
13	Climate Action	Can reduce energy consumption and lower carbon footprints.	✓	By optimizing workload distribution and minimizing unnecessary resource usage, it contributes to greener and more sustainable cloud operations.
14	Life Below Water	Focuses on cloud infrastructure, access control, and workload management, which are unrelated to marine ecosystems, ocean conservation, or aquatic biodiversity.		
15	Life on Land	Powering digital platforms for environmental monitoring and land management.		
16	Peace, Justice, and Strong Institutions	Supports Peace, Justice, and Strong Institutions by providing secure, transparent, and controlled access to digital government and legal platforms.	✓	By using attribute-based access control, it helps ensure data integrity, privacy, and accountability in systems that support public services and institutions.
17	Partnerships for the Goals	Can be shared across organizations to support collaborative digital initiatives.	✓	It facilitates interoperability and resource sharing, helping governments, institutions, and private sectors collaborate through trusted digital platforms.

7. BIBLIOGRAPHY

Appendix-A

URL Listing

Websites	Data collected
https://wikipedia.org	Searching of any information that will be used in documentation.
https://dev.sqlserver.com/doc	SQL server it performing in mainly depending on the one of the database using.
https://www.answers.com	Answers.com, online dictionary, encyclopedia and much more.
https://google.co.in	Any information searching and downloading.
https://training-classes.com	Designing part information as gathered.

REFERENCES

K. Graves, Ceh: Official certified ethical hacker review guide: Exam 312-50. John Wiley & Sons, 2007.

R. Christopher, “Port scanning techniques and the defense against them,” SANS Institute, 2001.

M. Baykara, R. Das , and I. Karadoğan, “Bilgisayar sistemlerinde kullanılan araçların incelenmesi,” in 1st International Symposium on Digital Forensics and Security (ISDFS13), 2013, pp. 231–239.

Rashmi T V. “Predicting the System Failures Using Machine Learning Algorithms”. International Journal of Advanced

Scientific Innovation, vol. 1, no. 1, Dec. 2020,
doi:10.5281/zenodo.4641686.

S. Robertson, E. V. Siegel, M. Miller, and S. J. Stolfo, “Surveillance detection in high bandwidth environments,” in DARPA Information Survivability Conference and Exposition, 2003. Proceedings, vol. 1. IEEE, 2003, pp. 130–138.

K. Ibrahimi and M. Ouaddane, “Management of intrusion detection systems based-kdd99: Analysis with lda and pca,” in Wireless Networks and Mobile Communications (WINCOM), 2017 International Conference on. IEEE, 2017, pp. 1–6.

Girish L, Rao SKN (2020) “Quantifying sensitivity and performance degradation of virtual machines using machine learning.”,Journal of Computational and Theoretical Nanoscience, Volume 17, Numbers 9-10, September/October 2020, pp.4055-4060(6) <https://doi.org/10.1166/jctn.2020.9019>

L. Sun, T. Anthony, H. Z. Xia, J. Chen, X. Huang, and Y. Zhang, “Detection and classification of malicious patterns in network traffic using benford’s law,” in Asia-Pacific Signal and Information Processing Association Annual Summit and Conference (APSIPA ASC), 2017. IEEE, 2017, pp. 864–872.

S. M. Almansob and S. S. Lomte, “Addressing challenges for intrusion detection system using naive bayes and pca algorithm,” in Convergence in Technology (I2CT), 2017 2nd International Conference for. IEEE, 2017, pp. 565–568.

Appendix-B

• GLOSSARY

- GUI : Graphical User Interface
- UML : Unified Modeling Language
- API : Application Programming Interface
- HTML : Hyper Text Markup Language
- URL : Uniform Resource Locator
- ODBC : Open Database Connectivity

Appendix-C

List of Figures

S.NO	Description	Page No	Chapter
1	Architecture	07	Introduction
2	Product Perspective	08	SRS
3	Characteristics of Product Perspective	09	SRS
4	SDLC	15	SRS
5	Water fall model	19	SRS
6	Working of Java Program	25	SRS
7	Implementation of java virtual machine	26	SRS
8	Program running on the java platform	27	SRS
9	ER-Model	41	System Design
10	Use case diagram for overall project	48	System Design
11	Class diagram for overall project	49	System Design
12	Sequence diagram for overall project	51	System Design
13	Activity diagram for overall project	53	System Design
14	Component diagram for For overall project	54	System Design
15	Collaboration diagram for For overall project	56	System Design
16	Deployment diagram for For overall project	57	System Design

List of Tables

S. No	Table Name	Page No	Chapter
1	1 NF	44	System Design
2	2 NF	45	System Design
3	Test Cases	66	Testing

List of Screens

S.No.	Screen No.	Screen Name	Page No	Chapter
1	5.1.1	Home Page	67	Implementation
2	5.1.2	Data Owner Registration	68	Implementation
3	5.1.3	Data Owner Login	69	Implementation
4	5.1.4	Welcome to Data Owner	70	Implementation
5	5.1.5	Owner Details	71	Implementation
6	5.1.6	Upload Data	72	Implementation
7	5.1.7	Verify Files	73	Implementation
8	5.1.8	Delete Files	74	Implementation
9	5.1.9	User Registration	75	Implementation
10	5.1.10	User Login	76	Implementation
11	5.1.11	Cloud Server Login	77	Implementation
12	5.1.12	Welcome to CS1 Control Panel	78	Implementation
13	5.1.13	User available files	79	Implementation
14	5.1.14	Cloud Files	80	Implementation
15	5.1.15	Request Files	81	Implementation
16	5.1.16	Download files	82	Implementation

Appendix-D

Help Document

1. Installation of java:

- Go to <http://www.oracle.com/technetwork/java/javase/downloads/index.html>.
- click on JDK DOWNLOAD button. run the exe file and then follow the instruction given in wizard.
- **To set up the path:-**

- Right click on my pc and then go to my properties

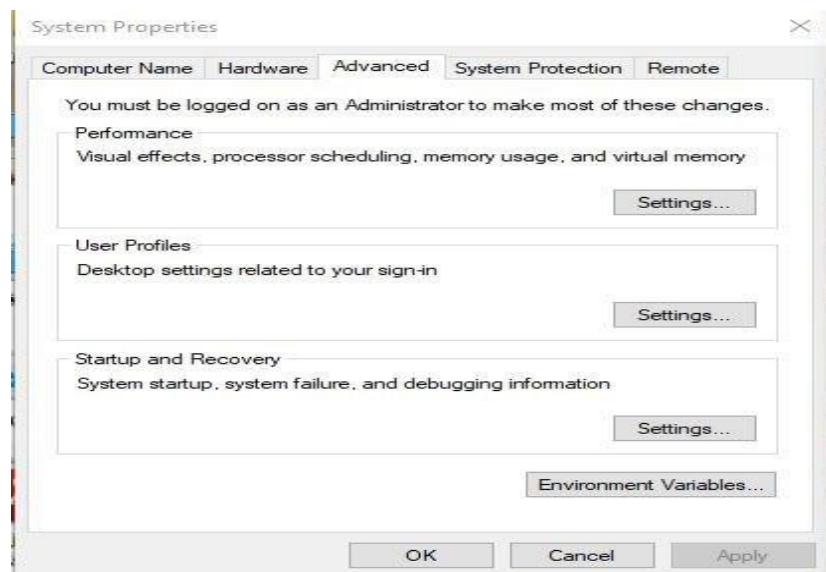


Fig : Properties wizard

- Go to advanced settings and then click on environment variables.
- create a class path and copy the path of the java folder where it is located in program files.

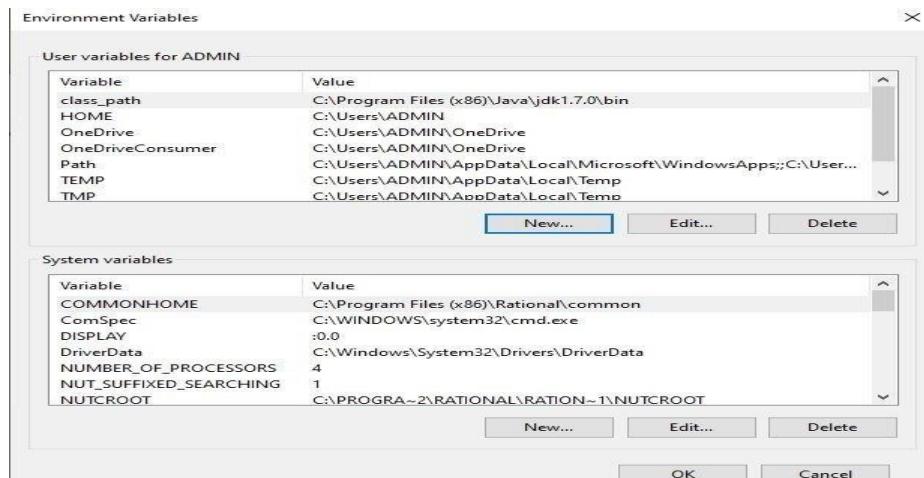


Fig : Path setting for java

2. Installation and setup of Apache Tomcat:

- Go to <http://tomcat.apache.org/index.html> and click on download latest versions.
- Run the exe file and click on next and follow the wizard instructions.



Fig : Welcome Page of Tomcat

- Click on install with port number 8090 with username and password as aits and aits.
- Mention the connection port as 8090 and then click on next and finally click on finish.

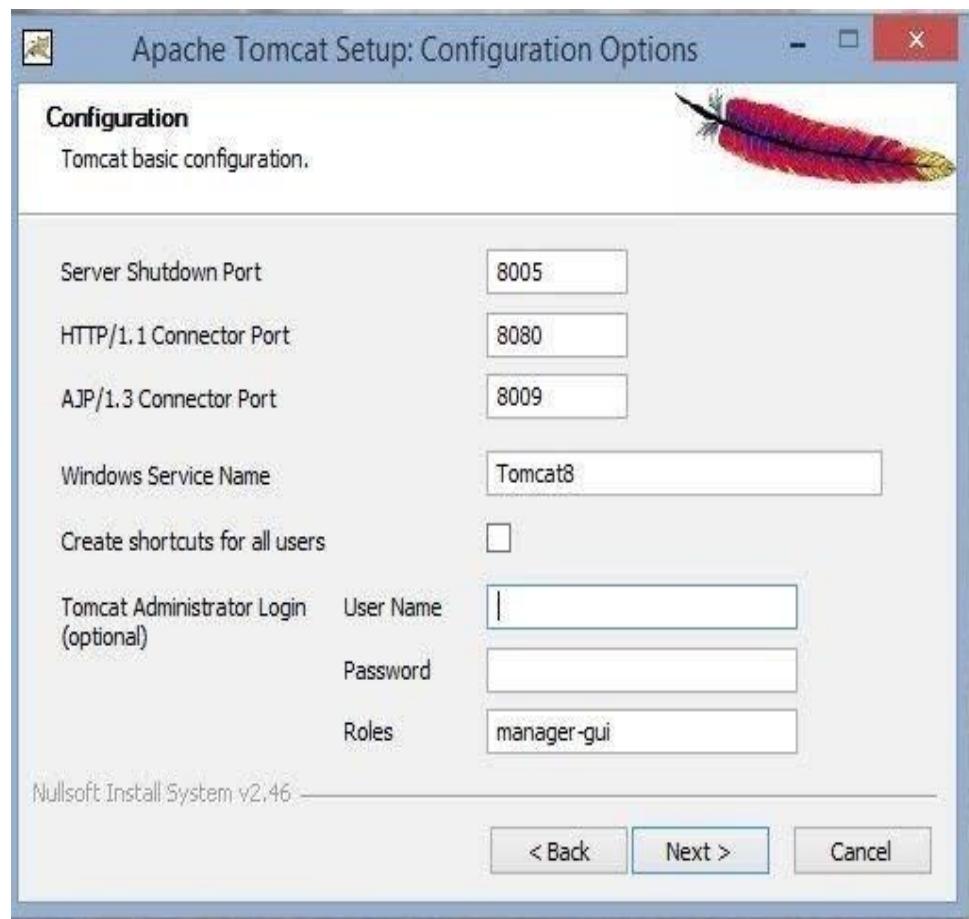


Fig : Tomcat Configuration Options Page

- Click on I agree button in license agreement in order to accept the terms and condition.

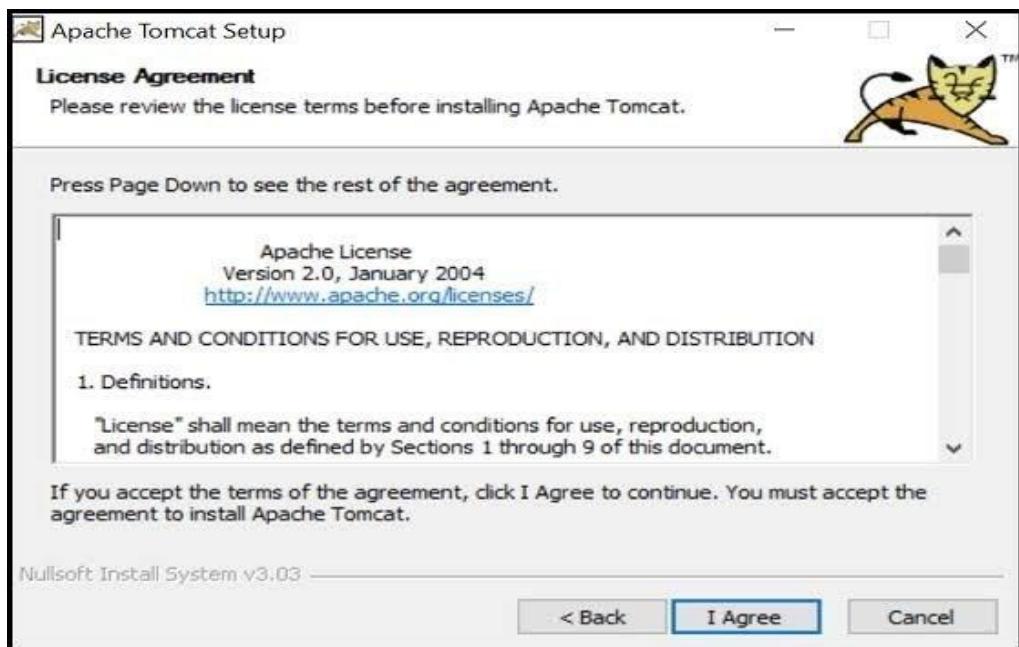


Fig : Tomcat License Agreement

3. Installation and setup of SQL:

- Go to <http://dev.mywql.com/downloads/> . and click on install button.
- After completion of installation, click on exe file and then click on next.
- Run the MySQL setup and click on next and follow the instruction in wizard.



Fig : Welcome wizard of MySQL

- Conform the type as typical and then click on next and follow the instructions.

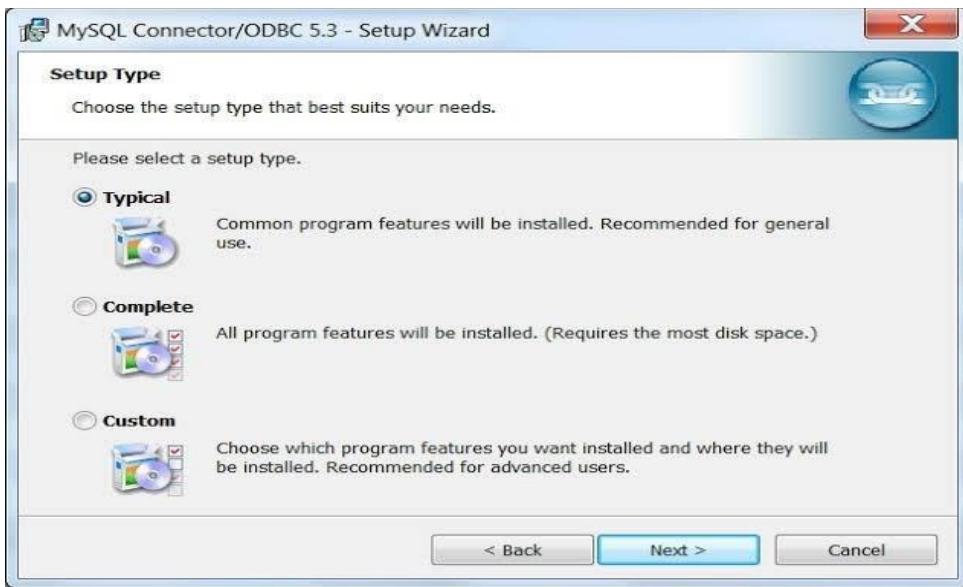


Fig : SQL Setup Wizard

- Now confirm the password as root in system settings field and then click on finish.

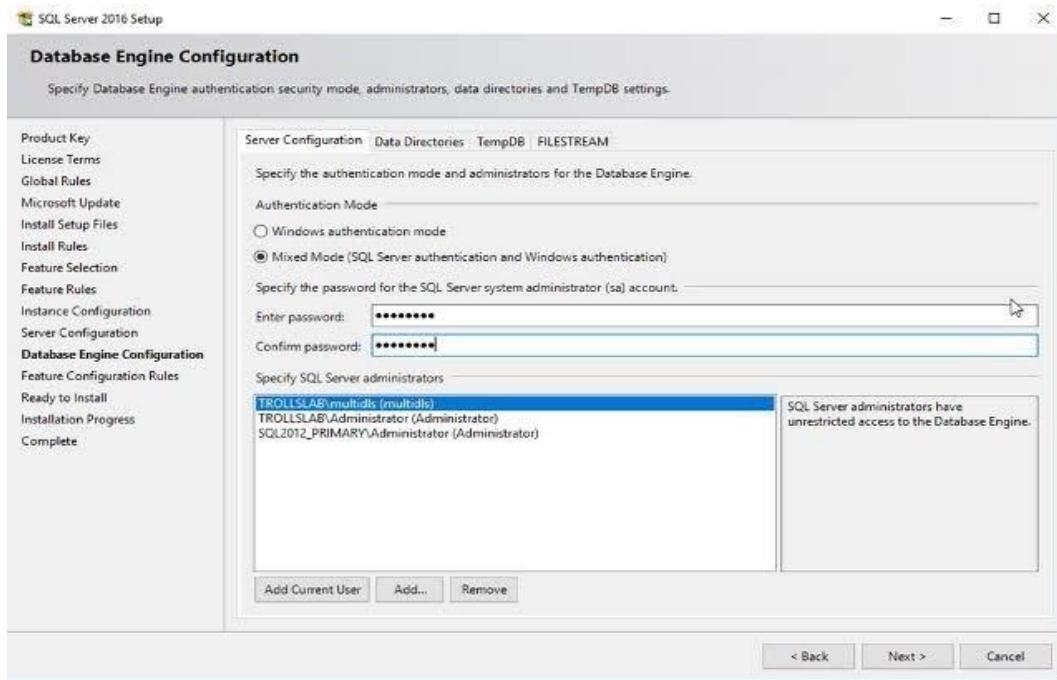


Fig : Database Configuration Engine

Attribute-Based Management of Secure Kubernetes Cloud Bursting

MAURO FEMMINELLA^{1,2} (Member, IEEE), MARTINA PALMUCCI³,
GIANLUCA REALI^{1,2}, AND MATTIA RENGO¹

¹Department of Engineering, University of Perugia, 06125 Perugia, Italy

²Consorzio Nazionale Interuniversitario per le Telecomunicazioni, 43124 Parma, Italy

³Consortium GARR, 00185 Rome, Italy

CORRESPONDING AUTHOR: M. FEMMINELLA (e-mail: mauro.femminella@unipg.it)

This work was supported in part by the European Union - NextGenerationEU under the Italian Ministry of University and Research (MUR) National Innovation Ecosystem under Grant ECS00000041 (VITALITY), and in part by the MUR Extended Partnerships under Grant PE0000001 (RESTART).

ABSTRACT In modern cloud computing, the need for flexible and scalable orchestration of services, combined with robust security measures, is paramount. In this paper, we propose an innovative approach for managing secure cloud bursting in Kubernetes, combining Attribute-Based Encryption (ABE) with Kubernetes labeling. Our model addresses the challenges of complexity, cost, and data protection compliance by leveraging both Kubernetes and ABE. We introduce an attribute-based bursting component that uses Kubernetes labels for orchestration, and an encryption component that employs ABE for data protection. This unified management model ensures data confidentiality while enabling efficient cloud bursting. Our approach combines the strengths of label-based orchestration with fine-grained encryption, providing a technologically advanced yet user-friendly solution for secure cloud bursting. We present a proof-of-concept implementation that demonstrates the feasibility and effectiveness of our model. Our approach offers a unified solution that complies with security and privacy laws while meeting the needs of contemporary cloud-based systems.

INDEX TERMS Cloud bursting, orchestration, attribute-based encryption, Kubernetes.

I. INTRODUCTION

IN RECENT years, the proliferation of virtualization and containerization technologies has led to a significant increase in the complexity of distributed systems, as cloud computing systems. As organizations strive to achieve efficient resource management and scalability, Kubernetes (K8s) [27] has emerged as the most popular solution for orchestrating resources in such complex systems. For example, it is integrated into platforms such as Amazon Elastic Kubernetes Service (EKS) [1], Google Kubernetes Engine (GKE) [2], Azure Kubernetes Service (AKS) [50], IBM Cloud Kubernetes Service [3], and Oracle Container Engine for Kubernetes (OKE) [5].

This paper aims at exploring the challenges associated with cloud bursting, which allows private cloud services to use public cloud resources when local resources are exhausted or for any other reasons. Specifically, we address

the configuration issues associated with service request load management over a hybrid cloud system including both private and public components. The orchestration of such a heterogeneous system presents a number of challenges, such as optimal management of typically large volume of resources, variable operating conditions, security issues, and compliance with local regulations. In order to address these challenges, resorting to attribute-based management policies is regarded as a valid approach [9].

Attributes are intrinsic characteristics or properties related to the entities, workload, or resources to manage. They can refer to different aspects, such as computational requirements, security levels, data location, and other relevant information. Recent findings indicate that the related management challenges can be effectively addressed through the utilization of an attribute-based approach, which has been found to be preferred over the conventional

role-based methods. Rather than depending solely on pre-defined roles, attributes can include a broader array of qualities and attributes associated with users, resources, or data. This level of flexibility, adaptability, and precision in access control renders them more suitable for scenarios characterized by a diverse, dynamic, and intricate range of access requirements [16], [62]. In particular, in this paper, we leverage the attribute-based approach to easily orchestrate load distribution and resource allocation between private and public clouds during the cloud bursting process.

Attribute-based policies can be enforced by different technologies. In this paper we make use of Kubernetes, since today it is one of the most appreciate tools for managing distributed systems, especially in the context of cloud computing. It provides flexibility though different built-in components and tools. Among them, the usage of labels and label selectors can be exploited to simplify cloud bursting operations. While Kubernetes best practices recommend that labels be assigned semantic meanings before being used [9], there is currently no standardized method for enforcing this practice. Our goal is to develop a systematic approach in the context of cloud bursting that ensures semantic meaning associated with the generic Kubernetes label concept.

Furthermore, it emerged that Kubernetes management does not suitably address all the security aspects related to data confidentiality and access controls, which are central in cloud bursting [6]. Kubernetes incorporates access management, but it requires separate configuration processes that are decoupled from the logic of the orchestrated functions. Moreover, the existing access management mechanisms in Kubernetes have certain limitations in terms of managing complex authorization scenarios and are constrained by their policy scope. Hence, these limitations are challenging for achieving comprehensive and secure resource management in the context of cloud bursting. To overcome these limitations, in this paper we propose an architectural solution to address the security challenges of cloud bursting that integrates the Kubernetes orchestration with attribute-based encryption.

In more detail, our specific contributions are as follows:

- 1) Association of semantic meaning with key labels, giving them the role of attributes. This approach adds context to the cloud bursting configuration and improves label comprehension.
- 2) Leveraging the Attribute-Based Encryption (ABE) technology, deployed through a cloud service, to improve security levels, in terms of data privacy, confidentiality, and access control, through fine-grained policies. This ABE component works seamlessly within the Kubernetes environment, aligning with attribute logic and improving overall system security.
- 3) Simplification and speed-up of configuration based on a unified management layer that handles holistically all attributes, including the ones directly used by Kubernetes and by the ABE module. This unification

ensures transparency and ease of use for administrators, eliminating the need for separate configurations and additional harmonization functions, as well as a streamlined experience.

In synthesis, we propose a secure cloud bursting scheme that improves both efficiency and security in resource management over legacy solutions by incorporating semantic labels, cryptographic technology, and a unified attribute configuration approach.

The organization of the paper is outlined as follows. In Section II, we delve into an examination of the existing literature within the domain. Moving on to Section III, we furnish a comprehensive introduction to the foundational concepts. Our proposed methodology is detailed in Section IV. The outcomes of the validation of the proposed approach carried out through a proof of concept implementation are presented in Section V. Finally, Section VI includes our closing remarks and conclusions.

II. RELATED WORKS

A. SECURITY SERVICES FOR CLOUD SYSTEMS

When it is necessary to move data to the cloud, it is critical to ensure security and flexible, granular control over file access. This can be efficiently done through ABE. However, user revocation is a significant issue in ABE. In [47], the authors propose a ciphertext-policy ABE (CP-ABE) scheme with efficient user revocation for cloud storage system. User revocation is handled by introducing the concept of user group, with the rule of updating private keys of the users remaining in the group when any other user leaves it. In addition, since the computation cost of CP-ABE grows linearly with the complexity of the access structure, in order to mitigate it they propose to offload high computation demand to cloud service providers without leaking file content and secret keys. They prove that the proposed scheme can withstand collusion attack performed by the revoked users cooperating with the remaining ones. A similar approach, which requires the update of the unrevoked users' keys, is proposed in [66]. It is based on the use of a group manager to accomplish this task. It also applies re-encryption technology to prevent the revoked users from decrypting ciphertexts.

However, since the correctness of outsourcing computing results is difficult to guarantee, this approach often requires resorting to the blockchain technology for obtaining such guarantees [53]. Blockchain is used also in [41] to solve the key escrow problem by replacing the traditional key authority with a blockchain. Keys are generated collaboratively by users and the blockchain, thus preventing the latter from obtaining them alone. Alternatively, multi-authority solutions can be used to securely delete data in cloud [48], where data sharing policies can be a challenging issue.

The recent proposal by Chen et al., RABE-DI [33], is an ABE scheme capable of addressing a different problem, namely protecting data integrity after user revocation, ensuring better efficiency compared to state-of-the-art proposals.

The proposal in [44] is based on a different approach. It consists of an efficient and provably secure cloud data sharing scheme (ABEDS-RR) using CP-ABE. The scheme makes use of a semi-trusted proxy party to transform part of the ciphertext with a conversion key. Compared with most existing schemes, when the attributes of a user are changed, only the private key and conversion key of that user need to be modified, leaving the other users' keys and all ciphertext unchanged. This improves efficiency of attribute revocation and update. ABEDS-RR eliminates key escrow by generating keys for users jointly by the key authority and the cloud service provider. Hence, neither party can obtain the private key of the user without collusion with each other. In addition, the weight attributes are also used to enhance the expression ability of attributes. However, this approach seems to be more suitable to a resource-constrained environment.

In order to address the timely decryption of data stored in cloud servers, the authors in [32] propose a cooperative approach based on CP-ABE to decrypt the ciphertext. Semi-authorized users cooperatively perform decryption task, making the scheme very efficient in term of computing resources and storage cost.

Bera et al. [22] propose an approach named Attribute-Based verifiable Data Storage and data Retrieval Scheme (ABDSRS) for cloud environments. It employs an attribute-based online-offline mechanism, in which only authorized data owners can anonymously upload data to the cloud. In addition, a data user can perform searching operations over encrypted data by using keyword policy. ABDSRS allows a user to verify the correctness of the search results in cloud without interacting with any authority. Thus, in addition of being lightweight, provably secure, and guaranteeing fine-grained access control, it offers further functions in terms of anonymity, privacy, secure search and verification of results.

Ahuja and Mohanty [15] propose an extension of CP-ABE in order to provide shared access privileges among users and delegation of access privileges in a flexible manner, without sacrificing scalability and fine-grained access control. The proposed solution merges CP-ABE with a hierarchical structure [64] to achieve scalability by decentralizing the key issuing authority at different levels of hierarchy. In more detail, lower level users get secret keys from the users that have a higher position in the hierarchy. The scheme results to be resistant to cheating and collusion attacks.

Repetto et al. [55] proposed a methodological approach for designing and implementing heterogeneous security services for distributed systems. The framework utilizes a hybrid architecture based on Attribute-Based Access Control (ABAC), ABE, and blockchain technology to provide secure and efficient access control in decentralized and distributed environments. ABE cryptographic primitives are specifically used to extend the ABAC functions. They implement an online authorization procedure, support time-limited authorization, protect against collusion attacks, and protect user privacy. Such features had previously been investigated by Sciancalepore et al. in the paper [59].

Lu et al. [49] propose a policy-driven approach to secure data sharing using an integration of ABAC and ABE. Private data is shared in ciphertext form between edge nodes to mitigate potential security problems such as privacy leakage.

All the papers [49], [55], [59] propose integrating ABE and ABAC to improve security of existing solutions. However, none of them fits into the context of resource and service orchestration.

Finally, Chiquito et al. in [34] survey attribute-based approaches for access control to data, focusing on policy management and enforcement. The paper aims at identifying the key properties provided by ABAC and ABE that can be used to control data access to prevent leakage to unauthorized users while providing easy-to-manage policies. To achieve this goal, they identify knowledge gaps. As for ABE, the main identified limitations consist in implementation difficulty and lack of expressive and easy-to-manage access structures. Covering these gaps with pure cryptographic approaches is challenging due to the limited information and mechanisms that can be embedded in ciphertext and keys. To address these gaps, an integration with ABAC is considered with a policy decision point, in order to take advantage of flexibility and expressiveness of the ABAC policy languages. In this case, the main issue results in the technical limitations of translating more complex policies into the ABE access structures, beyond other technical weaknesses. In addition, also in this case, no direct mapping with service and resource orchestration exists.

B. CLOUD BURSTING

A number of proprietary solutions for cloud bursting are made available by vendors. However, most of them just focus on integration between on-premise Kubernetes clusters and elastic public computing services provided by the vendor, mainly dealing with configuration issues [10], [39], [67]. All these solutions, which clearly suffer from vendor lock-in issues, overlook some of the main security issues, especially related to data privacy and protection. Most of them provide basic yet effective solutions for securing the communication between on-premise and public clouds. A popular approach consists of using the TLS-based encryption of transmitted data in both user and control plane (see, e.g., [13]). An interesting add-on for hybrid deployment is represented by Aporeto [12], now integrated in the Prisma Cloud [8]. It proposes a Zero Trust security solution based on Kubernetes Network Policies and Role-Based Access Control (RBAC). Even the whitepaper by Forrester [28], which addresses security issues and best practices for Kubernetes, focuses on identity, network, and container security, neglecting other security weaknesses on data management. The key point is that Kubernetes is not secure by design, but security is supported through the collaboration across cloud-native ecosystems.

Below, we examine the approaches used by the vendors for implementing cloud bursting in Kubernetes. Virtual Kubelet [14] is an Open Source project within the Cloud

Native Computing Foundation (CNCF)'s sandbox project. It provides a Kubernetes Kubelet that masquerades a remote cloud provider, thus extending the computing capacity of the on-premise cluster as a *virtual* worker node. When a pod is scheduled onto this virtual worker node, Virtual Kubelet makes a public cloud instance available on-demand to run it. This workaround is necessary since most on-premise enterprise Kubernetes platforms do not implement nor support the capabilities of Cluster Autoscaler [11], used to dynamically add or remove worker nodes or node pools for matching the current cluster utilization. As such, manual operation for adding a worker node in a public cloud provider to an on-premise Kubernetes cluster often does not match the requirements of cloud bursting operations. Virtual Kubelet allows enforcing access control rules with RBAC, whereas security of communications depends on capabilities of public cloud providers.

A different approach is followed by KubeSlice [57], which creates a virtual cluster across multiple clusters (on-premise and public clouds). This is realized by creating logical application *slice* boundaries, which allow pods and services to seamlessly communicate. It provides centralized control management through KubeSlice Hub, and communications between remote pods of the same slice are encrypted by the KubeSlice gateways running on different clusters. Access control is enforced via RBAC.

A fully centralized management is provided by Alibaba Cloud through its Distributed Cloud Container Platform for Kubernetes (ACK One) [67]. It is a distributed cloud container platform able to jointly control on-premise clusters, Alibaba Cloud resources, and third party resources. The management platform runs in the ACK cloud. It makes use of RBAC for access control and encrypted communications between clusters, in both control and user plane.

CloudBees Build Acceleration follows a different approach. Although it relies on centralized management of all clusters, including those on-premise and in public clouds, the Cluster Manager can be installed in *any* server. Computing clusters are called *cloudburst agents*, over which resources are deployed on demand (*agent cloud bursting* operation). Each worker in the cluster has to run an *Electric Agent* to correctly execute bursting operations. It supports access control via RBAC and TLS-based encryption between agents.

Finally, Anthos [7] is a solution provided by Google which can integrate multiple clusters. It is based on the GitOps concepts, which is an operational framework that takes DevOps best practices, developed for application development, such as version control, collaboration, compliance, and CI/CD, and uses them for infrastructure automation. As such, it synchronizes cluster configurations via a component named Anthos Config Management, which can manage bursting operations. Anthos makes use of RBAC for managing access control to resources. Communication security of both control and user plane is enforced by encrypted service mesh via TLS.

As a final note, commercial solutions for cloud bursting typically rely on a limited set of monitored metrics like CPU/RAM occupation, and are actionable when one or more of them exceeds user-defined thresholds. This aspect is addressed in the literature related to cloud bursting in Kubernetes, which is mainly focused on the criteria used to offload workload to public clouds [17]. This decision can be based on either heuristic procedures, or the solution of a specific optimization problem, even found by a machine learning engine, or pre-determined criteria. In any case, a portion of the *critical* workload is handled on-premise, and the other is assigned to a public cloud [18], [56]. Significant results have been achieved in task scheduling, specifically in the allocation of computing resources among various Kubernetes clusters. Although most approaches have been designed according to the edge-to-cloud offloading paradigm, many of these proposals can be applied also to cloud bursting operations between on-premise and public clouds. The most relevant proposals about Kubernetes scheduling in these multi-clusters scenarios can be classified in two main categories: a) proposals based on machines learning/artificial intelligence, used to take decisions about the replicas or the portion of workload to be offloaded to the public cloud, and b) proposals based on classic optimization algorithms. An interesting example of the first category is taken from mobile networks, in particular the edge-to-cloud deployment in 5G networks [36]. In this paper, Faticanti et al. make use of machine learning techniques to realize a proactive offloading of tasks by anticipating peak utilization of Kubernetes pods based on pattern recognition from historical data. Also reinforcement learning can be used to anticipate overload conditions and schedule further replicas in advance [20]. As for the second category, Carmona-Cejudo et al. [29] address the offloading of computing tasks in Kubernetes in a multi-tier architecture through the solution of an optimization problem. The recent survey in [43] provides a systematic overview of offloading solutions using both traditional optimization and machine learning techniques. Finally, the interested reader can find further details on different approaches in Kubernetes scheduling in the survey [30].

To sum up, a crucial step towards achieving a secure orchestration approach that covers both resource orchestration and enhanced security is the development of a solution that combines these two topics. Our contribution aims at filling the gap in achieving both aspects. This can be done by integrating the two topics through a solution that leverages Kubernetes labeling and ABE. In particular, access to resources and services is granted by using the same attributes that are used for their orchestration.

III. BACKGROUND

In this section, we illustrate three fundamental components of our proposal. The first component pertains to Ciphertext-Policy Attribute-Based Encryption (CP-ABE), which is a

cryptographic primitive that allowed us to provide confidentiality of data related to services deployed in the cloud. The second component concerns *Cloud bursting*, which is a technique that allowed us to increase the available computing capacity in the primary cloud environment by resorting to public cloud resources. Finally, we discuss K8s, which is a platform designed for automating deployment, scaling, and management of containerized applications. Collectively, these three technologies form the knowledge foundation for our model.

A. CIPHERTEXT-POLICY ATTRIBUTE-BASED ENCRYPTION

ABE [58] is an extension of the public key encryption that allows for precise control over access to encrypted data. By ABE, access policies are established based on user attributes or characteristics, which are usually represented as sets of attributes instead of individual user identities. These attributes can consist of any information regarding users, including their position, membership, or clearance level. For example, only users with the correct combination of attributes that satisfy the established access policy can be allowed decrypting and accessing data. ABE provides fine-grained access control based on user attributes. Unlike other access control mechanisms, ABE ensures confidentiality of data through the use of encryption. Consequently, ABE requires less reliance on third parties data storage services. Moreover, ABE functions enables a high level of user independence, as authorized users can independently obtain access by decrypting ciphertexts once they possess the appropriate secret key. In this way, users can avoid interacting further with any policy-enforcing entities that may be unavailable at the time of an access request [63].

CP-ABE [23] is one of the two main types of ABE. In CP-ABE, access structures are integrated into the ciphertexts, with the encrypting user being in charge of defining the access structures. The Key Generation Authority (KGA) generates secret keys that are linked to specific user attributes. Decrypting users can decrypt a ciphertext only if they possess an authorized secret key. If the set of attributes associated with a secret key satisfies the access policy attached to the ciphertext, then the set is authorized.

Alternatively, in key-policy ABE (KP-ABE) [45] data are encrypted over a set of attributes and user keys allow accessing a tree which can distinguish attributes. However, since in KP-ABE data include only user attributes, anyone can access data, whereas CP-ABE have a direct control of data access. Given that our application scenario involves the association of attributes with users and access policies to cloud resources, our focus is on CP-ABE.

In spite of its advantages, ABE has the limitation of high computational complexity. In order to mitigate this limitation, a hybrid mode can be implemented using CP-ABE. In this approach, the message is initially encrypted with a randomly generated symmetric secret key. Only this session key is then encrypted using CP-ABE under the access

policy. This technique aims at minimizing the computational complexity of the overall encryption process while still enabling fine-grained access control.

In conclusion, it is clear that, by leveraging the ABE technology, it is possible to improve security levels in various aspects. Specifically, ABE provides enhanced data privacy by allowing fine-grained access control based on user attributes, ensuring that only authorized users can access encrypted data. Additionally, ABE ensures data confidentiality by protecting sensitive data from unauthorized access, even when they are stored or transferred through untrusted media. Moreover, ABE attribute-based access control enables dynamic and flexible access management, allowing administrators to define complex access policies that can adapt to variable user attributes and security requirements. Deploying an ABE module as a cloud service allows its seamless interworking within Kubernetes clusters, aligning with attribute logic and improving overall system security.

B. CLOUD BURSTING

Cloud bursting is a popular strategy used to manage a significant increase in the demand for computational resources. The primary components of a cloud bursting solution include:

- a local environment, also referred to as *local* or *on-premise (on-prem)* cloud, which hosts the organization's resources;
- one or more secondary clouds, that are typically public clouds, providing on-demand, additional resources to expand the local environment when overloaded;
- a decision engine that autonomously manages cloud bursting while adhering to the company's policies.

The need of service flexibility and scalability in response to changes in cloud capacity requirements is the major driver of cloud bursting. Cloud bursting has to be executed in background, so that users do not encounter any service disruptions.

Some milestone papers [40], [42], [61] illustrate the evolution of cloud bursting. However, the solutions therein proposed became rapidly obsolete as they essentially rely on virtual machines, whilst current approaches have evolved to the use of containers and functions. Although their limitations, those pioneering papers still offer a robust logical foundation for the development of novel solutions.

Due to its capabilities in resource orchestration [24], cloud bursting strategies may include Kubernetes as a flexible and scalable solution to handle the process of expanding service deployment to the cloud as well as dynamically assigning resources in response to changing requirements [60]. A number of Web resources and technical reports clearly describe cloud bursting operations using Kubernetes (see, e.g., [37], [38]). Most of them discuss the general features that a cloud bursting solution should offer and relevant issues, which are typically related to networking and security aspects.

C. KUBERNETES

Kubernetes, also referred to as K8s, is an open-source platform for container orchestration. It was developed by Google based on its internal system Borg [27]. The aim was to simplify the deployment, scaling, and management of containerized applications. Users utilize declarative configuration files to specify the state of their application, and K8s is in charge of maintaining the desired state. In this way, by the use of Kubernetes, several management issues of complex and distributed applications have reduced significantly.

A K8s cluster is a distributed set of master and worker nodes. One or more master nodes manage the cluster's state and configuration, while worker nodes run containerized applications. Due to its flexibility, Kubernetes has emerged as the most popular and widely used orchestrator among numerous alternatives.

In addition, to ensure compliance with the user-defined cluster state, specified in the architectural YAML manifest files or *blueprint*, Kubernetes also includes a centralized management interface for managing regular orchestrator tasks. The process is divided into some steps, such as setting up the required network and storage resources, deploying service containers, monitoring their health, and replacing the compromised ones. Furthermore, Kubernetes offers automatic load balancing, service discovery, and horizontal scaling capabilities, which allows it to seamlessly route external traffic to the appropriate services within the cluster.

IV. ATTRIBUTE-BASED MANAGEMENT OF SECURE KUBERNETES CLOUD BURSTING

Our research targets the development of a service orchestration model using Kubernetes, which involves two key innovative aspects. First, we use Kubernetes labeling to design and implement an attribute-based orchestration system for cloud bursting. Labels, which are associated with services, are used to orchestrate service deployment in cluster nodes. Some nodes may have multiple labels with preferred policies to run multiple services; others may have just one label, while others may not even have any at all. This approach shapes a cloud bursting policy tailored to each service, which can make an optimal use of the local cluster resources before requesting additional paid resources in external secondary clouds. Second, we use ABE policies defined by referring to the *same* labels, as mentioned above. By using the different labels as attributes in ABE, we have improved security without adding any extra complexity and also obtained a more refined and selective cloud bursting. This approach not only provides confidentiality, but also enables seamless integration of ABE with the existing label-based orchestration system. The final result is a sophisticated orchestration system that conforms to data protection regulations and relies on attributes as the basis of its operations, which we denote Attributed-Based management of secure K8s cloud bursting.

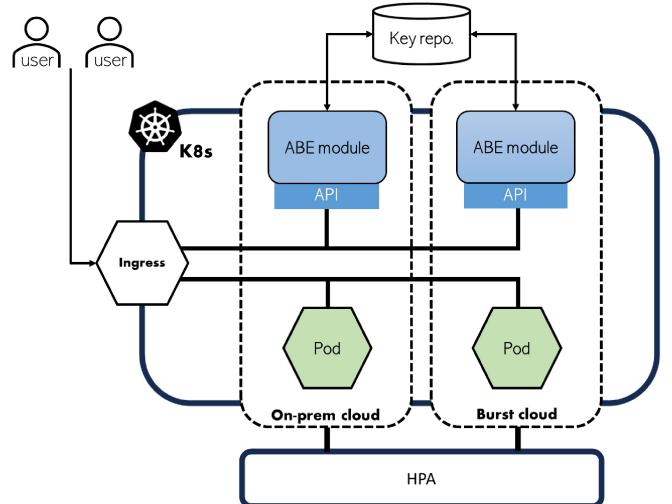


FIGURE 1. High-level diagram that illustrates the primary stakeholders and their interconnections of our model.

It is worth highlighting that K8s access control, which is the basic element of K8s security, includes both RBAC and ABAC. However, although the ABAC's flexibility is appealing in different contexts, RBAC is typically preferred over ABAC. The main reason for this is that ABAC is difficult to manage and understand in very complex distributed environments. Our proposal, which has simplicity of implementation as its key element, allows us reintroducing the flexibility of using attribute-based management in a context managed mainly with RBAC.

Below, we illustrate the structure of our proposed model, capable of achieving the aforementioned objectives.

Figure 1 depicts a high-level diagram that illustrates the primary elements of the system considered and their interconnections:

- *Users*. They access services through the K8s *Ingress* API and are represented by the user icon on the left side of Figure 1. Further, each user has a set of attributes that enable them to access only specific services, which will be discussed in what follows.
- *Clouds*. A cloud is any distributed and virtualized networked computing infrastructure that allows making use of shared pools of reconfigurable computing resources on-demand, which can be easily provisioned and withdrawn. Clouds integration within K8s provide the capability to deploy and manage containerized applications in various public cloud environments, such as Amazon Web Services (AWS), Google Cloud Platform (GCP), and Microsoft Azure. By leveraging cloud-native services, Kubernetes facilitates seamless scaling, high availability, and the efficient utilization of underlying infrastructure resources.
- *Kubernetes*. As mentioned above, this open-source and flexible platform for container orchestration is efficient for automating container operations and simplifies the

orchestration of multi-container applications, ensuring that they are always running and matching the desired state. For this reason, it is a central component of our solution. Below, we outline a step-by-step workflow of how our proposed system interacts with key Kubernetes and ABE components:

- **Label Assignment and Management:** Labels, as key-value pairs, are assigned to resources (nodes and services) in Kubernetes clusters. Such labels are used not only for resource grouping and management but also play a crucial role in the attribute-based bursting and encryption process.
- **Kubernetes Metrics Server:** This component provides real-time data on resource usage by pods and nodes. It acts as the primary source of information for the HPA to make scaling decisions.
- **HorizontalPodAutoscaler (HPA):** It monitors resource usage and performance metrics of pods. Based on these metrics and the labels assigned, it makes informed decisions about scaling pod replicas up or down. This is where our attribute-based approach intersects with the Kubernetes native scaling mechanism that adjusts the number of replicas in a Deployment or ReplicaSet to meet the required performance. HPA analyzes the metrics provided by the *Metrics Server* or external monitoring systems in order to determine whether the current resource allocation aligns with the configured scaling thresholds.
- **Ingress:** It acts as an application-level load balancer that manages external access to services deployed over a cluster. Specifically, it runs an HTTP and HTTPS reverse proxy service that routes network traffic from clients to the appropriate backend services. By using Ingress APIs, complex network configurations can be managed easily, since numerous routing rules can be combined into a single resource. Ingress resources utilize an ingress controller, such as NGINX or HAProxy, to manage the routing of HTTP and HTTPS requests to backend services. Additionally, the introduction of the Ingress resource simplifies the implementation of features like SSL termination, load balancing, and name-based virtual hosting.
- **ABE Policy Enforcement:** For each service request routed through the ingress to services inside pods, the ABE module checks the attributes (aligned with the Kubernetes labels) against the defined ABE policies. This step ensures that only authorized users or services can access or interact with the data, thus maintaining data confidentiality.
- **Key Management and Data Encryption:** The Key repository plays an essential role in managing the encryption keys used in the ABE scheme. It interacts with the ABE module to provide keys for

encrypting and decrypting data, thus securing the data at rest and in transit.

- **Data Flow and Access Control:** In summary, request routing within Kubernetes Ingress adheres to the K8s's label-based rules, while data access security is managed by the ABE module leveraging the label based attributes. This dual application of labeling enhances the overall operation efficiency. In fact, while the Ingress effectively manages routing of requests, ensuring they reach the correct service endpoints, the ABE module, using these same labels, determines whether a request should be granted access to the data based on the attributes of the requesting entity. This ensures that each service and the information within are accessible only under the appropriate conditions, adhering to both performance and security requirements.

- **ABE module.** This custom cloud service implements the ABE scheme in the proposed solution. It handles all the functions of an ABE scheme and manages the necessary attribute-based encryption keys. Furthermore, it establishes communication with the key repository.
- **Key repository.** It maintains the keys used to protect the volumes of services. It communicates with the ABE module.

In the subsequent sections, we describe the two primary components of our proposed solution: attribute-based bursting and encryption components.

A. ATTRIBUTE-BASED BURSTING COMPONENT

In complex systems, multiple service layers are typically involved. Their management often requires cross-cutting operations that can disrupt their strict hierarchical structure, particularly when that structure is rigidly determined by the underlying infrastructure rather than by user needs. Kubernetes uses labeling to manage such systems. Our model leverages Kubernetes labeling not just as mere identifiers, but as pivotal elements in orchestrating resource allocation and managing service demands. Labels are assigned to both services and nodes, and dictate how resources are allocated and managed within our cloud bursting ecosystem. Labeling is particularly important when it refers to autoscaling and load balancing. It allows the HPA to differentiate between resources and apply scaling policies based on the characteristics of the labels assigned to each resource.

In our proposal, the labels used in the HPA are also used as user attributes by the ABE module. In order to guarantee that services fulfill their requirements, both HPA and cloud nodes are configured with the same label values. For this purpose, it was necessary to take advantage of the concepts of *affinity* and *anti-affinity* in Kubernetes, as they influence scheduling of pods on cluster nodes:

- **Affinity** refers to the preference for scheduling pods on nodes that meet certain criteria. For example, a particular pod could have affinity to nodes with specific labels or that are located in a particular availability zone.

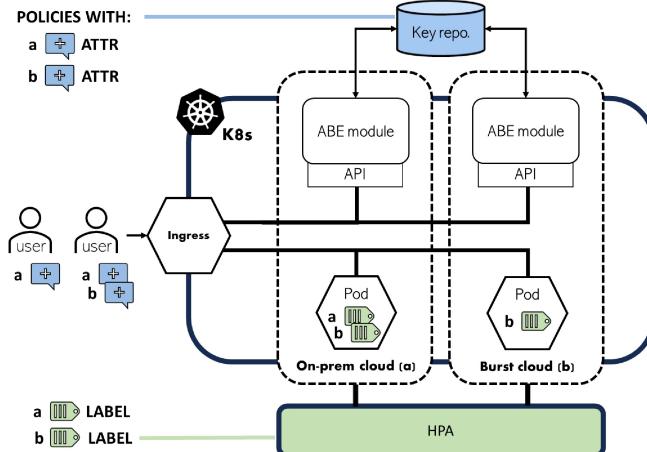


FIGURE 2. Attribute-based bursting and encryption components.

- *Anti-affinity*, on the other hand, refers to the preference for pods not to be scheduled on nodes that meet certain criteria. For example, it may be desirable that two replicas of a pod are not scheduled onto the same node for availability purposes.

Kubernetes allows setting affinity and anti-affinity values to express these types of preferences.

With reference to our system, a node with a specific label (e.g., ‘On-Prem’, as in Figure 2) is selected by the HPA to run the deployed services until there are no resources left in the local cloud environment, ensuring application demands are met while keeping operational expenses as low as possible.

In our proposal, the labels used for configuring Kubernetes orchestration for cloud bursting are also used as attributes used for defining ABE policies. Understanding the QoS needs for each service is crucial in any orchestration model, and in our case the Kubernetes labels are strategically used to manage both the QoS and security aspects of each service. Labels dictate not only allocation and management of resources according to QoS requirements but also form the basis of the corresponding ABE policies. This ensures seamless integration of performance management and security. For example, if two services, namely Service S_1 and Service S_2 , are designed with distinct QoS and security requirements, this is reflected in their respective HPA labels and ABE policies.

This concept can be better illustrated by referring to Figure 2. In this scenario, two services are designed with variable bursting and replication requirements: service S_1 , with loose requirements, and service S_2 , with more strict requirements. The system setup includes one local and one external cloud, namely *On-prem cloud* and *Burst cloud*. The latter is utilized for cloud bursting. The HPA is configured to meet the specific requirements of each service, which is achieved by associating labels with each service based on their requirements. It is important to note that multiple labels may be applied to each service in practical scenarios. Label a is applied to service S_1 , labels a and b to service S_2 . In

order to guarantee that services meet their requirements, it is necessary to configure pod affinity and anti-affinity on nodes within the cluster. In Figure 2, the nodes in the local cloud are assigned both labels, namely a and b , to enable all services to run on local nodes with affinity to every label in the system. On the other hand, nodes in Burst cloud are only assigned the label b , indicating that the nodes in Burst cloud are configured with affinity to the label b . As a result, a service that matches the required affinity only with labels assigned to local nodes can only scale out using local resources without any bursting. In contrast, a service deployment that matches affinity with both labels can be executed, duplicated and possibly leverage cloud bursting. It is worth noting that anti-affinity is not used in this example.

From a different perspective, we can say that S_1 has the strong requirement of being run only in the *on-prem cloud*, even at the cost of potentially degraded performance. By operating exclusively within the on-prem cloud environment, S_1 benefits from inherently enhanced security, as its data remain in the controlled on-prem environment. This service is tailored for operations where data security and sovereignty are paramount. Its ABE policy is accordingly structured to leverage the inherent security benefits of the on-prem environment, using a set of attributes optimized for controlled access while maintaining operational efficiency. Service S_2 , on the other hand, is designed with a focus on providing superior scalability in order to keep its performance compliant with its requirements. Thus, this service is capable of operating not just on-prem, but also in external cloud environments, allowing it to scale resources dynamically based on demand. Its ABE policy is more complex, incorporating a wider range of attributes to ensure secure access across diverse environments while maintaining high availability and performance standards. This approach allows orchestration and authorization to share the same attributes and solve the issues caused by managing multiple sets of attributes. Thus, usage of labels allows implementing sophisticated and heterogeneous service policies.

As mentioned above, one of the main features of our solution is the introduction of a semantic meaning to be associated with a set of labels used in Kubernetes. In this way, these labels can be used also as ABE attributes. Clearly, it is not necessary for all labels used in Kubernetes to be used as attributes in ABE. Our system is designed to leverage such mapping onto labels whose meaning is relevant to both security and orchestration management. For example, consider a Kubernetes label called a “node-type”. This label can take values like “on-prem”, “burst”, “a”, “b” and so on. These values can be used directly as attributes in the ABE system. This approach enables the formulation of ABE security policies in accordance with the Kubernetes orchestration context. This mapping between selected Kubernetes labels and ABE attributes enables simplified management of both resource orchestration and data security, giving system administrators integrated control of these aspects.

By using the example of S_1 and S_2 above and Figure 2, we detail how our cloud-bursting architecture with ABE-based security works:

- 1) A user requests access to a service through the Kubernetes Ingress.
- 2) The Ingress controller routes the request to the appropriate service based on its label.
- 3) The ABE policy associated with the service label is used to authorize the user access. If the service is scaled to the burst cloud, the HPA scales the service based on its label and the current resource utilization. The ABE policy associated with the service label is used to authorize the user access.
- 4) The Kubernetes Ingress continues routing the user to the correct services regardless of whether it is deployed within the hybrid cloud, in the on-prem cloud or in the burst cloud.

Our cloud bursting architecture with ABE-based security offers several advantages, including:

- Performance and scalability: By dynamically scaling resources to the burst cloud, our architecture can improve performance and scalability of applications with fluctuating workloads.
- Enhanced security: ABE policies enable fine-grained access control to services, even across multiple clouds, protecting sensitive data from unauthorized access.
- Simplified management: By leveraging joint labeling for both orchestration and authorization, our architecture simplifies management and reduces complexity.
- Increased flexibility: Labels provide a versatile mechanism for implementing sophisticated and heterogeneous service policies, meeting the diverse needs of modern cloud-based applications.

B. ATTRIBUTE-BASED ENCRYPTION COMPONENT

In the previous sections, we mentioned that our objective is not only to propose a unified orchestration, but also to offer a solution able to address some security issues, such as confidentiality, access control, and compliance to General Data Protection Regulation (GDPR) rules [35]. For example, when bursting is carried out using the infrastructure of a public cloud provider, it is important to protect user data, in order to avoid privacy issues for those services that require them. In this regard, ABE natively supports the decoupling of encryption keys from third parties (i.e., the infrastructure provider), thus ensuring that only users with the right attributes can access the protected service data, even if hosted on public clouds and regardless of local regulations that may affect the cloud service provider policies. Thus, the joint management of labels for both security and orchestration management can help mitigating confidentiality leakage for critical data. In addition, since ABE allows accessing contents through policies based on the owned attributes and not on the identity, the proposed approach can be used not

only to ensure confidentiality, but also to protect anonymity while accessing data outside the private cloud deployment. The level of this protection depends on several factors, such as user population, number of attributes, derived policies, and so on.

1) HYBRID CP-ABE FOR SECURE AND EFFICIENT CLOUD VOLUME ENCRYPTION

To address this objective, we propose utilizing CP-ABE in a hybrid mode in order to secure sensitive volumes. CP-ABE offers advantages such as fine-grained access control and adaptable policies. The hybrid mode balances security and efficiency by combining asymmetric and symmetric encryption, providing robust data protection while minimizing computational overhead.

In simple terms, the process illustrated below is followed. Managers of a specific service encrypt the entire volume or a portion of the associated volume. This encryption employs a symmetric cipher, with the encryption key referred to as the data encryption key (DEK). The service provider then securely stores the DEK in a designated key repository. Importantly, the DEK is not stored in plain text; rather, it is encrypted using CP-ABE. The policy selected during DEK encryption dictates who can decrypt the key, thereby obtaining access to the corresponding service.

On the user side, system participants possess various attributes. Utilizing these attributes, the system provides them with distinct decryption keys — referred to as *user keys* — enabling access to specific portions of the key repository. Essentially, only users possessing attributes that satisfy the encryption policy are authorized to access the DEK and subsequently retrieve data from the volumes.

The sequence diagram of ABE setup and operation is depicted in Figure 3. This protocol facilitates secure communication and information exchange among users, the ABE module, the key repository, and services. It serves as a safeguard, shielding users from untrusted service solutions. In fact, the only entity with which the user interacts and shares sensitive information is the ABE module — a trusted party.

Users must recover the DEKs stored in the key repository to access encrypted data within service volumes. This process consists of three phases: service setup, user registration, and service data access.

During the *service setup* phase, the service first generates a DEK, encrypts it using ABE, and stores it securely in the Key Repository within the ABE cloud service. Concurrently, the *user setup* phase consists of user registration, during which the user provides the attributes obtained from a system authority. The service may perform some operational actions that we call generic AAA actions, indicated as Admission control in the figure. They are additional security related actions that do not fall in the scope of our proposal, but often used in operation. For example, some preliminary traffic filtering can be used to mitigate potential attacks, such as

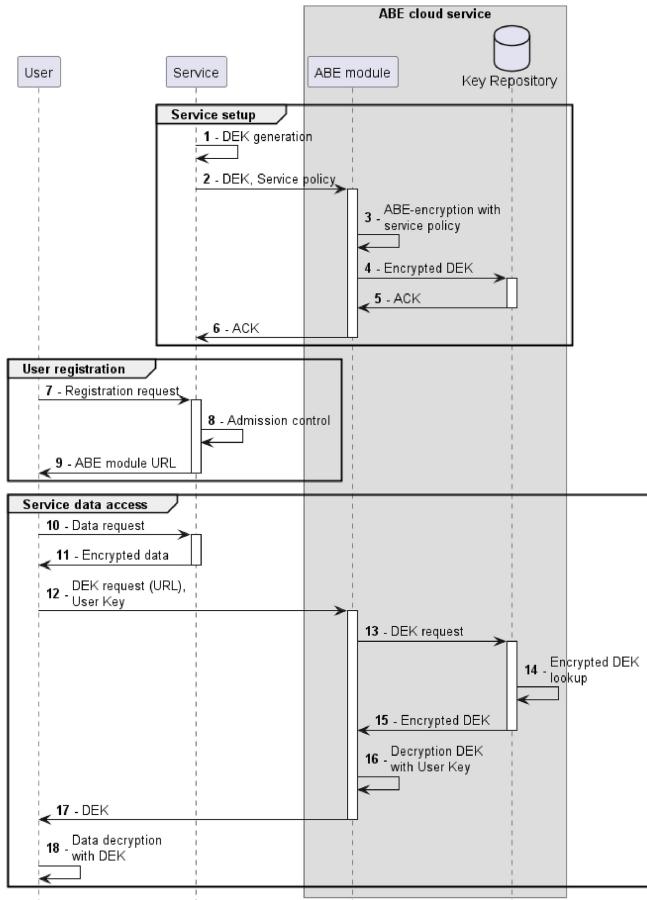


FIGURE 3. Sequence diagram of ABE setup and operation.

Denial of Service (DoS). Then, the service sends the ABE module URL to the user for subsequent exchanges.

Assuming the service has encrypted its data with a DEK, the subsequent *service data access* phase follows. In this phase, when a user requests data, the service provides encrypted data. Having a valid User Key, obtained during the User setup phase, the user can request the DEK from the Key Repository using the ABE module URL received earlier. The ABE module decrypts the DEK, allowing the user to decrypt and access the requested data. Once successfully decrypted, users can proceed retrieving data from services. This integrated process ensures secure and controlled data access within the cloud environment.

Moreover, cost-effectiveness is achieved by predominantly employing symmetric encryption for data encryption, which is notably faster and less computationally demanding than asymmetric encryption.

2) UNIFIED ATTRIBUTE MATCHING: SIMPLIFYING CONFIGURATION AND MANAGEMENT OF CLOUD SECURITY

Referring back to Figure 2, we can delve into another concept for further clarification. Let us turn our attention to the ABE module, the key repository, and the attribute

values assigned to users. As previously mentioned, the DEKs encrypted with a specific policy find their place in the key repository. In accordance with our proposed approach, these policies are not composed of arbitrary attributes; instead, they consist of specific attributes values that match with the same set of values of the labels in the HPA.

Examining Figure 2, we observe that the policies within the key repository are formed exclusively with attributes *a* and *b*. Notably, these same attributes assigned to users also possess values within the same set. Remarkably, these are the same values employed earlier for the HPA labels.

This concept is a pivotal aspect of our proposal. Matching of label values with attribute values is essential, since this matching streamlines and expedites configuration and management. By introducing a unified management layer that comprehensively handles all attributes, including those from Kubernetes, and integrating it with the ABE module, we achieve transparency and ease of administration. This integration eliminates the need for distinct configurations and additional harmonization processes, resulting in a streamlined user experience.

3) EXPANDING POLICY EXPRESSIVENESS WITH ABE

The third key concept is related to the type of policy that can be expressed using ABE. In contrast to simple label-based policies, the ABE schema can be selected to articulate policies with greater expressive capacity. This enables the inclusion of logical expressions involving conjunctions, disjunctions, and negations. This expanded capability is outlined in the taxonomy of ABE schemes presented in [63].

The most sophisticated policies fall under the category of Non-Monotone Span Program (NMSP). For example, basic policies could be represented as follows:

- Policy = *a*
- Policy = *b*

Alternatively, more complex policies can encompass conjunctions, disjunctions, and/or negations:

- Policy = NOT *a*
- Policy = *a* OR *b*
- Policy = *a* AND *b*
- Policy = *a* OR NOT *b*
- Policy = NOT *a* AND *b*

and so forth, encompassing all feasible combinations. These expressive capabilities enable ABE to accommodate a diverse range of policies, enhancing its versatility and applicability.

4) THE OPERATIONAL ASPECTS OF THE ABE SCHEME FOR CLOUD SECURITY

Finally, we outline the ABE scheme selected for our experiments. We opted for the well-known scheme proposed by Lewko and Waters [46]. Our decision was guided by the scheme flexibility in adapting to the dynamic nature of cloud environments, allowing for the integration of multiple authorities in a decentralized system without requiring coordination between them. Nonetheless, it is

TABLE 1. Symbols.

Symbol	Description
λ	Security parameter
GP	Global parameters
SK	Authority's Private Key
PK	Authority's Public Key
GID	User's Global identity
i	i -th attribute
$K_{i,\text{GID}}$	User's Secret Key share of i attribute
M	Plaintext
(A, ρ)	Access matrix (Policy)
CT	Ciphertext

worth noting that this choice is not a constraint, and any other ABE scheme can be adapted in our system and procedures.

Lewko and Waters' Multi-Authority ABE scheme [46] eliminates the need of a central authority. Decentralizing key generation and issuance allows users to encrypt data by using boolean equations based on attributes received from different authorities. The key idea presented in [46] consists of using a hash function $H(\cdot)$ on the user's global identity, GID. This allows handling collusion resistance across multiple key generations by different authorities. In more detail, $H(\cdot)$ hashes each identity to a bilinear group element (for details see Appendix A and [25], [26]). $H(\text{GID})$ is used to link together keys issued to the same user by different authorities [31]. Although this way is more challenging than the single authority case, it has the great advantage of avoiding relying on a single entity.

We now consider the distinct functions constituting a multi-authority scheme, as demonstrated by Lewko and Waters [46]. The following description is not a detailed explanation of all the mathematical details of the LW-ABE scheme, which are available in the original document, but it serves to contextualize the functions of this scheme in the proposed application environment. Figure 4 shows a visual representation of the key functions of ABE within the context of our application although, for the sake of clarity, the figure includes a single authority. It is essential to note that multiple authorities can also be involved in providing users with private key. We clarify this concept below. Additionally, Table 1 reports the symbols utilized in what follows and their meaning. The Appendices contain the essential mathematical background to understand the basic operation of the functions used.

The fundamental function initiating the process is the global system setup. This step establishes the global parameters (GP), serving as the basis for the subsequent operations.

Step 0: Global Setup(λ) \rightarrow GP. The global setup algorithm receives the λ security parameter as input and outputs global parameters for the system. In the global setup, a bilinear group G of order $N = p_1 p_2 p_3$ is selected, where N is the product of 3 primes. The output GP consists of N and a

generator g_1 of G_{p_1} .¹ In addition, the hash function $H(\cdot)$ that maps GID to elements of G is published.

Once the global parameters are established, one or more authorities can be set up. Each of them is responsible for an attribute or a set of attributes.

Step 1: Authority Setup(GP) \rightarrow SK, PK. Each authority runs the authority setup algorithm with GP as input to produce its own secret key (SK) and public key (PK) pair. For each attribute i belonging to an authority, it selects two random numbers $\alpha_i, y_i \in \mathbb{Z}_N$ and publishes $\text{PK} = \{e(g_1, g_1)^{\alpha_i}, g_1^{y_i} \forall i\}$, keeping secret $\text{SK} = \{\alpha_i, y_i \forall i\}$. The authorities independently generate portions of a user private key in a multi-authority system. Each user holds attributes from distinct authorities, and the absence of the need for coordination between them is a noteworthy characteristic.

The user global identity (GID) ties together the portions of the secret key generated by different authorities [31]. This is related to two crucial functions: (1) ensures that a user must use all the pieces of the private key together for successful decryption, (2) prevents users from exchanging key portions, thus mitigating collusion attacks [63].

Step 2: KeyGen(GID, GP, i , SK) \rightarrow $K_{i,\text{GID}}$. The key generation algorithm receives as inputs an identity GID, the global parameters, an attribute i belonging to some authority, and the secret key SK for this authority. Then, it produces a key $K_{i,\text{GID}} = g_1^{\alpha_i} H(\text{GID})^{y_i}$ for this attribute-identity pair.

A data owner can encrypt a message or plaintext M by using a policy expressed as a (A, ρ) -pair. A is an $n \times l$ matrix, defined as the share-generating matrix for the linear secret-sharing scheme, which is the mathematical theory underlying the overall machinery.² ρ denotes a mapping function associating each row of A with a specific attribute. In our application, which uses ABE in hybrid mode, the plaintext M assumes the role of the DEK, and the data owner, represented by the service, can encrypt content using a DEK. This encrypted content is then made exclusively available to specific users who meet the criteria outlined in the selected policy.

Step 3: Encrypt(M , (A, ρ) , GP, {PK}) \rightarrow CT. The encryption algorithm takes in a message M , an access matrix (A, ρ) , the set of public keys for the relevant authorities, and the global parameters GP. It outputs a ciphertext CT. In more detail, the encrypter selects a random $s \in \mathbb{Z}_N$ and a random vector $v \in \mathbb{Z}_N$ with s as its first entry. We let $\lambda_x = A_x \cdot v$, where A_x is row x of A . It also selects a random vector $w \in \mathbb{Z}_N$ with 0 as its first entry. We let $\omega_x = A_x \cdot w$. For each row A_x of A , it selects a random $r_x \in \mathbb{Z}_N$. The ciphertext is computed as:

- $C_0 = M e(g_1, g_1)^s,$
- $C_{1,x} = e(g_1, g_1)^{\lambda_x} e(g_1, g_1)^{\alpha_{\rho(x)} r_x} \quad \forall x,$
- $C_{2,x} = g_1^{r_x} \quad \forall x,$
- $C_{3,x} = g_1^{\beta_{\rho(x)} r_x} g_1^{\omega_x} \quad \forall x.$

1. The interested reader can find details and references about bilinear groups in Appendix A.

2. The interested reader can find the basic concepts on the underlying theory of linear secret-sharing schemes in the Appendix B.

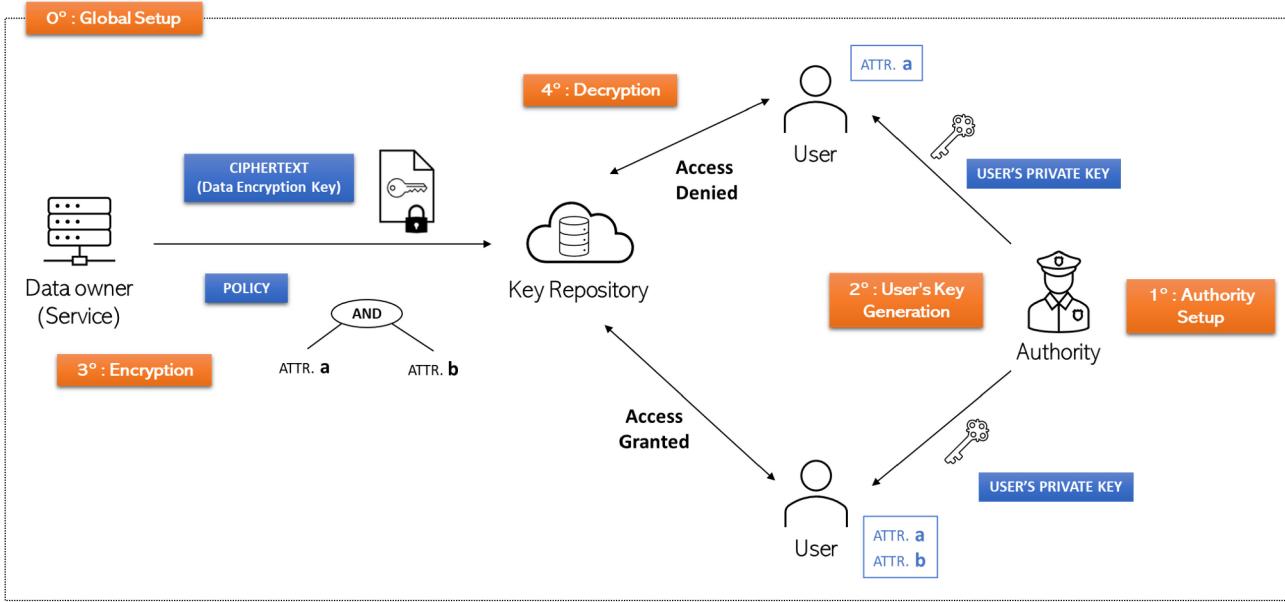


FIGURE 4. ABE scheme functions.

The successful decryption of a ciphertext depends on whether all the key shares of a user adhere to the policy linked to the ciphertext. In our application, DEKs are stored in the Key Repository of the ABE cloud service. As depicted in Figure 4, access to a particular ciphertext is granted exclusively to users who meet the specified policy criteria.

Step 4: Decrypt(CT, GP, {K_{i,GID}}) → M. The decryption algorithm takes in the global parameters GP, the ciphertext CT, and a collection of keys {K_{i,GID}} corresponding to pairs (attribute, identity), all with the same identity GID. It outputs the message M when the collection of attributes i satisfies the access matrix A corresponding to the ciphertext. Otherwise, decryption fails. From a mathematical viewpoint, the decryptor first obtains H(GID). If the decryptor has the secret keys {K_{ρ(x),GID}} for a subset of rows A_x of A such that f = (1, 0, ..., 0) is in the span of these rows, then it proceeds according to the following steps. For each such x, the decryptor computes:

$$\begin{aligned} C_{1,x} \cdot e(H(\text{GID}), C_{3,x}) / e(K_{\rho(x),\text{GID}}, C_{2,x}) &= \\ &= e(g_1, g_1)^{\lambda_x} e(H(\text{GID}), g_1)^{\omega_x}. \end{aligned}$$

Then, it selects constants c_x ∈ ℤ_N | ∑_x c_xA_x = f and computes:

$$\prod_x (e(g_1, g_1)^{\lambda_x} e(H(\text{GID}), g_1)^{\omega_x})^{c_x} = e(g_1, g_1)^s.$$

In fact, λ_x = A_x · v and ω_x = A_x · w, where v · f = s and w · f = 0. Consequently, the message can then be obtained as M = C₀/e(g₁, g₁)^s.

5) SECURITY ASPECTS OF THE ABE SCHEME: MODEL DEFINITION AND FORMAL SECURITY ANALYSIS

The schema preserves security and resilience through a game played between a challenger and an attacker. The assumption

here is that adversaries can only corrupt authorities statically, but key queries can be made adaptively. This slightly deviates from the standard static model, allowing an adversary to independently choose the public keys of corrupted authorities instead of having them initially generated by the challenger.

For the sake of clarity, consider the steps of the challenge, as illustrated in Figure 5. In this context, the set of authorities is represented as S, and the universe of attributes as U, with each attribute assigned to a specific authority. The security game unfolds as follows:

- 1) **Global Setup:** The global setup algorithm is initiated. The attacker designates a set S' of corrupt authorities out of the total set S. For non-corrupt authorities in S - S', the challenger generates public-private key pairs by using the authority setup algorithm. The public keys are provided to the attacker.
- 2) **Key Query Phase 1:** The attacker makes key queries, submitting pairs (i, GID) to the challenger, where i is an attribute from a non-corrupt authority, and GID is an identity. The challenger responds by providing the corresponding key K_{i,GID}.
- 3) **Challenge Phase:** The attacker specifies two messages, M₀, M₁, and an access matrix (A, ρ). The matrix must satisfy constraints to ensure that the attacker cannot construct a set of keys allowing decryption in conjunction with keys obtained from corrupt authorities. The challenger encrypts M_β under access matrix (A, ρ), where β is a random coin flip (β ∈ {0, 1}).
- 4) **Key Query Phase 2:** The attacker may submit additional key queries (i, GID), adhering to constraints on the challenge matrix (A, ρ).
- 5) **Guess:** The attacker submits a guess β' for the coin flip β. The attacker wins if β = β'. The attacker's advantage in this game is defined to be Pr[β = β'] - 1/2.

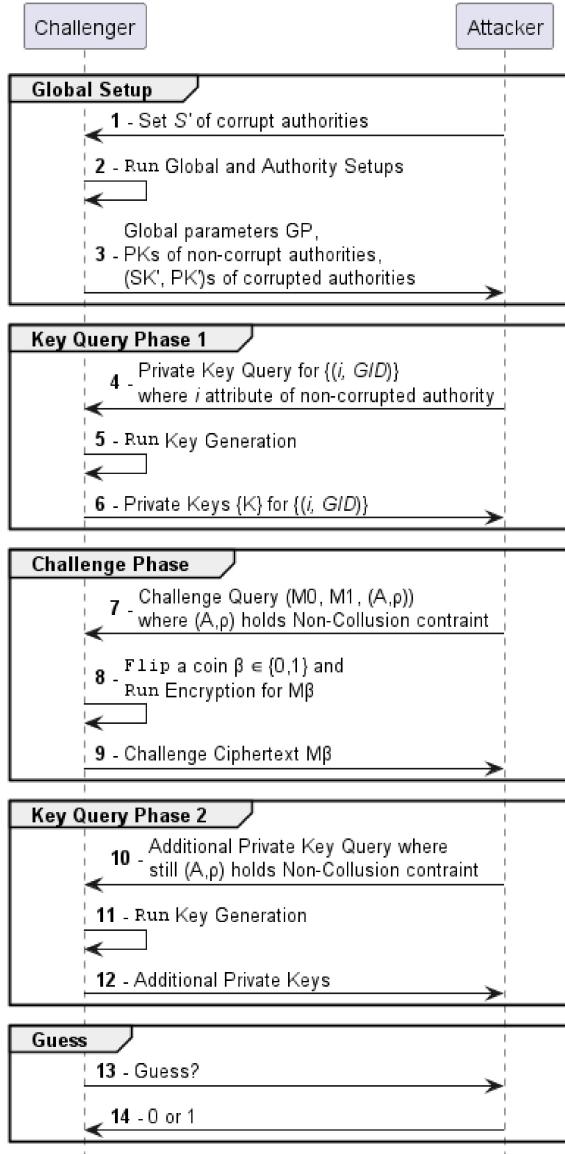


FIGURE 5. Security game sequence diagram.

A Multi-Authority CP-ABE system is secure (against static corruption of authorities) if all polynomial time attackers have at most a negligible advantage in this security game.

In order to prevent collusion, the main strategy consists of using a hash function on the user's global identity. The output of this function $H(GID)$ is the element of a bilinear group (see Appendix A) that ties keys together. Based on the theory of the linear secret-sharing schemes (introduced in Appendix B), the key idea is to handle the decryption process on the nodes x of the access tree, such that it is possible for a user to recover a group element of the form $e(g, g)^{\lambda_x} \cdot e(g, H(GID))^{w_x}$. The first element of this group element is a share λ_x of the secret s to recover. Thus, these shares need to be combined for deciphering the message. However, they are masked by a share w_x of the 0 element in the exponent of the second factor, that is the one with base

$e(g, H(GID))$. The algorithm in [46] allows simultaneously reconstructing the secret s and unblinding it. In fact, if a user identified by GID satisfies the access tree, he can obtain s by reconstructing it in the exponent by raising group elements to the proper base. At the same time, this operation will also reconstruct the shares of 0 and thus, in the end, the terms $e(g, H(GID))$ will disappear. In particular, the encryption algorithm ciphers the message M with $e(g_1, g_1)^s$, where g_1 is a generator of the subgroup G_{p_1} , and s is a randomly selected value in \mathbb{Z}_N , with $N = p_1 p_2 p_3$. The value s is then split into shares λ_x through the A matrix, and the value 0 is split into shares w_x . Hence, in order to recover the information message M , the decryption algorithm has to recover the blinding factor $e(g_1, g_1)^s$, so introducing terms in the form $e(g_1, H(GID))^{w_x}$. If the decryptor has a satisfying set of key s with the same identity GID, these additional terms will cancel from the final result, since the w_x elements are shares of 0. Instead, if two users with identities GID and GID' attempt to collude and combine their keys, then there are two types of terms, those of the form $e(g_1, H(GID))^{w_x}$ and others of the form $e(g_1, H(GID'))^{w_x}$. These different types of terms do not cancel with each other, thereby preventing to recover $e(g_1, g_1)^s$ and thus to decipher the message M .

The formal proof of security of this model against collusion attacks, making use of the dual system encryption methodology, is quite lengthy, and it is given in [46, Secs. IV-A and 4.2, Appendix C]. Here, we provide an outline of the strategy adopted, inviting the interested reader to refer to the original document for all the details, as well as to [65]. By employing a form of the dual system encryption technique, the authors of [46] address the challenges specific to the multi-authority setting. In this framework, keys and ciphertexts can be either normal or semi-functional, with distinct decryption capabilities. The security proof involves a series of games employing a hybrid argument, where the challenge ciphertext is first transformed into a semi-functional form, followed by a gradual transformation of the keys to a semi-functional state. The key challenge lies in ensuring indistinguishability between these games, preventing the simulator from discerning the transformation of keys from normal to semi-functional during the testing process. To tackle this, the authors adopt an approach involving nominal semi-functionality, where both key and ciphertext have semi-functional components that cancel out upon decryption. Moreover, in the presence of multiple authorities lacking coordination, the authors introduce innovative solutions, such as the use of temporary "blinding factors" to conceal nominal semi-functionality. These factors, active for one key at a time, are strategically switched between two subgroups in the multi-authority case, preserving semi-functionality while averting information leakage about the subgroup in which the semi-functional components operate.

In conclusion, this section highlights our dual objective, consisting in both unified management system for orchestration services and a robust security solution addressing confidentiality and compliance concerns. To achieve the

TABLE 2. Comparison with other Kubernetes cloud bursting solutions.

Solution	Vendor	Supported public cloud providers	Approach	Security		
				Access control	Network security	Data privacy
KIP [39]	Elotl	EKS, GKE	Virtual Kubelet	RBAC	Dedicated interconnect or VPN	N/A
KubeSlice [57]	Avesha	EKS, GKE, AKS	Virtual slice across multiple clusters	RBAC	Secure VPN tunnels + zero trust	N/A
ACK One [67]	Alibaba Cloud	ACK (mandatory) + many others	Centralized management via registration to ACK portal	RBAC	TLS-based encryption	N/A
CloudBees Build Acceleration [10]	CloudBees	EKS, GKE, AKS	Cluster Manager installed anywhere + Electric Agents running in worker nodes	RBAC	TLS encryption between agents	N/A
Anthos [7]	Google	GKE, EKS	Configuration synchronization via GitOps	RBAC	TLS-based service mesh	N/A
Proposed solution	Open-source	Multiple	Label-based operations for both orchestration and security	ABAC + ABE	TLS-based encryption	ABE

latter, we advocate for employing CP-ABE in a hybrid mode, providing fine-grained access control, adaptable policies, and an optimized balance between security and efficiency. This approach, as depicted in Figure 3, enables authorized users to access encrypted data stored within service volumes, streamlining data retrieval.

We have explored the correlation between attribute values, label values, and encryption policies, highlighting the critical role of congruence for streamlined configuration and management. This unification enhances transparency and user experience by eliminating redundant configurations and harmonization processes.

Furthermore, the discussion on the expressive power of ABE policies, encompassing logical expressions, underlines the versatility of ABE schemes and their potential to accommodate a broad spectrum of access control scenarios.

Finally, we describe the ABE scheme used in our research, motivated by its features, and analyzed its security. Importantly, we emphasize again that the selection of this specific scheme is not a rigid constraint. This flexibility reflects the dynamic nature of the security solutions, allowing for adaptation to different requirements. For these reasons, we believe that these concepts can improve security, efficiency, and versatility in the considered framework.

C. COMPARISON WITH OTHER CLOUD BURSTING APPROACHES

In this section, our proposal is compared with available counterparts presented in Section II. A schematic comparison is reported in Table 2. This table summarizes the main features of cloud bursting solutions of interest for our purposes, i.e., support by public cloud vendors, strategic approach, and security features. Whilst different solutions follow different approaches, all of them seem suitable to offload computing workload in public clouds. However, for what concerns security, it is handled separately from

orchestration functions. In fact, whereas communications security is typically handled via TLS-based encryption, access control is enforced only via RBAC, which suffers from known limitations. Finally, none of the considered solutions natively supports strong data encryption, which is a significant shortcoming. Although it is possible to add encryption modules for handling data security as additional services in virtual clusters running on public resources, they have to be managed manually by developers. This increases the complexity of service management when cloud bursting is necessary. Instead, by jointly carrying out orchestration, including bursting, and data protection through labeling, which is supported natively by K8s, this feature comes for free, without any additional burden for developers.

V. VALIDATION

In this section, we present the findings from our proof of concept (POC) for the proposed attribute-based management model for secure K8s cloud bursting. The main research question addressed by this POC is: Can attributes be used to establish a centralized system capable of simultaneously providing data confidentiality and facilitating efficient cloud bursting?

We evaluated the feasibility and effectiveness of this attribute-based management model through a specific proof of concept. To showcase the achievable performance, we produced a POC available on our GitHub repository [54]. This POC offers a straightforward scenario for assessing our proposal.

A. PROOF OF CONCEPT SET UP

The experimental setup included two interconnected clouds. It emulates a system either where a service is hosted on a local cloud only or where a local cloud connects to a public cloud through a dedicated virtual connection service, such as Azure ExpressRoute. The ABE service, crucial to our

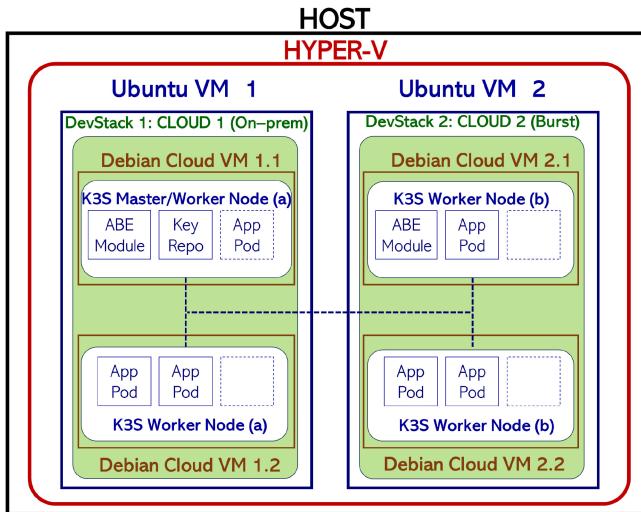


FIGURE 6. Test-bed components.

approach, is hosted within this cluster. It serves to encrypt, decrypt, and store data of string type, illustrating the core functions of our model.

The POC environment was hosted on a personal computer featuring an Intel i7-8700 CPU, 32GB RAM, and a 1TB SSD. The test environment is illustrated in Figure 6. It encompassed two distinct Ubuntu 22.04 LTS virtual machines (VMs), each one running a DevStack instance, thus effectively creating two separate cloud deployments (On-prem and Burst cloud). Each VM was allocated 6 virtual cores (vCPU), 12GB RAM, and 64GB SSD, ensuring a balanced testing environment. In turn, each DevStack cloud instance ran two Debian Cloud VMs, each one equipped with 2 vCPUs, 4GB RAM, and 20GB SSD. Each Debian VM ran a Kubernetes instance using K3s multi-cloud, one acting as both cluster master and worker node, and the other three as worker nodes. Thus, all Kubernetes instances were able to run application pods. Figure 6 shows the label (in brackets) of all Kubernetes nodes. In more detail, in order to configure the *on-premise/burst* attribute on the Kubernetes nodes, we labeled each node using kubectl by the label *a* for the on-premise nodes inside the first cloud and by the label *b* for the burst nodes running inside the second cloud, respectively. The commands used for labeling nodes as *on-premise* and *burst* are as follows:

```
kubectl label nodes <nn> node-type=a
where nn="master1a worker1a"
```

```
kubectl label nodes <nn> node-type=b
where nn="worker2b worker2b1"
```

These labels are then used in the *affinity* section of the Kubernetes Deployment manifest to automatically guide scheduling of pods.

We produced some YAML manifest files to configure the Kubernetes orchestration in the proposed scenario.

```
apiVersion: v1
kind: Service
metadata:
  name: kaboom
  labels:
    app: kaboom
spec:
  selector:
    app: kaboom
  ports:
    - name: http
      port: 8000
      targetPort: 8000
      nodePort: 30800
  type: NodePort
```

FIGURE 7. Yaml of Kubernetes Service.

The first manifest file, shown in Figure 7, defines a Kubernetes Service of type NodePort. The Service is named *kaboom* and is labeled with *app: kaboom*. Its purpose is to expose the application to the external. In particular, this Service exposes port 8000 and maps it to the node port 30800 and targets Pods labeled *app: kaboom*. When a user's request accesses the node on port 30800, the Service forwards the traffic to one of the pods listening on port 8000.

The second manifest file, shown in Figure 8, defines a Kubernetes Deployment. It creates a pod labeled with *app: kaboom*, runs the container *mrcolorrain/kaboom:latest*, and exposes port 8000 on the container.

The Readiness Probe is used to check the health of the application. The resources section sets resource requests for the pod. The affinity section specifies that the pod is expected to be scheduled on the nodes labeled with *node-type: a*, i.e., on-premise nodes. The strategy is set to Rolling Update, which means that the old pods are gradually replaced with new pods. This strategy ensures that there is no downtime during the update process and that the application remains available to users.

The last manifest file, shown in Figure 9, defines an autoscaler for the *kaboom* deployment. The autoscaler targets the *kaboom* deployment and scales between a minimum of 1 and a maximum of 40 replicas. It scales based on CPU utilization, targeting an average of 50%. The scale-down stabilization window is set to 90 seconds. Note that the horizontal pod autoscaler targets a maximum of 40 replicas for the *kaboom* deployment. This is due to an important aspect of resource utilization: considering that each *kaboom* pod consumes 0.128 CPU, a maximum of 40 replicas would require up to 5.12 CPUs in total. Given that the two VMs hosting the on-prem worker nodes are configured with only 2 vCPUs each, and one of them also runs the Kubernetes master, any deployment beyond 32 replicas would certainly trigger cloud bursting to meet the resource requirements.

By using the manifest files shown, we created a test computing environment where the *kaboom* application is exposed to the external world, automatically scales based on

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: kaboom
  labels:
    app: kaboom
spec:
  replicas: 1
  selector:
    matchLabels:
      app: kaboom
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
  revisionHistoryLimit: 1
  template:
    metadata:
      labels:
        app: kaboom
    spec:
      containers:
        - name: kaboom
          image: mrcolorrain/kaboom:latest
          ports:
            - containerPort: 8000
          readinessProbe:
            httpGet:
              path: /
              port: 8000
            initialDelaySeconds: 5
            periodSeconds: 5
          resources:
            requests:
              cpu: 128m
              memory: 128Mi
          affinity:
            nodeAffinity:
              preferredDuringScheduling:
                - ignoredDuringExecution:
                  weight: 100
                  preference:
                    matchExpressions:
                      - key: node-type
                        operator: In
                        values:
                          - a

```

FIGURE 8. Yaml of Kubernetes Deployment.

CPU utilization, and is preferably scheduled on nodes *on-premise* labeled as a. This is actually the cloud architecture characterizing the POC scenario illustrated above, enhanced by the integration of the Kubernetes management functions with ABE for secure cloud bursting.

This implementation can be regarded as the starting point for additional research in a more complex situations, such as strengthening security of multi-tenant Kubernetes environments loaded with multiple applications. In these challenging situations, it is critical to achieve a fine control of application deployment which is compliant with policies and regulations. Our approach, based on node affinity and pod anti-affinity rules combined with ABE policies, can help

```

apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: kaboom
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: kaboom
  minReplicas: 1
  maxReplicas: 40
  metrics:
    - type: Resource
      resource:
        name: cpu
      target:
        type: Utilization
        averageUtilization: 50
  behavior:
    scaleDown:
      stabilizationWindowSeconds: 90

```

FIGURE 9. Yaml of Autoscaler.

to achieve this control. In fact, it allows protecting access to data and isolating target applications from both specific nodes and other services. For instance, if certain applications should be prevented from being scheduled on cloud nodes or near another specific application, it is possible to use the manifest files shown as blueprint to include a node affinity rule and a pod anti-affinity rule for comprehensive scheduling control, as shown in Figure 10.

In this example, the node affinity rule confines the kaboom application to on-premises nodes, while the pod anti-affinity rule guarantees that kaboom pods do not share an environment with anti-kaboom pods. This dual-layered approach offers robust isolation and fine-grained control over Kubernetes deployments.

B. ANALYSIS OF ABE FUNCTIONS

The ABE critical components were developed by using the Rust programming language. Specifically, the ABE module was crafted by using the Rocket framework version 5.0 [21]. Cryptographic primitives integral to the AW11 scheme implementation are based on the Rabe library version 0.2.7 [4]. We emulated user interactions through both the Web interface and the API calls issued through a terminal. The ABE components run as Kubernetes pods, as shown in Figure 6. The ABE module can be accessed either through the Web interface or the terminal, and it presents a menu comprising the following options:

- 1) Show encrypted storage
- 2) Decrypt storage
- 3) Update storage

1) OPTION 1: SHOW ENCRYPTED STORAGE

The first option, accessible to anyone, allows accessing and viewing the encrypted storage. As shown in Figure 11,

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: kaboom
spec:
  replicas: 3
  selector:
    matchLabels:
      app: kaboom
  template:
    metadata:
      labels:
        app: kaboom
    spec:
      affinity:
        nodeAffinity:
          requiredDuringScheduling:
            - IgnoredDuringExecution:
            - nodeSelectorTerms:
              - matchExpressions:
                - key: node-type
                  operator: NotIn
                  values:
                    - b
        podAntiAffinity:
          requiredDuringScheduling:
            - IgnoredDuringExecution:
            - labelSelector:
              matchExpressions:
                - key: app
                  operator: In
                  values:
                    - anti-kaboom
        topologyKey:
          - "kubernetes.io/hostname"

```

FIGURE 10. Example showing the use of affinity and anti-affinity to change cloud bursting and apps deployment behavior.



Encrypted storage

1. Encrypted Entry -- Policy "a"
2. Encrypted Entry -- Policy "b"
3. Encrypted Entry -- Policy "a" and "b"
4. Encrypted Entry -- Policy "a" or "b"

[Go back to index](#)

FIGURE 11. Example of encrypted storage (screenshot of the Web interface).

which presents a screenshot of the Web interface, this list reveals the number of encrypted entries within the ABE volume. Additionally, it provides an instructive glimpse into the associated policies for each entry.

In Figure 11, the example showcases the dynamic interplay of policies. The first two entries are linked to simple policies encompassing a single system attribute. In contrast, the subsequent entries are tied to slightly more intricate policies that incorporate not only attributes but also logical operators.

```

{
  "_gid": "bob",
  "_attr": [
    [
      "b",
      {
        "x": [
          5712868850552490035,
          3340114564258679341,
          7948731980494642711,
          194928839779717989
        ],
        "y": [
          11833180424976684086,
          17503833986257901125,
          16512980884993783250,
          2315106061390818652
        ],
        "z": [
          13701724635261595198,
          7942669771980227583,
          14101786126588513581,
          3102591607448728316
        ]
      }
    ]
  }
}

```

FIGURE 12. Example of user key.

This visual demonstration underscores the ABE potential to extend policy complexity, enhancing the expressive capabilities of the protection system.

2) OPTION 2: DECRYPT STORAGE

The second option, labeled as *Decrypt Storage*, requires the introduction of a user key before the decryption process can proceed.

In Figure 12, we can observe an example of a user key utilized in our tests. This private key is formatted in JSON, containing references to the key owner, in this case, bob, and the attributes associated with it. In this instance, it holds a single attribute b, defined with specific coordinates. This key functions as a passkey, instructing the system on what content to decrypt and what to keep concealed.

Once the user key has been entered, the screen depicted in Figure 13 is displayed. In accordance with the policies outlined in Figure 11, the user possession of the b attribute grants him access solely to the second and fourth entries. Consequently, the user can retrieve two DEKs (see also Figure 3), providing him with the means to securely interact with data services within the cluster.

3) OPTION 3: UPDATE STORAGE

The third option concludes the menu by enabling the addition of a new entry to the encrypted storage. This procedure involves supplying two crucial pieces of information: the plaintext content and the policy used for encryption. Once

← → ⌂ ① 127.0.0.1:8000/decrypt-storage

Decrypted storage

Welcome bob!

1. Plaintext =
2. Plaintext = second_dek
3. Plaintext =
4. Plaintext = fourth_dek

[Go back to index](#)

FIGURE 13. Example of decrypted storage (screenshot of the Web interface).

TABLE 3. Performance of ABE-related functions.

		#2 attributes	#3 attributes
Authority Setup	Average	0.4906	0.7307
	SD	0.0096	0.0122
	Min	0.4748	0.7134
	Max	0.5329	0.7782
User Private Key Generation	Average	0.0594	0.0885
	SD	0.0016	0.0019
	Min	0.057	0.0855
	Max	0.0656	0.0974
Encryption	Average	1.2276	1.6089
	SD	0.1225	0.118
	Min	1.1779	1.5565
	Max	2.1592	2.4436
Decryption	Average	0.55	0.554
	SD	0.0073	0.0133
	Min	0.5431	0.5443
	Max	0.5838	0.6454

these details are provided, the system undertakes the encryption process automatically, followed by the seamless storage of the encrypted entry within the ABE volume.

C. ABE-FUNCTIONS PERFORMANCE

This section presents the performance results obtained by using the ABE scheme. Experiments have been carried out on one of the Ubuntu VMs used in the POC (see Section V-A). We used the Rabe library, which is a Rust-based library implementing ABE. We selected the AW11 scheme [4], [46], as mentioned above. The encryption process converts 256-bit plaintext segments, managed as a symmetric key, according to Section IV-B. We collected the performance metrics related to the four main functions of an ABE scheme: authority setup, user private key generation, encryption, and decryption. All of them are sketched in Figure 4.

In performance evaluation, We varied the number of attributes used in the policy and the user's secret key. In particular, we used 2 and 3 attributes. The experiment was repeated 100 times for each attribute count. Numerical

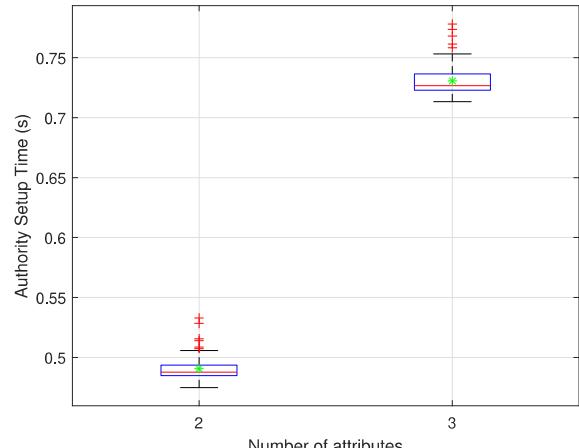


FIGURE 14. Box plot of the authority setup time as a function of the number of attributes. The green asterisk indicates the average value.

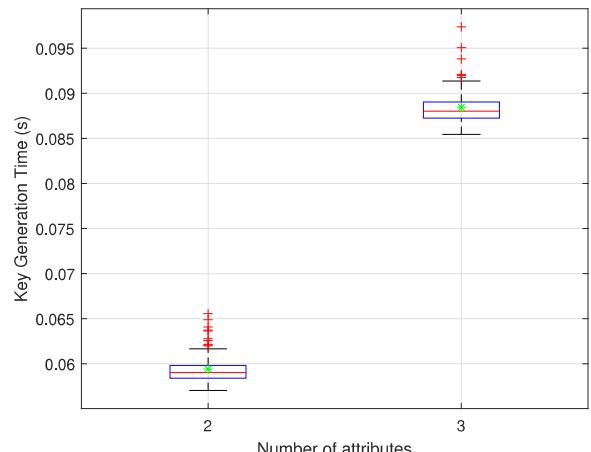


FIGURE 15. Box plot of the user private key generation time as a function of the number of attributes. The green asterisk indicates the average value.

results are reported in Table 3, in terms of average latency, its standard deviation, minimum and maximum values. All results are expressed in seconds. In addition, for each of these functions, we included a box plot diagram, in order to better appreciate their statistical properties. These diagrams also show the average value, marked by a green asterisk. All latency values are grouped in a quite compact range (interquartile range), showing only a few significant outliers for encryption (Figure 16) and decryption (Figure 17) functions. Differently, authority setup times (Figure 14) and key generation times (Figure 15) show more grouped values.

Observing Table 3 and box plots diagrams, it can be observed a generalized increase of the processing time when switching from 2 to 3 attributes, which indicates an increase in complexity. Notably, the increase with the number of attributes in the authority setup time (+48.9% on average), user private key generation (+48.9% on average), and encryption (+31% on average) functions is more pronounced compared to the decryption function, where it is less steep (+0.66% only on average).

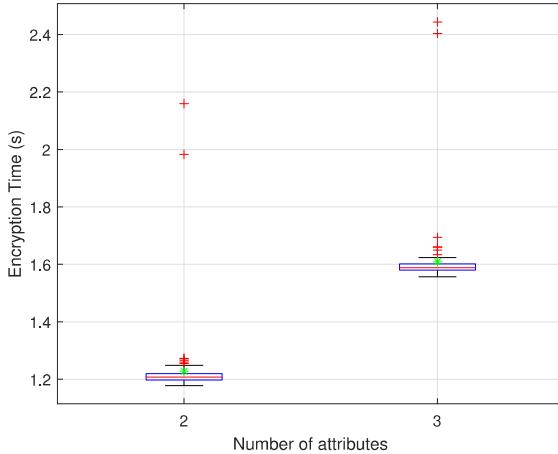


FIGURE 16. Box plot of the encryption time as a function of the number of attributes. The green asterisk indicates the average value.

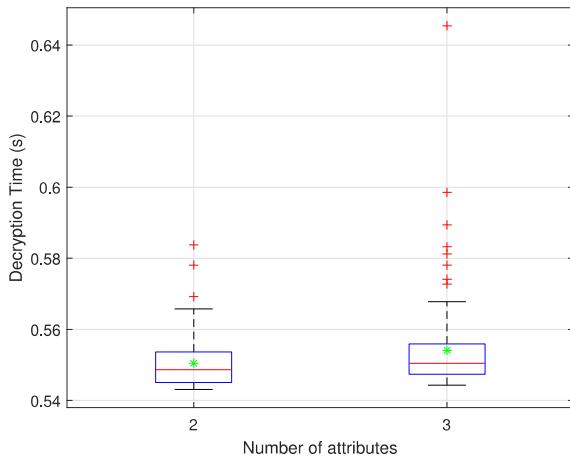


FIGURE 17. Box plot of the decryption time as a function of the number of attributes. The green asterisk indicates the average value.

In order to evaluate the average payoff due to the presence of ABE, we evaluated the added latency per user session, depicted in Figure 3 (service data access). In that figure it appears that, in general, only the decryption function is invoked for completing a new session by the user, whereas the other functions are used only occasionally (setup and encryption of the DEK). Thus, in order to evaluate the average added latency, we measured the added contribution due to the decryption function as well as the protocol used for interacting with the ABE module, which is accessible through a REST API. Figure 18 shows these results, identifying the two above delay contribution: DEK decryption and REST interaction. The overall added latency due to our solution is slightly less than 0.6 seconds, with a negligible 95% confidence intervals. This latency is essentially due to the DEK decryption, and does not depend in a significant way from the number of attributes used to implement the policy. It seems to be a reasonable price to obtain an improved security.

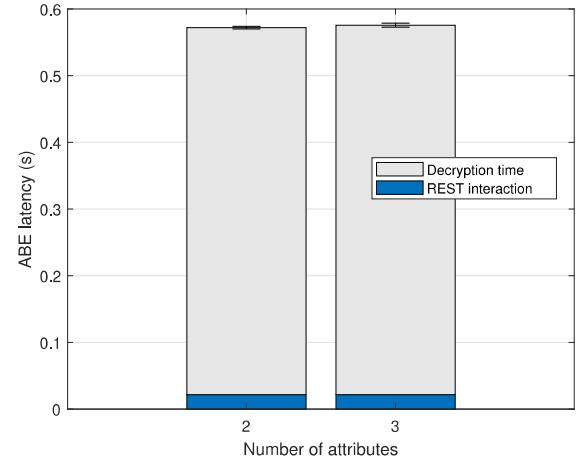


FIGURE 18. Latency associated with ABE function with the relevant 95% confidence intervals. Relative weights of decryption function and REST calls contributions are highlighted.

Although we explored some key use-cases, an exhaustive exploration of attribute usage is beyond the scope of this paper. Mosteiro-Sánchez et al. [51] have already conducted comprehensive experiments in this regard. Our emphasis lies in the broader applicability of the insights obtained. While we specifically implemented the AW11 scheme, [51] serves not only to deepen the performance understanding of AW11, but is a rather valuable reference also for assessing the performance of different ABE libraries.

D. DISCUSSION

We validated through a POC the system's ability to adjust cloud resources based on computational requirements using the cluster setup described in previous blueprints, which is critical in cloud bursting scenarios.

The ABE module revealed its capabilities through the menu options. *Show Encrypted Storage* illustrated the potential of the module to manage intricate access policies, *Decrypt Storage* emphasized secure data access through user keys, and *Update Storage* showcased the seamless integration of encryption processes. The ABE user-centric module design promotes secure data interaction. It aligns with modern data security paradigms, maintaining data privacy even during burst operations. The visual policy representation empowers administrators with effective access control.

In summary, the usage of labels common to Kubernetes configuration (nodes and app pods) and ABE policies allows *jointly* configure cloud bursting and security features, leveraging the potential of the attribute-based approach.

In terms of implications and future prospects, our experiments validate the viability of our attribute-based management model for secure cloud bursting. The convergence of Kubernetes and ABE presents a promising solution for organizations seeking secure and scalable cloud deployments. Future work can address scalability challenges and explore advanced security layers.

VI. CONCLUSION

In this paper, we introduce a robust orchestration scheme tailored for secure cloud bursting. This scheme addresses the complexities, cost challenges, and stringent data protection compliance requirements by harnessing the combined capabilities of K8s and ABE.

By incorporating ABE, our approach achieves granular encryption and provides cloud resources with suitable confidentiality. At the same time, cloud bursting empowers the extension of computational tasks beyond the scope of a local primary cloud environment. The synergy of these two technologies establishes a cohesive management framework, guaranteeing secure access to bursting services and streamlined deployment of excess workloads to the cloud, all facilitated by Kubernetes.

Significantly, our contributions encompass the design blueprint for an attribute-based cloud bursting authorization system within Kubernetes. Through the application of ABE, we enhance data security to adhere to regulatory mandates and mitigate data-related apprehensions. By bridging this critical gap, our approach stands as a holistic solution that aligns with both security and privacy regulations, meeting the contemporary requisites of cloud-driven systems.

Practical feasibility of our proposal is demonstrated by the implementation of a proof of concept, that shows its potential to be seamlessly integrated into real-world scenarios.

Future research will consider the integration of the proposed solution with artificial intelligence algorithms to proactively infer the need of resorting to bursting operations. In fact, purely relying on a reactive approach the latency of the process can either lead to temporary violation of service level agreements - loose approach - or allocation of excessive resources - conservative approach.

APPENDIX

In these appendices, we report some useful definitions, some notions about bilinear groups as well as the very basics of linear secret-sharing schemes. These concepts are used in the paper of Lewko and Waters [46], and are instrumental for a better understanding of the mathematics about ABE used in this manuscript.

A. BILINEAR GROUPS

Bilinear groups of composite order are groups with an efficient bilinear map where the group order is a product of two large primes. Such groups are constructed from pairing friendly curves over a finite field [25], [26]. The following assumptions are formulated for a bilinear group G of order N , where $N = p_1 p_2 p_3$ is the product of 3 different primes. Now, let $e : G \times G \rightarrow G_T$ denote a bilinear map, that is a function with the following properties:

- 1) $\forall g, h \in G, a, b \in \mathbb{Z}_N, e(g^a, h^b) = e(g, h)^{ab}$ (bilinear),
- 2) $\exists g \in G | e(g, g)$ has order n in G_T (non-degenerate).

A group generator \mathcal{G} is an algorithm which takes a security parameter λ as input and produces a description of a bilinear

group G , i.e., $(p_1, p_2, p_3, G, G_T, e)$, where both G and G_T are cyclic groups of order N , and e is a bilinear map. It is assumed that the group operations in G and G_T and the map e can be computed in polynomial time with respect to λ , and that the group descriptions include generators of the respective cyclic groups, according to the above definition. Let G_{p_1} denote the subgroup of order p_1 in G . If $g_i \in G_{p_1}$ and $g_j \in G_{p_j}$ for $i \neq j$, then $e(g_i, g_j) = 1$. The construct $g_1 \xleftarrow{R} G_{p_1}$, indicates that g_1 is a random generator of G_{p_1} . Thus, it is not the identity element. The interested reader can find additional details on composite order bilinear group in [46, Appendix A] and references therein.

According to the class of General Subgroup Decision Assumptions described in [19], in a bilinear group G of order $N = p_1 p_2 \dots p_n$, a subgroup of order $\prod_{i \in S} p_i$ exists for each subset $S \subseteq \{1, \dots, n\}$. Let S_0, S_1 denote two distinct subsets, assuming that it is hard to distinguish a random element from the subgroup associated with S_0 from another of the subgroup associated with S_1 . This is true even though it is possible to have random elements from subgroups associated with several subsets Z_i , which satisfy one of the following two mutually exclusive conditions: $S_0 \cap Z_i = S_1 \cap Z_i = \emptyset$, or $S_0 \cap Z_i \neq \emptyset \neq S_1 \cap Z_i$.

In [46], the entire system is confined to the subgroup G_{p_1} in G except $H(\cdot)$, which instead maps identities onto random group elements in G . The subgroups G_{p_2} and G_{p_3} are used in the security proof, which makes use of the dual system encryption technique, where keys and ciphertexts can be either normal or semi-functional. In the approach proposed in [46] and used in this work, normal keys and ciphertexts are contained in the subgroup G_{p_1} , while semifunctional keys and ciphertexts use elements of G_{p_2} and G_{p_3} . Thus, G_{p_2} and G_{p_3} form the semi-functional space orthogonal to G_{p_1} .

B. LINEAR SECRET-SHARING SCHEME

Given a set of parties $\mathcal{P} = \{P_1, \dots, P_n\}$, a collection is defined as $\mathbb{A} \subseteq 2^{\{P_1, \dots, P_n\}}$. An access structure is a collection \mathbb{A} of non-empty subset, i.e., $\mathbb{A} \subseteq 2^{\{P_1, \dots, P_n\}} - \emptyset$. The sets in \mathbb{A} are called authorized sets, those outside \mathbb{A} are called unauthorized sets.

A Linear Secret-Sharing Scheme (LSSS) [52] can be defined as follows. Given a set of parties \mathcal{P} , a secret sharing scheme Λ is defined *linear* over \mathcal{P} (over \mathbb{Z}_p) if

- The shares of each party form a vector in \mathbb{Z}_p .
- An $l \times n$ matrix A exists with the role of share-generating matrix for Λ . Given a function $\rho(x) : x \in \{1, \dots, l\} \rightarrow \mathcal{P}$, each row x of A is labeled by it. If the secret to be shared is $s \in \mathbb{Z}_p$ and $r_2, \dots, r_n \in \mathbb{Z}_p$ are random values, then it is possible to obtain the column vector $v = (s, r_2, \dots, r_n)$ such that the product Av is the vector of l shares of the secret s according to Λ . The x -th share of Av , denoted as $(Av)_x = A_{x \cdot} v$, belongs to party $\rho(x)$.

From this definition, the linear reconstruction property follows. Under the assumption that Λ is an LSSS for \mathbb{A} , if S

is an authorized set, it is possible to define $I = \{x | \rho(x) \in S\}$, with $I \subseteq \{1, \dots, l\}$. Thus, the vector $f = (1, 0, \dots, 0)$ is in the set of rows of A indexed by I . Furthermore, a set of constants $\{\omega_x \in \mathbb{Z}_p\}_{x \in I}$ exists such that $s = \sum_{x \in I} \omega_x \lambda_x$ for any valid shares $\{\lambda_x\}$ of the secret s according to Λ . The above constants $\{\omega_x \in \mathbb{Z}_p\}$ can be found in polynomial time with respect to the size of A . This also means that $\lambda_x = A_x \cdot v$ where $v \cdot f = s$.

For the construction of the composite order group, given an LSSS matrix A over \mathbb{Z}_N , where $N = p_1 p_2 p_3$ is the product of 3 different primes, S is an authorized set if the rows of A that are labeled by the elements in S through the function ρ have the vector f as defined above in their span modulo N . The access policies can be described through monotonic boolean formulas and there are standard techniques to derive a LSSS matrix from any monotonic boolean formula. The boolean formula can be represented as an access tree, where interior nodes consists of AND and OR gates, whereas the leaves correspond to attributes. Clearly, the number of rows in a LSSS matrix is equal to the number of leaves in the access tree. Additional details can be found in [46, Appendix G].

ACKNOWLEDGMENT

The authors acknowledge Università degli Studi di Perugia and MUR for support within the projects VITALITY and RESTART.

REFERENCES

- [1] (Amazon Web Serv., Inc., Seattle, WA, USA). *Amazon Elastic Kubernetes Service (Amazon EKS)*. Accessed: Jun. 8, 2023. [Online]. Available: https://aws.amazon.com/eks/?nc1=h_ls
- [2] (google, Mountain View, CA, USA). *Google Kubernetes Engine*. Accessed: Jun. 8, 2023. [Online]. Available: <https://cloud.google.com/kubernetes-engine>
- [3] (Int. Bus. Mach. Corp., Armonk, NY, USA). *Ibm Kubernetes Service*. Accessed: Jun. 8, 2023. [Online]. Available: <https://www.ibm.com/cloud/kubernetes-service>
- [4] “Module aw11 rabe.” Accessed: Jan. 5, 2024. [Online]. Available: <https://docs.rs/rabe/latest/rabe/schemes/aw11/index.html>
- [5] (Oracle Comput. Softw. Co., Austin, TX, USA). *Oracle Cloud Native Services—Container Engine for Kubernetes*. Accessed: June 8, 2023. [Online]. Available: <https://www.oracle.com/cloud/cloud-native/container-engine-kubernetes/>
- [6] “Security best practices for Kubernetes deployment.” Mar. 2019. [Online]. Available: <https://kubernetes.io/blog/2016/08/security-best-practices-kubernetes-deployment/>
- [7] “Anthos.” Dec. 2023. [Online]. Available: <https://cloud.google.com/anthos>
- [8] (Palo Alto Networks, Inc., Santa Clara, CA, USA). *Cloud-Native Applications Protection Platform*. Dec. 2023. [Online]. Available: <https://www.paloaltonetworks.com/prisma/cloud/cloud-native-application-protection-platform>
- [9] “Configuration best practices: Using labels.” Jul. 2023. [Online]. Available: <https://kubernetes.io/docs/concepts/configuration/overview/#using-labels>
- [10] (CloudBees Softw. Co., San Jose, CA, USA). *Configuring Cloudbees Build Acceleration for Agent Cloud Bursting*. (Dec. 2023). [Online]. Available: <https://docs.cloudbees.com/docs/cloudbees-build-acceleration/latest/configuration-guide/config-accelerator-agents-for-cloud-burst>
- [11] “Kubernetes autoscaler.” github. Dec. 2023. [Online]. Available: <https://github.com/kubernetes/autoscaler>
- [12] (Microsoft Corp. Technol. Corp., Redmond, WA, USA). *Securing Kubernetes Workloads In Hybrid Settings With Aporeto*. (Dec. 2023). [Online]. Available: <https://cloudblogs.microsoft.com/opensource/2018/08/31/securing-kubernetes-workloads-hybrid-cloud-aporeto/>
- [13] (Amazon Web Serv., Inc., Seattle, WA, USA). *TLS-Enabled Kubernetes Clusters With ACM Private CA and Amazon EKS*. (Dec. 2023). [Online]. Available: <https://aws.amazon.com/it/blogs/security/tls-enabled-kubernetes-clusters-with-acm-private-ca-and-amazon-eks-2/>
- [14] “Virtual kubelet.” Dec. 2023. [Online]. Available: <https://virtual-kubelet.io/>
- [15] R. Ahuja and S. K. Mohanty, “A scalable attribute-based access control scheme with flexible delegation cum sharing of access privileges for cloud storage,” *IEEE Trans. Cloud Comput.*, vol. 8, no. 1, pp. 32–44, Mar. 2020.
- [16] S. Ameer, J. Benson, and R. Sandhu, “An attribute-based approach toward a secured smart-home IoT access control and a comparison with a role-based approach,” *Information*, vol. 13, no. 2, p. 60, 2022.
- [17] D. Balouek-Thomert, E. G. Renart, A. R. Zamani, A. Simonet, and M. Parashar, “Towards a computing continuum: Enabling edge-to-cloud integration for data-driven workflows,” *Int. J. High Perform. Comput. Appl.*, vol. 33, no. 6, pp. 1159–1174, 2019.
- [18] L. Baresi, D. F. Mendonça, M. Garriga, S. Guinea, and G. Quattrochi, “A unified model for the mobile-edge-cloud continuum,” *ACM Trans. Internet Technol.*, vol. 19, no. 2, pp. 1–21, Apr. 2019.
- [19] M. Bellare, B. Waters, and S. Yilek, “Identity-based encryption secure against selective opening attack,” *Cryptol. ePrint Arch.*, IACR, Bellevue, WA, USA, Rep. 2010/159, 2010.
- [20] P. Benedetti, M. Femminella, G. Reali, and K. Steenhaut, “Reinforcement learning applicability for resource-based auto-scaling in serverless edge applications,” in *Proc. IEEE Int. Conf. Pervasive Comput. Commun. Workshops Other Affil. Events (PerCom Workshops)*, 2022, pp. 674–679.
- [21] S. Benitez, “Meet rocket.” Accessed: Jan. 5, 2024. [Online]. Available: <https://rocket.rs/>
- [22] S. Bera, S. Prasad, Y. Sreenivasa Rao, A. K. Das, and Y. Park, “Designing attribute-based verifiable data storage and retrieval scheme in cloud computing environment,” *J. Inf. Secur. Appl.*, vol. 75, Jun. 2023, Art. no. 103482.
- [23] J. Bethencourt, A. Sahai, and B. Waters, “Ciphertext-policy attribute-based encryption,” in *Proc. IEEE Symp. Security Privacy (SP’07)*, 2007, pp. 321–334.
- [24] S. Böhm and G. Wirtz, “Cloud-edge orchestration for smart cities: A review of Kubernetes-based orchestration architectures,” *EAI Endorsed Trans. Smart Cities*, vol. 6, no. 18, p. e2, May 2022.
- [25] D. Boneh, “Bilinear groups of composite order,” in *Proc. 1st Int. Conf. Pair-Based Cryptogr.*, 2007, p. 1.
- [26] D. Boneh, E. Goh, and K. Nissim, “Evaluating 2-DNF formulas on ciphertexts,” in *Proc. Theory Cryptogr. Conf.*, 2005, pp. 325–341.
- [27] B. Burns, B. Grant, D. Oppenheimer, E. Brewer, and J. Wilkes, “Borg, Omega, and Kubernetes,” *Commun. ACM*, vol. 59, no. 5, pp. 50–57, 2016.
- [28] S. Carielli and L. Sustar (Forrester Res. Inc., Cambridge, MA, USA). *Best Practices: Kubernetes Security—Coordinate Across Identity, Network, and Container Security to Protect K8S*. (Jun. 2023). [Online]. Available: <https://reprints2..com/#/assets/2/1941/RES179415/report>
- [29] E. Carmona-Cejudo, F. Iadanza, and M. S. Siddiqui, “Optimal offloading of Kubernetes pods in three-tier networks,” in *Proc. IEEE Wireless Commun. Netw. Conf. (WCNC)*, 2022, pp. 280–285.
- [30] C. Carrión, “Kubernetes scheduling: Taxonomy, ongoing issues and challenges,” *ACM Comput. Surv.*, vol. 55, no. 7, pp. 1–37, Dec. 2022.
- [31] M. Chase, “Multi-authority attribute based encryption,” in *Proc. 4th Theory Cryptogr. Conf.*, 2007, pp. 515–534.
- [32] N. Chen, J. Li, Y. Zhang, and Y. Guo, “Efficient CP-ABE scheme with shared decryption in cloud storage,” *IEEE Trans. Comput.*, vol. 71, no. 1, pp. 175–184, Jan. 2022.
- [33] S. Chen, J. Li, Y. Zhang, and J. Han, “Efficient revocable attribute-based encryption with verifiable data integrity,” *IEEE Internet Things J.*, early access, Oct. 23, 2023, doi: [10.1109/IJOT.2023.3325996](https://doi.org/10.1109/IJOT.2023.3325996).
- [34] A. Chiquito, U. Bodin, and O. Schelén, “Attribute-based approaches for secure data sharing in industrial contexts,” *IEEE Access*, vol. 11, pp. 10180–10195, 2023.
- [35] *Regulation (EU) 2016/679 of the European Parliament and of the Council*, European Parliament, Brussel, Belgium, 2016

- [36] F. Faticanti et al., "Distributed cloud intelligence: Implementing an ETSI manu-compliant predictive cloud bursting solution using openstack and Kubernetes," in *Proc. 17th Int. Conf. Econ. Grids, Clouds, Syst., Serv.*, 2020, pp. 80–85.
- [37] B. Ghosh. "Hybrid Kubernetes model: Bridging clouds and on-premises." Medium. Sep. 2017. [Online]. Available: <https://medium.com/@bijit211987/hybrid-kubernetes-model-bridging-clouds-and-on-premises-4206cd430d50>
- [38] B. Ghosh. "The evolution of Kubernetes clusters in multicloud and hybrid cloud." Medium. Apr. 2023. [Online]. Available: <https://addozhang.medium.com/the-evolution-of-kubernetes-clusters-in-multi-cloud-and-hybrid-cloud-b243aa74eab>
- [39] C. Gocul. "Cloud bursting with virtual kubelet and KIP (kloud instance provider)." LinkedIn. May 2020. [Online]. Available: <https://www.linkedin.com/pulse/cloud-bursting-virtual-kubelet-kip-kloud-instance-provider-chandra/>
- [40] T. Guo, U. Sharma, T. Wood, S. Sahu, and P. J. Shenoy, "Seagull: Intelligent cloud bursting for enterprise applications," in *Proc. USENIX Annu. Tech. Conf.*, 2012, pp. 361–366.
- [41] Y. Guo, Z. Lu, H. Ge, and J. Li, "Revocable blockchain-aided attribute-based encryption with escrow-free in cloud storage," *IEEE Trans. Comput.*, vol. 72, no. 7, pp. 1901–1912, Jul. 2023.
- [42] M. A. Ibrahim, G. A. Ebrahim, and H. K. Mohamed, "A modern cloud bursting framework," in *Proc. 12th Int. Conf. Comput. Eng. Syst. (ICCES)*, 2017, pp. 148–153.
- [43] B. Kar, W. Yahya, Y. Lin, and A. Ali, "Offloading using traditional optimization and machine learning in federated cloud-edge-fog systems: A survey," *IEEE Commun. Surveys Tuts.*, vol. 25, no. 2, pp. 1199–1226, 2nd Quart., 2023.
- [44] C. Lan, L. Liu, C. Wang, and H. Li, "An efficient and revocable attribute-based data sharing scheme with rich expression and escrow freedom," *Inf. Sci.*, vol. 624, pp. 435–450, May 2023.
- [45] J. Lee, S. Oh, and J. W. Jang, "A work in progress: Context based encryption scheme for Internet of Things," *Procedia Comput. Sci.*, vol. 56, pp. 271–275, 2015. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S1877050915016890>
- [46] A. Lewko and B. Waters, "Decentralizing attribute-based encryption," in *Proc. 30th Annu. Int. Conf. Theory Appl. Cryptogr. Techn.*, 2011, pp. 568–588.
- [47] J. Li, W. Yao, Y. Zhang, H. Qian, and J. Han, "Flexible and fine-grained attribute-based data storage in cloud computing," *IEEE Trans. Services Comput.*, vol. 10, no. 5, pp. 785–796, Sep.–Oct. 2017.
- [48] J. Li et al., "Multiauthority attribute-based encryption for assuring data deletion," *IEEE Syst. J.*, vol. 17, no. 2, pp. 2029–2038, Jun. 2023.
- [49] H. Lu, R. Yu, Y. Zhu, X. He, K. Liang, and W. Cheng-Chung Chu, "Policy-driven data sharing over attribute-based encryption supporting dual membership," *J. Syst. Softw.*, vol. 188, Jun. 2022, Art. no. 111271.
- [50] (Microsoft Corp. Technol. Corp., Redmond, WA, USA). *Introduction to Azure Kubernetes Service (AKS)*. Accessed: Jun. 8, 2023. [Online]. Available: <https://learn.microsoft.com/en-us/azure/aks/intro-kubernetes>
- [51] A. Mosteiro-Sanchez, M. Barcelo, J. Astorga, and A. Urieta, "Too many options: A survey of abe libraries for developers," 2022, *arXiv:2209.12742*.
- [52] C. Padro, "Lecture notes in secret sharing," *Cryptol. ePrint Arch.*, IACR, Bellevue, WA, USA, Rep. 2012/674, 2012.
- [53] X. Qin, Z. Yang, Q. Li, H. Pan, Z. Yang, and Y. Huang, "Attribute-based encryption with outsourced computation for access control in IoTs," in *Proc. 3rd Asia Serv. Sci. Softw. Eng. Conf.*, 2022, pp. 66–73.
- [54] M. Rengo. "Project onehundredten: Cloud bursting approaches based on Kubernetes ." GitHub. Accessed: Mar. 29, 2023. [Online]. Available: <https://github.com/MRColorR/onehundredten.git>
- [55] M. Repetto, D. Striccoli, G. Piro, A. Carrega, G. Boggia, and R. Bolla, "An autonomous cybersecurity framework for next-generation digital service chains," *J. Netw. Syst. Manage.*, vol. 29, no. 4, p. 37, 2021.
- [56] S. Risco, G. Moltó, D. M. Naranjo, and I. Blanquer, "Serverless workflows for containerised applications in the cloud continuum," *J. Grid Comput.*, vol. 19, no. 3, Jul. 2021.
- [57] R. Ronan. "Simplify your hybrid/multi-cluster, multi-cloud Kubernetes deployments with KubeSlice." Developers Blog. Apr. 2022. [Online]. Available: <https://blogs.oracle.com/developers/post/simplify-your-hybridmulti-cluster-multi-cloud-kubernetes-deployments-with-kubeslice>
- [58] A. Sahai and B. Waters, "Fuzzy identity-based encryption," in *Proc. 24th Annu. Int. Conf. Theory Appl. Cryptogr. Techn.*, Aarhus, Denmark, 2005, pp. 457–473.
- [59] S. Sciancalepore, G. Piro, D. Calderola, G. Boggia, and G. Bianchi, "On the design of a decentralized and multiauthority access control scheme in federated and cloud-assisted cyber-physical systems," *IEEE Internet Things J.*, vol. 5, no. 6, pp. 5190–5204, Dec. 2018.
- [60] V. Sharma, "Managing multi-cloud deployments on Kubernetes with Istio, rometheus and Grafana," in *Proc. 8th Int. Conf. Adv. Comput. Commun. Syst. (ICACCS)*, 2022, pp. 525–529.
- [61] O. Tomarchio, D. Calcaterra, and G. D. Modica, "Cloud resource orchestration in the multi-cloud landscape: A systematic review of existing frameworks," *J. Cloud Comput.*, vol. 9, pp. 1–24, Sep. 2020.
- [62] S. Touw (Immuta Inc., Boston, MA, USA). *Role-Based Access Control Vs. Attribute-Based Access Control*. (Apr. 2023). [Online]. Available: <https://www.immuta.com/blog/attribute-based-access-control/>
- [63] M. Venema, G. Alpár, and J. Hoepman, "Systematizing core properties of pairing-based attribute-based encryption to uncover remaining challenges in enforcing access control in practice," *Des., Codes Cryptogr.*, vol. 91, no. 1, pp. 165–220, Sep. 2023.
- [64] Z. Wan, J. Liu, and R. H. Deng, "Hasbe: A hierarchical attribute-based solution for flexible and scalable access control in cloud computing," *IEEE Trans. Inf. Forensics Security*, vol. 7, pp. 743–754, 2012.
- [65] B. Waters, "Dual system encryption: Realizing fully secure IBE and HIBE under simple assumptions," in *Proc. 29th Annu. Int. Cryptol. Conf.*, 2009, pp. 619–636.
- [66] R. Zhang, J. Li, Y. Lu, J. Han, and Y. Zhang, "Key escrow-free attribute based encryption with user revocation," *Inf. Sci.*, vol. 600, pp. 59–72, Jul. 2022.
- [67] Y. Zhuang (Alibaba Cloud Comput. Softw. Com., Hangzhou, China). *Enhancing Self-Created Kubernetes With Cloud Elasticity to Cope With Traffic Bursts*. (Oct. 2023). [Online]. Available: https://www.alibabacloud.com/blog/enhancing-self-created-kubernetes-with-cloud-elasticity-to-cope-with-traffic-bursts_600491



MAURO FEMMINELLA received the master's and Ph.D. degrees in electronic engineering from the University of Perugia in 1999 and 2003, respectively, where he has been an Associate Professor with the Department of Engineering since July 2022. He is also the Representative of the University of Perugia in consortium CNIT. He is a coauthor of more than 120 papers in international journals and refereed international conferences. His current research interests focus on molecular communications, big data systems, and network management solutions for 5G/6G networks.



MARTINA PALMUCCI was born in Jesi, Italy, in 1997. She received the bachelor's degree in mathematics from the University of Camerino, Italy, in 2020, and the master's degree in computer engineering from the University of Perugia, Italy, in 2022.

She broadened her academic exposure through Erasmus programs with the University of Salamanca, Spain, in 2017 and the Vrije Universiteit Brussel, Belgium, in 2022. She currently works as a Research Fellow with the GARR Consortium, Rome, Italy. Her contributions center on the exploration of cryptography, reflecting her unwavering commitment to enhancing cybersecurity practices. Her primary focus lies in applied cryptography and cybersecurity.



GIANLUCA REALI received the Ph.D. degree in telecommunications from the University of Perugia in 1997, where he has been an Associate Professor with the Department of Engineering since January 2005. From 1997 to 2004, he was a Researcher with the Department of Electronic and Information Engineering, University of Perugia. In 1999, he visited the Computer Science Department, UCLA. His research activities include resource allocation over packet networks, wireless networking, network management, multimedia services, big data management, and nanoscale communications. He is also a member of CNIT.



MATTIA RENGO was born in Terni, Italy, in 1995. He received the bachelor's degree in computer and electronic engineering with a focus on the Computer Science curriculum from the University of Perugia in 2020, and the master's degree in computer and robotics engineering, specializing in data science, from the University of Perugia in 2023.

Currently working as a DevOps Specialist, he is also active on GitHub. Some of his personal projects have garnered more than a hundred stars, showcasing their impact and usefulness in the tech community. His professional and academic commitments reflect his dedicated focus on computer engineering, particularly in cloud technologies and data science.

Open Access funding provided by ‘Università degli Studi di Perugia’ within the CRUI CARE Agreement